

INTRODUCTION TO

ARTIFICIAL INTELLIGENCE

Theoretical Learning Notes



WHAT YOU'LL LEARN

-  AI Fundamentals
-  Machine Learning
-  Deep Learning
-  Neural Networks
-  LLMs
-  Responsible AI

AI DOMAINS COVERED



NLP



Computer Vision



Robotics



Business AI



Generative AI



Agentic AI

Prepared by



KASHIF MAQBOOL JOIYA

Data Scientist | AI Engineer



GitHub

<https://github.com/KashifMaqbool>

“

The journey into Artificial Intelligence begins with understanding its fundamentals

”

Table of Contents

1	<u>INTRODUCING ARTIFICIAL INTELLIGENCE (AI)</u>	28
1.1	LEARNING OBJECTIVES	28
1.2	HISTORY OF ARTIFICIAL INTELLIGENCE	28
1.3	WHAT IS ARTIFICIAL INTELLIGENCE?	29
1.4	TECHNOLOGIES DRIVING AI GROWTH	29
1.4.1	INTERNET	29
1.4.2	DISTRIBUTED COMPUTING	29
1.4.3	INTERNET OF THINGS (IoT)	29
1.4.4	SOCIAL NETWORKING	29
1.5	AI AS AUGMENTED INTELLIGENCE	29
1.6	INNATE INTELLIGENCE	30
1.7	HOW AI LEARNS	30
1.8	CATEGORIES OF AI	30
1.8.1	WEAK AI (NARROW AI)	31
1.8.2	STRONG AI (GENERALIZED AI)	31
1.8.3	SUPER AI (CONSCIOUS AI)	31
1.9	FIELDS THAT CONTRIBUTE TO AI	32
1.10	AI IN EVERYDAY LIFE	32
1.11	KEY TAKEAWAYS	33
2	<u>HUMAN INTELLIGENCE, ARTIFICIAL INTELLIGENCE, AND AUGMENTED INTELLIGENCE</u>	33
2.1	UNDERSTANDING DIFFERENT FORMS OF INTELLIGENCE	33
2.1.1	HUMAN INTELLIGENCE	33
2.1.2	ARTIFICIAL INTELLIGENCE	33
2.1.3	AUGMENTED INTELLIGENCE	34
2.2	ARTIFICIAL INTELLIGENCE VS. AUGMENTED INTELLIGENCE	34
2.3	STRENGTHS OF MACHINES	35
2.4	STRENGTHS OF HUMANS	35
2.5	WHY AUGMENTED INTELLIGENCE IS THE SWEET SPOT	35
2.6	KEY TAKEAWAYS	36
3	<u>INTRODUCING GENERATIVE AI AND ITS USE CASES</u>	36
3.1	LEARNING OBJECTIVES	36
3.2	INTRODUCTION TO GENERATIVE AI	36
3.3	LARGE LANGUAGE MODELS (LLMs)	37
3.4	CAPABILITIES OF GENERATIVE AI	37
3.5	ECONOMIC IMPACT OF GENERATIVE AI	37
3.6	DOMAIN-SPECIFIC USE CASES OF GENERATIVE AI	37

3.6.1	MARKETING	37
3.6.2	CREATIVE FIELDS	38
3.6.3	PRODUCT DEVELOPMENT	38
3.7	INDUSTRY-SPECIFIC USE CASES OF GENERATIVE AI	38
3.7.1	HEALTHCARE AND PRECISION MEDICINE	38
3.7.2	GAMING	38
3.7.3	FASHION	38
3.7.4	EDUCATION AND TRAINING	38
3.8	KEY TAKEAWAYS	39

4 INTRODUCTION **39**

4.1	OBJECTIVES	39
4.2	AI TYPES	39
4.2.1	DIAGNOSTIC/DESCRIPTIVE AI	39
4.2.1.1	Capabilities	40
4.2.2	PREDICTIVE AI	40
4.2.2.1	Capabilities	40
4.2.3	PRESCRIPTIVE AI	40
4.2.3.1	Capabilities	40
4.2.4	GENERATIVE/COGNITIVE AI	40
4.2.4.1	Capabilities	40
4.2.5	REACTIVE AI	41
4.2.5.1	Capabilities	41
4.2.6	LIMITED MEMORY AI	41
4.2.6.1	Capabilities	41
4.2.7	THEORY OF MIND AI	41
4.2.7.1	Capabilities	41
4.2.8	SELF-AWARE AI	41
4.2.8.1	Capabilities	41
4.2.9	NARROW AI (WEAK AI)	42
4.2.9.1	Capabilities	42
4.2.10	GENERAL AI (STRONG AI)	42
4.2.10.1	Capabilities	42
4.3	SUMMARY	42

5 HOW TRADITIONAL AI DIFFERS FROM GENERATIVE AI **43**

5.1	INTRODUCTION	43
5.2	TRADITIONAL AI ARCHITECTURE	43
5.2.1	REPOSITORY	43
5.2.2	ANALYTICS PLATFORM	43
5.2.3	APPLICATION LAYER	43
5.3	THE ROLE OF THE FEEDBACK LOOP	44
5.4	GENERATIVE AI ARCHITECTURE	45
5.5	LARGE LANGUAGE MODELS (LLMs)	45

5.6	PROMPTING AND TUNING	45
5.7	APPLICATION LAYER IN GENERATIVE AI	46
5.8	FEEDBACK LOOP IN GENERATIVE AI	46
5.9	TRADITIONAL AI VS. GENERATIVE AI	46
5.10	KEY TAKEAWAYS	46

6 ARE WE THERE YET? UNDERSTANDING THE CURRENT STATE OF ARTIFICIAL INTELLIGENCE **47**

6.1	INTRODUCTION	47
6.2	WHAT IS ARTIFICIAL INTELLIGENCE?	47
6.3	ARTIFICIAL GENERAL INTELLIGENCE (AGI)	47
6.4	DOES SOLVING MATH INSTANTLY MEAN INTELLIGENCE?	47
6.5	IS MEMORIZATION THE SAME AS INTELLIGENCE?	48
6.6	CHESS AS A MEASURE OF INTELLIGENCE	48
6.6.1	IBM DEEP BLUE	48
6.7	THE TURING TEST	49
6.7.1	HOW THE TURING TEST WORKS	49
6.7.2	HAS THE TURING TEST BEEN PASSED?	49
6.8	CHATGPT AND MODERN CONVERSATIONAL AI	49
6.9	THE MOVING FINISH LINE OF AI	49
6.10	SPECIALIZED AI VERSUS GENERAL INTELLIGENCE	50
6.11	THE RAPID DEVELOPMENT OF AI	50
6.12	ARE WE THERE YET?	50
6.13	EVERYDAY USES OF AI	51
6.13.1	SELF-DRIVING CAR	51
6.13.2	AI IMAGE GENERATION	51

7 AI IN DAILY LIFE **51**

7.1	COMMON USE CASES OF AI IN DAILY LIFE	51
7.2	BENEFITS OF AI IN DAILY LIFE	52
7.3	AI FOR AUTOMATION AND CONVENIENCE	52
7.4	AI-POWERED PERSONALIZED RECOMMENDATIONS	52
7.4.1	STREAMING PLATFORMS	52
7.4.2	SOCIAL MEDIA	53
7.4.3	E-COMMERCE	53
7.5	AI FOR SECURITY	53
7.6	AI IN HEALTHCARE	53
7.7	AI IN SMART DEVICES	54
7.7.1	SMARTPHONE CAMERAS	54
7.7.2	NAVIGATION APPLICATIONS	54
7.7.3	KEYBOARDS AND MESSAGING APPLICATIONS	54
7.8	SUMMARY	54

8 AI CHATBOTS AND SMART ASSISTANTS **55**

8.1	OVERVIEW	55
8.2	EVOLUTION OF AI CHATBOTS	55
8.3	EXAMPLES OF AI CHATBOT PLATFORMS AND SMART ASSISTANTS	55
8.3.1	AI CHATBOT PLATFORMS	55
8.3.2	SMART ASSISTANTS	56
8.3.3	GENERATIVE AI CHATBOTS	56
8.4	HOW AI CHATBOTS WORK	56
8.4.1	UNDERSTANDING USER INPUT	56
8.4.2	INTENT AND CONTEXT RECOGNITION	56
8.4.3	DIALOGUE MANAGEMENT	56
8.4.4	RESPONSE GENERATION	57
8.5	EXAMPLE: WEATHER INFORMATION	57
8.6	BENEFITS OF AI CHATBOTS	57
8.7	APPLICATIONS OF AI CHATBOTS	57
8.7.1	CUSTOMER SERVICE	57
8.7.2	E-COMMERCE	57
8.7.3	HEALTHCARE	58
8.7.4	EDUCATION	58
8.8	FUTURE TRENDS OF AI CHATBOTS	58
8.9	KEY TAKEAWAYS	58

9 UNDERSTANDING CHATBOTS **59**

9.1	INTRODUCTION	59
9.2	EXAMPLE: FLOWER SHOP CHATBOT	59
9.3	EXAMPLE: BANKING CHATBOT	60
9.3.1	CUSTOMER REQUEST	60
9.3.2	AUTHENTICATION PROCESS	60
9.3.3	COMMUNICATION WITH THE BANK	60
9.4	WAYS WE CAN INTERACT WITH A CHATBOT	61
9.4.1	WEBSITE CHATBOT	61
9.4.2	VOICE-BASED CHATBOT	61
9.4.3	INSTANT MESSAGING PLATFORMS	61
9.5	HOW A CHATBOT WORKS BEHIND THE SCENES	62
9.5.1	STEP 1: USER SENDS A REQUEST	62
9.5.2	STEP 2: REQUEST IS SENT TO THE CLOUD	62
9.5.3	STEP 3: CHATBOT CHECKS BUSINESS SYSTEMS	62
9.5.4	STEP 4: ORDER CONFIRMATION	63
9.6	TECHNOLOGIES USED IN CHATBOTS	63
9.7	BENEFITS OF USING CHATBOTS	64
9.7.1	SAVES TIME	64
9.7.2	SIMPLICITY	64
9.7.3	FAST TIME TO LAUNCH	64
9.8	KEY TAKEAWAYS	65

10 AI ASSISTANT: USER CONVERSATIONS **65**

10.1	LEARNING OBJECTIVES	65
10.2	INTRODUCTION	65
10.3	EXERCISE 1: FINANCIAL AI ADVISOR	65
10.3.1	PROMPT INSTRUCTIONS	65
10.3.2	PROMPT 1	66
10.3.3	PROMPT 2	66
10.4	EXERCISE 2: HEALTH AI ASSISTANT	66
10.4.1	PROMPT INSTRUCTIONS	66
10.4.2	PROMPT 1	66
10.4.3	PROMPT 2	66
10.5	SUMMARY	66

11 APPLICATIONS OF AI IN DIFFERENT INDUSTRIES **66**

11.1	LEARNING OBJECTIVES	66
11.2	AI ADOPTION ACROSS INDUSTRIES	67
11.3	PREDICTIVE MAINTENANCE	67
11.4	AI IN MANUFACTURING	67
11.4.1	AI-DRIVEN ROBOTICS AND AUTOMATION	67
11.4.2	IMAGE RECOGNITION FOR QUALITY INSPECTION	68
11.4.3	ENERGY OPTIMIZATION	68
11.4.4	FOOD AND BEVERAGE MANUFACTURING	68
11.5	AI IN HEALTHCARE	68
11.5.1	MEDICAL IMAGING ANALYSIS	68
11.5.2	PREDICTIVE ANALYTICS	69
11.5.3	OPERATIONAL EFFICIENCY	69
11.5.4	DRUG DISCOVERY AND DEVELOPMENT	69
11.6	AI IN FINANCE	70
11.6.1	AI-POWERED CUSTOMER SUPPORT	70
11.6.2	INVESTMENT ANALYSIS	70
11.6.3	RISK MANAGEMENT	70
11.6.4	ROBO-ADVISORS	70
11.6.5	REAL-WORLD EXAMPLE: BANK OF AMERICA – ERICA	70
11.7	AI IN RETAIL	71
11.7.1	PERSONALIZED PRODUCT RECOMMENDATIONS	71
11.7.2	DEMAND FORECASTING	71
11.7.3	PERSONALIZED MARKETING	71
11.7.4	AI-POWERED CASHIERLESS STORES	72
11.8	FUTURE OF AI	72
11.9	KEY TAKEAWAYS	72

12 GENERATIVE AI TOOLS AND APPLICATIONS **72**

12.1	OVERVIEW	72
12.2	GENERATIVE AI FOR LANGUAGE	73
12.3	LARGE LANGUAGE MODELS AND THEIR CAPABILITIES	73

12.3.1	OPENAI GPT MODELS	73
12.3.2	GOOGLE PALM MODELS	73
12.3.3	GOOGLE GEMINI MODELS	74
12.3.4	OTHER MAJOR LARGE LANGUAGE MODELS	74
12.4	GENERATIVE AI FOR VISUAL ARTS AND DESIGN	74
12.4.1	STABLE DIFFUSION	74
12.4.2	DALL·E	74
12.4.3	STYLEGAN	74
12.4.4	SUPER RESOLUTION	75
12.5	GENERATIVE AI FOR VOICE AND MUSIC GENERATION	75
12.5.1	MURPH	75
12.5.2	OPENAI WHISPER	75
12.5.3	AI MUSIC GENERATION	75
12.5.4	JUKEDECK	75
12.5.5	AMPER MUSIC	75
12.5.6	AIVA	75
12.6	GENERATIVE AI FOR VIDEO GENERATION	76
12.6.1	GOOGLE'S IMOGEN VIDEO	76
12.6.2	OPENAI SORA	76
12.7	ADOPTION OF GENERATIVE AI BY LEADING COMPANIES	76
12.7.1	GOOGLE	76
12.7.2	SALESFORCE	76
12.7.3	ADOBE	77
12.7.4	IBM	77
12.8	KEY TAKEAWAYS	77
13	<u>MACHINE LEARNING USE CASES IN EVERYDAY LIFE</u>	<u>77</u>
13.1	NATURAL LANGUAGE PROCESSING (NLP)	78
13.2	CUSTOMER SERVICE	78
13.3	VOICE ASSISTANTS AND SPEECH RECOGNITION	78
13.4	MOBILE APPLICATIONS AND SMARTPHONES	78
13.5	FINANCIAL TRANSACTIONS AND FRAUD DETECTION	79
13.6	CYBERSECURITY	79
13.7	TRANSPORTATION	80
13.8	EMAIL SERVICES	80
13.9	HEALTHCARE	80
13.10	MARKETING AND SALES	81
13.11	MACHINE LEARNING VS. ARTIFICIAL GENERAL INTELLIGENCE (AGI)	81
14	<u>GENERATIVE AI TOOLS IN ACTION</u>	<u>81</u>
14.1	LEARNING OBJECTIVES	81
14.2	INTRODUCTION	82
14.3	EXPLORING GOOGLE GEMINI	82
14.3.1	SCENARIO: LEARNING A NEW GAME	82

14.3.2	SCENARIO: WRITING A BLOG POST ABOUT THE REVIVAL OF BOARD GAMES	82
14.4	EXPERIMENTING WITH CHATGPT	82
14.4.1	SCENARIO: SUMMARIZING TEXT	82
14.4.2	SCENARIO: PLANNING A WEEKEND TRIP	83
14.5	EXPERIMENTING WITH MICROSOFT COPILOT	83
14.5.1	SCENARIO: CREATING THE PLOT FOR A FICTIONAL STORY	83
14.6	SUMMARY	83
15	MODULE SUMMARY	84
15.1	ARTIFICIAL INTELLIGENCE (AI)	84
15.2	GENERATIVE AI	84
15.2.1	DOMAIN-SPECIFIC USE CASES OF GENERATIVE AI	84
15.2.2	INDUSTRY-SPECIFIC USE CASES OF GENERATIVE AI	85
15.3	BENEFITS OF AI	85
15.4	GENERATIVE AI USE CASES	85
15.4.1	TEXT GENERATION	85
15.4.2	IMAGE GENERATION	85
15.4.3	VOICE AND MUSIC GENERATION	85
15.4.4	VIDEO GENERATION	85
15.5	AI CHATBOTS AND SMART ASSISTANTS	86
15.5.1	BENEFITS OF AI CHATBOTS	86
15.5.2	APPLICATIONS OF AI CHATBOTS	86
15.6	AI'S IMPACT ON INDUSTRIES	86
16	COGNITIVE COMPUTING	87
16.1	OVERVIEW	87
16.2	HUMAN DECISION-MAKING PROCESS	87
16.3	CORE ELEMENTS OF COGNITIVE COMPUTING	88
16.3.1	PERCEPTION	88
16.3.2	LEARNING	88
16.3.3	REASONING	88
16.4	HOW COGNITIVE COMPUTING WORKS	88
16.5	BENEFITS OF COGNITIVE COMPUTING	89
16.6	APPLICATIONS OF COGNITIVE COMPUTING	89
16.7	REAL-WORLD EXAMPLES	89
16.7.1	IBM WATSON	89
16.7.2	GOOGLE	89
16.7.3	AMAZON	90
16.7.4	JPMORGAN CHASE AND WELLS FARGO	90
16.8	KEY TAKEAWAYS	90
17	TERMINOLOGIES AND RELATED CONCEPTS OF ARTIFICIAL INTELLIGENCE (AI)	91

17.1	OVERVIEW	91
17.1.1	REAL-WORLD EXAMPLE	91
17.2	ARTIFICIAL INTELLIGENCE (AI)	91
17.2.1	CATEGORIES OF ARTIFICIAL INTELLIGENCE	91
17.3	MACHINE LEARNING	92
17.4	DEEP LEARNING	92
17.5	NEURAL NETWORKS	92
17.6	KEY TAKEAWAYS	93

18	MACHINE LEARNING FUNDAMENTALS	93
-----------	--------------------------------------	-----------

18.1	INTRODUCTION	93
18.2	UNDERSTANDING MACHINE LEARNING WITH A HEART FAILURE PREDICTION EXAMPLE	93
18.3	TRADITIONAL PROGRAMMING VS. MACHINE LEARNING	94
18.3.1	TRADITIONAL PROGRAMMING	94
18.3.2	MACHINE LEARNING	94
18.4	HOW MACHINE LEARNING WORKS	94
18.5	TYPES OF MACHINE LEARNING	95
18.6	SUPERVISED LEARNING	95
18.7	UNSUPERVISED LEARNING	95
18.7.1	NETWORK TRAFFIC EXAMPLE	95
18.8	REINFORCEMENT LEARNING	96
18.8.1	APPLICATIONS	96
18.9	MACHINE LEARNING MODEL	96
18.10	KEY TAKEAWAYS	96

19	MACHINE LEARNING: TECHNIQUES AND TRAINING	97
-----------	--	-----------

19.1	LEARNING OBJECTIVES	97
19.2	MACHINE LEARNING CATEGORIES	97
19.2.1	SUPERVISED LEARNING	98
19.2.2	UNSUPERVISED LEARNING	98
19.2.3	REINFORCEMENT LEARNING	98
19.3	SUPERVISED LEARNING TASKS	98
19.3.1	REGRESSION	98
19.3.2	NEURAL NETWORKS	99
19.3.3	CLASSIFICATION	99
19.4	FEATURES AND DATA POINTS	99
19.5	MODEL TRAINING	100
19.5.1	EXAMPLE: SPAM EMAIL CLASSIFICATION	100
19.6	DATASET SPLITTING	100
19.6.1	TRAINING SET	100
19.6.2	VALIDATION SET	100
19.6.3	TEST SET	101
19.7	MODEL EVALUATION METRICS	101
19.8	KEY TAKEAWAYS	101

20	DEEP LEARNING FUNDAMENTALS AND APPLICATIONS	101
20.1	INTRODUCTION	101
20.2	HOW DEEP LEARNING WORKS	102
20.3	DEVELOPING AND TRAINING DEEP LEARNING MODELS	102
20.4	ADVANTAGES OF DEEP LEARNING	103
20.5	APPLICATIONS OF DEEP LEARNING	103
20.6	KEY TAKEAWAYS	103
21	NEURAL NETWORKS	103
21.1	OVERVIEW OF NEURAL NETWORKS	104
21.2	STRUCTURE OF A NEURAL NETWORK	104
21.2.1	INPUT LAYER	104
21.2.2	HIDDEN LAYERS	104
21.2.3	OUTPUT LAYER	104
21.3	HOW NEURAL NETWORKS LEARN	104
21.3.1	FORWARD PASS (FORWARD PROPAGATION)	104
21.3.2	BACKWARD PASS (BACKPROPAGATION)	105
21.4	TYPES OF NEURAL NETWORKS	105
21.4.1	PERCEPTRON NEURAL NETWORKS	105
21.4.2	FEED-FORWARD NEURAL NETWORKS	105
21.4.3	DEEP FEED-FORWARD NEURAL NETWORKS	106
21.4.4	MODULAR NEURAL NETWORKS	106
21.4.5	CONVOLUTIONAL NEURAL NETWORKS (CNNs)	106
21.4.6	RECURRENT NEURAL NETWORKS (RNNs)	106
21.5	KEY POINTS	106
22	MACHINE LEARNING VS. DEEP LEARNING	107
22.1	INTRODUCTION	107
22.2	MACHINE LEARNING	107
22.2.1	PIZZA ORDERING EXAMPLE	107
22.2.2	ASSIGNING WEIGHTS	108
22.2.3	CALCULATING THE OUTPUT	108
22.3	WHAT MAKES DEEP LEARNING DIFFERENT?	108
22.4	HUMAN INTERVENTION IN MACHINE LEARNING	108
22.5	DEEP LEARNING AND UNSUPERVISED LEARNING	109
22.6	FEEDFORWARD NEURAL NETWORKS	109
22.7	BACKPROPAGATION	110
22.8	KEY DIFFERENCES BETWEEN MACHINE LEARNING AND DEEP LEARNING	110
22.9	SUMMARY	110
23	GENERATIVE AI MODELS	111
23.1	OVERVIEW	111

23.2	TYPES OF GENERATIVE AI MODELS	111
23.3	VARIATIONAL AUTOENCODERS (VAEs)	111
23.3.1	APPLICATIONS OF VAES	112
23.3.2	EXAMPLE	112
23.4	GENERATIVE ADVERSARIAL NETWORKS (GANs)	112
23.4.1	APPLICATIONS OF GANS	112
23.4.2	EXAMPLE	113
23.5	AUTOREGRESSIVE MODELS	113
23.5.1	EXAMPLES	113
23.6	TRANSFORMERS	113
23.6.1	EXAMPLES	113
23.7	CATEGORIES OF GENERATIVE AI MODELS	114
23.7.1	UNIMODAL MODELS	114
23.7.1.1	Example	114
23.7.2	MULTIMODAL MODELS	114
23.7.2.1	Example: DALL·E	114
23.7.2.2	Example: Meta's ImageBind	115
23.8	FUTURE OF GENERATIVE AI MODELS	115
23.9	KEY TAKEAWAYS	115
24	<u>FOUNDATION MODELS AND LARGE LANGUAGE MODELS (LLMS)</u>	116
24.1	INTRODUCTION	116
24.2	FOUNDATION MODELS	116
24.3	WHY FOUNDATION MODELS CAN PERFORM MULTIPLE TASKS	116
24.4	FOUNDATION MODELS AND GENERATIVE AI	117
24.5	TUNING A FOUNDATION MODEL	117
24.6	PROMPTING AND PROMPT ENGINEERING	117
24.6.1	EXAMPLE: SENTIMENT CLASSIFICATION	118
24.7	ADVANTAGES OF FOUNDATION MODELS	118
24.7.1	HIGH PERFORMANCE	118
24.7.2	IMPROVED PRODUCTIVITY	118
24.8	DISADVANTAGES OF FOUNDATION MODELS	118
24.8.1	HIGH COMPUTE COSTS	118
24.8.2	TRUSTWORTHINESS CHALLENGES	119
24.9	IBM'S RESEARCH AND INNOVATIONS	119
24.10	FOUNDATION MODELS BEYOND LANGUAGE	119
24.10.1	VISION MODELS	120
24.10.2	CODE MODELS	120
24.10.3	IBM PRODUCTS USING FOUNDATION MODELS	120
24.11	FOUNDATION MODELS FOR SCIENTIFIC RESEARCH	120
24.11.1	CHEMISTRY	120
24.11.2	CLIMATE CHANGE	120
24.12	SUMMARY	121
25	<u>ARTIFICIAL INTELLIGENCE, MACHINE LEARNING, DEEP LEARNING, FOUNDATION MODELS, LARGE LANGUAGE MODELS, AND GENERATIVE AI</u>	121

25.1	ARTIFICIAL INTELLIGENCE (AI)	121
25.2	MACHINE LEARNING (ML)	122
25.2.1	CORE CATEGORIES OF MACHINE LEARNING	122
25.3	DEEP LEARNING (DL)	122
25.3.1	IMPORTANT NOTE	123
25.4	FOUNDATION MODELS	123
25.5	LARGE LANGUAGE MODELS (LLMs)	124
25.5.1	UNDERSTANDING THE TERM "LLM"	124
25.5.2	COMMON LANGUAGE TASKS PERFORMED BY LLMs	124
25.6	OTHER TYPES OF FOUNDATION MODELS	125
25.7	GENERATIVE AI	125
25.8	RELATIONSHIP BETWEEN THE AI TERMS	125

26 NATURAL LANGUAGE PROCESSING, SPEECH TECHNOLOGY, AND COMPUTER VISION **126**

26.1	OVERVIEW	126
26.2	NATURAL LANGUAGE PROCESSING (NLP)	126
26.2.1	NLP MARKET GROWTH	126
26.2.2	HOW NLP WORKS	126
26.3	SPEECH TECHNOLOGY	127
26.3.1	SPEECH-TO-TEXT (STT)	127
26.3.2	HOW STT WORKS	127
26.3.3	APPLICATIONS OF STT	127
26.4	TEXT-TO-SPEECH (TTS)	128
26.4.1	HOW TTS WORKS	128
26.5	INTEGRATION OF STT AND TTS	128
26.5.1	APPLICATIONS	128
26.5.1.1	Translation Services	128
26.5.1.2	Smart Home Devices	128
26.5.1.3	Customer Support	128
26.5.1.4	Accessibility	129
26.6	COMPUTER VISION	129
26.7	NEURAL NETWORKS IN COMPUTER VISION	129
26.7.1	IMAGE CLASSIFICATION	129
26.7.2	OBJECT DETECTION	129
26.7.3	IMAGE SEGMENTATION	130
26.8	INDUSTRY APPLICATIONS OF COMPUTER VISION	130
26.8.1	RETAIL	130
26.8.2	MANUFACTURING	130
26.8.3	AGRICULTURE	130
26.9	KEY TAKEAWAYS	130

27 WHAT IS NATURAL LANGUAGE PROCESSING (NLP)? **131**

27.1	UNSTRUCTURED TEXT VS. STRUCTURED DATA	131
-------------	--	------------

27.2	THE ROLE OF NLP	131
27.3	APPLICATIONS OF NLP	132
27.3.1	MACHINE TRANSLATION	132
27.3.2	VIRTUAL ASSISTANTS AND CHATBOTS	132
27.3.3	SENTIMENT ANALYSIS	132
27.3.4	SPAM DETECTION	133
27.4	HOW NLP WORKS	133
27.5	TOKENIZATION	133
27.6	STEMMING	134
27.7	LEMMATIZATION	134
27.8	PART-OF-SPEECH (POS) TAGGING	134
27.9	NAMED ENTITY RECOGNITION (NER)	135
27.10	NLP PROCESSING PIPELINE	135
27.11	KEY TAKEAWAYS	135
28	<u>SELF-DRIVING CARS AND THE ROLE OF COMPUTER VISION</u>	136
28.1	OVERVIEW	136
28.2	3D OBJECT DETECTION IN SELF-DRIVING CARS	136
28.3	WHY COMPUTER VISION IS CLOSELY ASSOCIATED WITH SELF-DRIVING CARS	137
28.4	LIMITATIONS OF HUMAN VISION	137
28.5	HOW COMPUTER VISION OVERCOMES HUMAN LIMITATIONS	137
28.6	THE FUTURE OF SELF-DRIVING CARS	137
28.7	REMAINING CHALLENGES	138
29	<u>AI AND CLOUD COMPUTING, EDGE COMPUTING, AND IOT</u>	138
29.1	OVERVIEW	138
29.2	INTERNET OF THINGS (IoT)	138
29.2.1	HOW IoT DEVICES WORK	139
29.3	CLOUD COMPUTING	139
29.3.1	EXAMPLE OF CLOUD COMPUTING	139
29.4	AI IN CLOUD COMPUTING	139
29.4.1	EXAMPLE: AI IN CLOUD-BASED EMAIL SERVICES	139
29.5	EDGE COMPUTING	140
29.5.1	EXAMPLE: SMART THERMOSTAT	140
29.6	EDGE AI	140
29.6.1	EXAMPLE OF EDGE AI	140
29.7	HOW AI, IoT, CLOUD COMPUTING, AND EDGE COMPUTING WORK TOGETHER	140
29.7.1	EXAMPLE: FITNESS TRACKER	140
29.7.2	ROLE OF CLOUD COMPUTING	141
29.7.3	ROLE OF CLOUD-BASED AI	141
29.8	REAL-WORLD APPLICATIONS	141
29.9	KEY TAKEAWAYS	142
30	<u>MODULE SUMMARY</u>	142

30.1	MACHINE LEARNING	142
30.1.1	TYPES OF MACHINE LEARNING	142
30.1.1.1	Supervised Learning	142
30.1.1.2	Unsupervised Learning	142
30.1.1.3	Reinforcement Learning	143
30.2	MODEL TRAINING PROCESS	143
30.3	DEEP LEARNING	143
30.4	NEURAL NETWORKS	143
30.4.1	TYPES OF NEURAL NETWORKS	143
30.5	GENERATIVE AI MODEL ARCHITECTURES	144
30.5.1	VARIATIONAL AUTOENCODERS (VAEs)	144
30.5.2	GENERATIVE ADVERSARIAL NETWORKS (GANS)	144
30.5.3	AUTOREGRESSIVE MODELS	144
30.5.4	TRANSFORMERS	144
30.6	TYPES OF AI MODELS	144
30.6.1	UNIMODAL MODELS	145
30.6.2	MULTIMODAL MODELS	145
30.7	COGNITIVE COMPUTING	145
30.8	NATURAL LANGUAGE PROCESSING (NLP)	145
30.9	SPEECH TECHNOLOGIES	145
30.9.1	SPEECH-TO-TEXT (STT)	145
30.9.2	TEXT-TO-SPEECH (TTS)	145
30.10	COMPUTER VISION	145
30.11	INTERNET OF THINGS (IoT)	146
30.12	CLOUD COMPUTING	146
30.13	EDGE COMPUTING	146
30.14	INTEGRATION OF AI, IoT, CLOUD COMPUTING, AND EDGE COMPUTING	146
30.14.1	REAL-WORLD APPLICATIONS	146

31 AI AGENTS **146**

31.1	INTRODUCTION	146
31.2	WHAT ARE AI AGENTS?	147
31.3	HOW AI AGENTS WORK	147
31.3.1	PERCEPTION	147
31.3.2	UNDERSTANDING	147
31.3.3	DECISION-MAKING	147
31.3.4	ACTION	148
31.3.5	LEARNING	148
31.4	CHARACTERISTICS OF AI AGENTS	148
31.4.1	SOCIAL ABILITY	148
31.4.2	AUTONOMY	148
31.4.3	REACTIVENESS	148
31.4.4	PROACTIVENESS	148
31.5	APPLICATIONS OF AI AGENT CHARACTERISTICS	149
31.5.1	HEALTHCARE	149
31.5.2	SELF-DRIVING CARS	149

31.5.3	SMART THERMOSTATS	149
31.5.4	PREDICTIVE MAINTENANCE	149
31.6	MULTI-AGENT SYSTEMS	149
31.7	APPLICATIONS OF MULTI-AGENT SYSTEMS	150
31.7.1	ONLINE MARKETPLACES	150
31.7.2	ROBOTIC COORDINATION	150
31.7.3	TRAFFIC MANAGEMENT	150
31.8	AI AGENTS IN REAL-WORLD PRODUCTS AND SERVICES	150
31.8.1	GOOGLE	150
31.8.1.1	YouTube	150
31.8.1.2	Gmail	151
31.8.1.3	Google Maps	151
31.8.2	AMAZON	151
31.8.2.1	Alexa	151
31.8.2.2	Amazon.com	151
31.8.2.3	Amazon Web Services (AWS)	151
31.9	FUTURE OF AI AGENTS	152
31.10	KEY TAKEAWAYS	152

32 AI AGENTS AND THE SHIFT FROM COMPOUND AI SYSTEMS TO AGENTIC SYSTEMS **152**

32.1	INTRODUCTION	152
32.2	FROM MONOLITHIC MODELS TO COMPOUND AI SYSTEMS	153
32.2.1	EXAMPLE: CHECKING REMAINING VACATION DAYS	153
32.3	HOW A COMPOUND AI SYSTEM SOLVES THE PROBLEM	153
32.4	CHARACTERISTICS OF COMPOUND AI SYSTEMS	154
32.5	RETRIEVAL-AUGMENTED GENERATION (RAG)	154
32.6	FROM COMPOUND AI SYSTEMS TO AI AGENTS	155
32.7	FAST THINKING VS. SLOW THINKING	155
32.7.1	FAST THINKING	155
32.7.2	SLOW THINKING	155
32.8	COMPONENTS OF LLM AGENTS	156
32.8.1	REASONING	156
32.8.2	ACTING THROUGH TOOLS	156
32.8.3	MEMORY	156
32.9	REACT AGENTS	157
32.9.1	REACT WORKFLOW	157
32.10	EXAMPLE: PLANNING A VACATION WITH AN AI AGENT	157
32.10.1	DETERMINE AVAILABLE VACATION DAYS	157
32.10.2	ESTIMATE TIME IN THE SUN	158
32.10.3	RETRIEVE HEALTH RECOMMENDATIONS	158
32.10.4	PERFORM CALCULATIONS	158
32.11	COMPOUND AI SYSTEMS VS. AGENTIC SYSTEMS	158
32.11.1	WHEN A PROGRAMMATIC APPROACH IS BETTER	158
32.11.2	WHEN AN AGENTIC APPROACH IS BETTER	159
32.12	CURRENT STATE OF AGENTIC SYSTEMS	159

32.13	COMPARISON OF MONOLITHIC, COMPOUND, AND AGENTIC AI SYSTEMS	159
32.13.1	ARCHITECTURE COMPARISON	159
32.13.2	SUMMARY OF CORE DIFFERENCES	161

33 ROBOTICS AND AUTOMATION **161**

33.1	INTRODUCTION	161
33.2	ROBOTICS AND AUTOMATION IN INDUSTRY	161
33.3	ROBOTICS	162
33.3.1	SENSORS	162
33.3.2	ACTUATORS	162
33.3.3	CONTROLLERS	162
33.4	AI TECHNOLOGIES USED IN ROBOTICS	162
33.5	EXAMPLES OF AI-POWERED ROBOTS	163
33.5.1	ROBOT VACUUM CLEANER	163
33.5.2	SMART LAWN MOWER	163
33.5.3	VOICE-ACTIVATED SMART SPEAKER ROBOT	163
33.6	SELF-OPERATING ROBOTS	163
33.6.1	SPACE EXPLORATION	164
33.6.2	DEEP SEA EXPLORATION	164
33.7	COLLABORATIVE ROBOTS (COBOTS)	164
33.8	APPLICATIONS OF ROBOTICS IN DIFFERENT INDUSTRIES	164
33.8.1	HEALTHCARE	164
33.8.2	AGRICULTURE	164
33.8.3	RETAIL	165
33.9	ROBOTIC PROCESS AUTOMATION (RPA)	165
33.9.1	DEFINITION	165
33.9.2	APPLICATIONS OF RPA	165
33.9.2.1	Finance and Accounting	165
33.9.2.2	Human Resources	166
33.10	KEY TAKEAWAYS	166

34 TRANSFORMING BUSINESSES THROUGH AI **166**

34.1	AI AS A BUSINESS TRANSFORMATION TOOL	166
34.2	AI FOR WORKFLOW AUTOMATION	166
34.3	AI IN CUSTOMER SERVICE	167
34.3.1	REAL-WORLD EXAMPLE: AIRHELP	167
34.4	AI IN RECRUITMENT	167
34.4.1	REAL-WORLD EXAMPLE: HILTON	167
34.5	AI IN ACCOUNTING	168
34.5.1	REAL-WORLD EXAMPLE: EY	168
34.6	AI FOR DATA ANALYSIS AND DECISION MAKING	168
34.7	AI IN MARKETING AND SALES	169
34.7.1	REAL-WORLD EXAMPLE: AMAZON	169
34.8	AI IN MANUFACTURING	169

34.8.1	REAL-WORLD EXAMPLE: BMW	169
34.9	AI IN HEALTHCARE AND RESEARCH	170
34.9.1	REAL-WORLD EXAMPLE: IBM	170
34.10	AI FOR CREATIVITY AND INNOVATION	170
34.10.1	REAL-WORLD EXAMPLE: CANVA	170
34.11	BENEFITS OF AI IN BUSINESS	171
34.12	KEY TAKEAWAYS	171
35	<u>RISE OF GENERATIVE AI FOR BUSINESS</u>	171
35.1	LEARNING OBJECTIVES	171
35.2	THE RISE OF GENERATIVE AI IN BUSINESS	171
35.3	GENERATIVE AI BEYOND THE DIGITAL WORLD	172
35.4	GENERATIVE AI ADOPTION ACROSS INDUSTRIES	172
35.5	HOW GENERATIVE AI IS TRANSFORMING BUSINESS	173
35.6	CONTENT GENERATION	173
35.6.1	REAL-WORLD EXAMPLE: PERSADO AND VANGUARD	173
35.7	DATA ANALYSIS AND BUSINESS INSIGHTS	173
35.7.1	REAL-WORLD EXAMPLE: SALESFORCE	174
35.8	CUSTOMER SERVICE	174
35.8.1	REAL-WORLD EXAMPLE: SEPHORA	174
35.9	PRODUCT DEVELOPMENT	175
35.9.1	REAL-WORLD EXAMPLE: NIKE	175
35.10	GENERATIVE AI FOR STARTUPS	175
35.10.1	REAL-WORLD EXAMPLE: STITCH FIX	176
35.11	FUTURE OF GENERATIVE AI IN BUSINESS	176
35.12	KEY TAKEAWAYS	176
36	<u>GENERATIVE AI FOR BUSINESS TRANSFORMATION</u>	177
36.1	LEARNING OBJECTIVES	177
36.2	INTRODUCTION	177
36.3	MARKETING AND TECHNOLOGY TRENDS	177
36.3.1	ANALYZE RECENT E-COMMERCE MARKET TRENDS	177
36.3.2	ANALYZE RECENT TECHNOLOGICAL INNOVATIONS	177
36.3.3	IMPLEMENTING AR/VR IN AN E-COMMERCE PLATFORM	178
36.3.4	BEST PRACTICES FOR IMPLEMENTING AR/VR	178
36.4	AI-DRIVEN DYNAMIC PRICING	178
36.4.1	GENERATE A DYNAMIC PRICING IMPLEMENTATION PLAN	178
36.4.2	KEY STEPS FOR IMPLEMENTING DYNAMIC PRICING	178
36.4.3	IDENTIFY DATA REQUIRED FOR DYNAMIC PRICING	178
36.5	SUMMARY	179
37	<u>FIRE, GENERATIVE AI, AND THE FOUNDATION OF AI VALUE CREATION</u>	179
37.1	FIRE AS THE FOUNDATION OF HUMAN PROGRESS	179

37.2	A LANDMARK MOMENT FOR GENERATIVE AI	179
37.3	TRADITIONAL AI	180
37.4	WHAT MAKES GENERATIVE AI DIFFERENT	180
37.5	FOUNDATION MODELS	180
37.6	ONE FOUNDATION MODEL, MANY APPLICATIONS	181
37.7	THREE MODES OF AI CONSUMPTION	181
37.7.1	EMBEDDED AI	181
37.7.2	API-BASED AI	182
37.8	PLATFORM AI: BECOMING OUR OWN AI VALUE CREATOR	183
37.9	FOUNDATION MODELS AS THE BASIS FOR AI VALUE CREATION	184
37.10	RISKS OF LOSING CONTROL OVER AI	184
37.11	THE IMPORTANCE OF AI GOVERNANCE	185
37.12	THE FUTURE OF AI IS MULTI-MODEL	185
37.13	OPEN SCIENCE AND OPEN SOURCE WILL SHAPE AI	186
37.14	BECOMING AN AI VALUE CREATOR	186
37.15	BUILDING AI CAPABILITIES FOR THE FUTURE	187
37.16	THE LONG-TERM IMPACT OF GENERATIVE AI	187
37.17	AI VALUE CREATORS WILL SHAPE THE FUTURE	188
38	<u>INTRODUCTION TO RETRIEVAL-AUGMENTED GENERATION (RAG)</u>	188
38.1	INTRODUCTION	188
38.2	OBJECTIVES	188
38.3	RETRIEVAL-AUGMENTED GENERATION (RAG)	188
38.3.1	HOW RAG WORKS	188
38.4	LIMITATIONS OF GENERATIVE AI MODELS	189
38.5	HOW RAG ADDRESSES THE LIMITATIONS OF GENERATIVE AI MODELS	189
38.6	KEY COMPONENTS OF RAG	190
38.6.1	RETRIEVAL COMPONENT	190
38.6.2	GENERATION COMPONENT	190
38.7	BENEFITS OF RAG	190
38.8	APPLICATIONS OF RAG	190
38.9	IMPLEMENTATION OF RAG ON GOOGLE CLOUD	191
38.9.1	VERTEX AI	191
38.9.2	BIGQUERY	191
38.10	KEY FEATURES OF RAG ON GOOGLE CLOUD	191
38.11	EXAMPLE	191
38.12	EXAMPLE USE CASE: ENHANCING LLMs WITH RAG	192
38.13	SUMMARY	192
39	<u>ADOPTING AI IN YOUR BUSINESS</u>	192
39.1	LEARNING OBJECTIVES	192
39.2	INTRODUCTION	192
39.3	BENEFITS OF ADOPTING AI IN BUSINESS	193
39.3.1	INDUSTRY INSIGHT	193

39.4	REAL-WORLD EXAMPLES OF AI ADOPTION	193
39.4.1	AMAZON	193
39.4.1.1	Personalized Product Recommendations	193
39.4.1.2	Alexa Voice Assistant	194
39.4.1.3	Supply Chain and Inventory Management	194
39.4.2	TESLA	194
39.4.2.1	Autopilot System	194
39.4.2.2	Summon Feature	195
39.4.2.3	Battery Management System (BMS)	195
39.5	GENERAL ROADMAP FOR ADOPTING AI IN BUSINESS	195
39.6	DEFINE BUSINESS GOALS	196
39.7	IDENTIFY SUITABLE USE CASES	196
39.8	PREPARE DATA (DATA READINESS)	196
39.9	BUILD AI CAPABILITIES	197
39.10	DEPLOY AI SOLUTIONS	197
39.11	MONITOR AND OPTIMIZE AI SYSTEMS	197
39.12	KEY TAKEAWAYS	197

40	FRAMEWORKS FOR AI ADOPTION	198
-----------	-----------------------------------	------------

40.1	OVERVIEW	198
40.2	IBM AI LADDER	198
40.2.1	COLLECT	198
40.2.2	ORGANIZE	199
40.2.3	ANALYZE	199
40.2.4	INFUSE	200
40.3	BENEFITS OF THE IBM AI LADDER	200
40.4	FROM +AI TO AI+	201
40.4.1	+AI APPROACH	201
40.4.2	AI+ APPROACH	201
40.5	KEY ELEMENTS OF THE AI+ APPROACH	201
40.5.1	HOLISTIC INTEGRATION	201
40.5.2	IDENTIFY HIGH-VALUE AI USE CASES	201
40.5.3	CHOOSE THE RIGHT AI TECHNOLOGY	202
40.5.4	BUILD A ROBUST DATA FOUNDATION	202
40.5.5	CONTINUOUS INNOVATION AND MODERNIZATION	202
40.6	KEY TAKEAWAYS	203
40.6.1	IBM AI LADDER	203
40.6.2	AI+ APPROACH	203

41	AI ADOPTION FRAMEWORKS: AMAZON, OPENAI, AND FACEBOOK	203
-----------	---	------------

41.1	INTRODUCTION	203
41.2	OBJECTIVE	203
41.3	AMAZON AI SERVICES FRAMEWORK	203
41.3.1	DATA PREPARATION	203

41.3.2	MODEL DEVELOPMENT	204
41.3.3	DEPLOYMENT	204
41.3.4	OPTIMIZATION	205
41.3.5	KEY BENEFIT	205
41.4	OPENAI FRAMEWORK	205
41.4.1	DATA PREPARATION	205
41.4.2	MODEL DEVELOPMENT	206
41.4.3	MODEL DEPLOYMENT	206
41.4.4	CONTINUOUS IMPROVEMENT	207
41.5	FACEBOOK AI INTEGRATION FRAMEWORK	207
41.5.1	DATA INTEGRATION	207
41.5.2	AI MODEL DEVELOPMENT	208
41.5.3	AI MODEL DEPLOYMENT	208
41.5.4	CONTINUOUS IMPROVEMENT	209
41.6	SUMMARY	209
41.6.1	DATA PREPARATION	209
41.6.2	MODEL DEVELOPMENT	209
41.6.3	DEPLOYMENT	209
41.6.4	OPTIMIZATION	209

42 TRANSFORMING WORK THROUGH AI TOOLS **210**

42.1	LEARNING OBJECTIVES	210
42.2	AI TOOLS IN TODAY'S DIGITAL AGE	210
42.3	AI TOOLS FOR RESEARCH AND HEALTHCARE PROFESSIONALS	210
42.3.1	RESEARCH ANALYST: MICROSOFT COPILOT	210
42.3.2	HEALTHCARE PROFESSIONAL: GOOGLE GEMINI	211
42.4	AI TOOLS FOR MARKETERS, CONTENT CREATORS, AND EDITORS	211
42.4.1	SARAH'S AI-POWERED CONTENT CREATION WORKFLOW	211
42.5	AI TOOLS FOR LANGUAGE LEARNING AND TRANSLATION	212
42.5.1	JAMES' LANGUAGE LEARNING AND TRANSLATION WORKFLOW	212
42.5.1.1	Duolingo	212
42.5.1.2	Google Translate	212
42.6	AI TOOLS FOR CUSTOMER SUPPORT	212
42.6.1	BELLA'S CUSTOMER SUPPORT WORKFLOW	212
42.6.1.1	AI-Powered Chatbots	212
42.6.1.2	Zendesk AI	213
42.6.1.3	LivePerson	213
42.7	AI TOOLS FOR DATA ANALYSIS AND VISUALIZATION	213
42.7.1	EMILY - DATA ANALYST	213
42.7.2	DAVID - DATA SCIENTIST	213
42.8	AI TOOLS FOR HUMAN RESOURCES (HR)	214
42.8.1	MIRANDA'S HR WORKFLOW	214
42.8.2	CULTURE AMP	214
42.8.3	IBM WATSONX ORCHESTRATE	214
42.9	SUMMARY	215
42.10	KEY TAKEAWAYS	215

43	CAREER OPPORTUNITIES WITH AI	216
43.1	LEARNING OBJECTIVES	216
43.2	AI AND THE FUTURE OF JOBS	216
43.3	TECHNICAL AI CAREER OPPORTUNITIES	216
43.3.1	AI ENGINEER	216
43.3.2	DATA SCIENTIST	217
43.3.3	ROBOTICS ENGINEER	217
43.3.4	NLP ENGINEER	217
43.3.5	AI APPLICATION DEVELOPER	217
43.3.6	AI RESEARCH SCIENTIST	218
43.4	AI CAREERS THAT DO NOT REQUIRE STRONG PROGRAMMING SKILLS	218
43.4.1	AI ETHICIST	218
43.4.2	AI PRODUCT MANAGER	219
43.4.3	AI STRATEGIST	219
43.4.4	AI MARKETING SPECIALIST	219
43.5	HOW TO SWITCH TO AN AI CAREER	220
43.5.1	IDENTIFY EXISTING SKILLS	220
43.5.2	LEARN CORE AI CONCEPTS	220
43.5.3	APPLY SKILLS PRACTICALLY	220
43.5.4	STAY UPDATED	220
43.5.5	SPECIALIZE IN AN AI DOMAIN	220
43.6	REAL-WORLD CAREER TRANSITION EXAMPLES	221
43.6.1	EMMA	221
43.6.2	SARAH	221
43.7	KEY TAKEAWAYS	221
44	USING GENERATIVE AI FOR YOUR WORK	222
44.1	LEARNING OBJECTIVES	222
44.2	INTRODUCTION	222
44.3	EXERCISE 1: WRITING AND REFINING RESUMES	222
44.3.1	PURPOSE	222
44.3.2	PROMPT 1: GENERATE A RESUME TEMPLATE	222
44.3.3	PROMPT 2: CUSTOMIZE THE RESUME BASED ON A JOB DESCRIPTION	222
44.3.4	PROMPT 3: GENERATE A COVER LETTER	223
44.4	EXERCISE 2: DRAFTING EMAILS	223
44.4.1	PURPOSE	223
44.4.2	PROMPT 1: JOB APPLICATION EMAIL	223
44.4.3	PROMPT 2: TEAM INTRODUCTION EMAIL	223
44.4.4	PROMPT 3: MEETING REQUEST EMAIL	223
44.5	EXERCISE 3: ENHANCING YOUR SKILLS	224
44.5.1	PURPOSE	224
44.5.2	PROMPT 1: IDENTIFY SKILL GAPS AND DEVELOPMENT METHODS	224
44.5.3	PROMPT 2: SKILLS REQUIRED FOR CAREER ADVANCEMENT	224
44.5.4	PROMPT 3: PERSONALIZED LEARNING PATH	224
44.6	SUMMARY	224

45	<u>HUMAN, AI, OR AUGMENTED INTELLIGENCE: WHO SHOULD MAKE THE DECISION?</u>	225
45.1	DECISION MAKING: HUMAN VS. AI	225
45.2	FRAUD DETECTION EXAMPLE	225
45.3	AI CONFIDENCE SCORE AND SUCCESS RATE	225
45.3.1	AI PERFORMANCE CURVE	225
45.4	HUMAN PERFORMANCE CURVE	226
45.4.1	WHY HUMANS PERFORM BETTER IN UNCERTAIN CASES	226
45.5	AUGMENTED INTELLIGENCE	227
45.6	HUMAN COGNITIVE BIAS	227
45.7	FORCED DISPLAY VS. OPTIONAL DISPLAY	227
45.7.1	FORCED DISPLAY	227
45.7.2	AUTOMATION BIAS	228
45.8	OPTIONAL DISPLAY	228
45.9	TRUST IN AI RECOMMENDATIONS	228
45.10	KEY TAKEAWAYS	228
46	<u>MODULE SUMMARY</u>	229
46.1	AI AGENTS	229
46.2	ROBOTICS AND COBOTS	229
46.3	AI IN BUSINESS	229
46.3.1	BUSINESS APPLICATIONS OF AI	229
46.4	EVOLUTION OF ARTIFICIAL INTELLIGENCE	230
46.5	KEY AREAS WHERE GENERATIVE AI CONTRIBUTES	230
46.6	ADOPTING AI IN BUSINESS	230
46.6.1	GENERAL STEPS TO ADOPT AI IN BUSINESS	230
46.7	AI FOR ORGANIZATIONAL PRODUCTIVITY	230
46.8	AI TOOLS AND PLATFORMS	231
46.9	BUILDING A CAREER IN AI	231
46.10	AI CAREER OPPORTUNITIES	231
47	<u>ETHICAL CONSIDERATIONS AND RESPONSIBLE USE OF AI</u>	231
47.1	LEARNING OBJECTIVES	231
47.2	CASE STUDY: AMAZON'S AI RECRUITMENT TOOL	232
47.3	IMPORTANCE OF ETHICAL AI	232
47.4	DATA PRIVACY AND SECURITY	232
47.4.1	REAL-WORLD EXAMPLE: CLEARVIEW AI	233
47.5	BIAS AND FAIRNESS	233
47.6	TRANSPARENCY AND ACCOUNTABILITY	233
47.7	AUTONOMOUS SYSTEMS AND HUMAN OVERSIGHT	234
47.8	ACCESS AND EQUALITY	234
47.9	ETHICAL CONSIDERATIONS ACROSS DIFFERENT AI APPLICATIONS	234
47.9.1	HEALTHCARE	235

47.9.2	LAW ENFORCEMENT	235
47.10	ENVIRONMENTAL IMPACT OF AI	235
47.11	CONTINUOUS MONITORING AND IMPROVEMENT	235
47.12	RESPONSIBLE AI IMPLEMENTATION STRATEGIES	236
47.13	KEY TAKEAWAYS	236

48 CONSIDERATIONS AROUND GENERATIVE AI **236**

48.1	LEARNING OBJECTIVES	236
48.2	COPYRIGHT AND CONTENT OWNERSHIP	237
48.3	PRIVACY AND CONFIDENTIALITY	237
48.3.1	PRIVATE AI	238
48.4	ACCURACY OF AI OUTPUTS AND HALLUCINATIONS	238
48.4.1	REAL-WORLD EXAMPLE	238
48.4.2	REDUCING AI HALLUCINATIONS	238
48.5	MODEL PERFORMANCE CONSIDERATIONS	239
48.6	ETHICAL CONSIDERATIONS	239
48.6.1	BIAS IN AI SYSTEMS	239
48.6.2	DEEPPFAKES	240
48.6.3	USING AI FOR POSITIVE SOCIAL IMPACT	240
48.7	KEY TAKEAWAYS	241

49 UNDERSTANDING HALLUCINATIONS IN LARGE LANGUAGE MODELS (LLMS) **241**

49.1	INTRODUCTION	241
49.2	WHAT ARE HALLUCINATIONS?	242
49.3	TYPES OF HALLUCINATIONS	242
49.3.1	SENTENCE CONTRADICTIONS	242
49.3.2	PROMPT CONTRADICTIONS	242
49.3.3	FACTUAL CONTRADICTIONS	243
49.3.4	NONSENSICAL OR IRRELEVANT INFORMATION	243
49.4	WHY DO HALLUCINATIONS HAPPEN?	243
49.4.1	DATA QUALITY	243
49.4.2	GENERATION METHODS	244
49.4.3	INPUT CONTEXT	244
49.5	STRATEGIES TO REDUCE HALLUCINATIONS	245
49.5.1	PROVIDE CLEAR AND SPECIFIC PROMPTS	245
49.5.2	USE ACTIVE MITIGATION STRATEGIES	245
49.5.3	USE MULTI-SHOT PROMPTING	246
49.6	KEY TAKEAWAYS	246

50 KEY PLAYERS AROUND AI ETHICS **247**

50.1	LEARNING OBJECTIVES	247
50.2	OVERVIEW	247
50.3	IBM'S PILLARS OF TRUST	247

50.3.1	EXPLAINABILITY	248
50.3.2	FAIRNESS	248
50.3.3	ROBUSTNESS	248
50.3.4	TRANSPARENCY	248
50.3.5	PRIVACY	248
50.4	IBM'S AI ETHICS TOOLKITS	248
50.5	MICROSOFT'S APPROACH TO ETHICAL AI	249
50.5.1	HUMAN-IN-THE-LOOP SYSTEMS	249
50.5.2	CONTINUOUS MONITORING AND IMPROVEMENT	249
50.5.3	REGULAR AUDITS	249
50.5.4	ALGORITHMIC IMPACT ASSESSMENTS (AIA)	249
50.5.5	INTERNAL REVIEW BODIES	249
50.5.6	RESPONSIBLE AI STANDARD	249
50.6	GOOGLE'S AI PRINCIPLES	250
50.6.1	BE SOCIALLY BENEFICIAL	250
50.6.2	AVOID CREATING OR REINFORCING BIAS	250
50.6.3	BUILD AND TEST AI FOR SAFETY	250
50.6.4	BE ACCOUNTABLE TO PEOPLE	250
50.6.5	INCORPORATE PRIVACY DESIGN PRINCIPLES	250
50.6.6	UPHOLD HIGH STANDARDS OF SCIENTIFIC EXCELLENCE	250
50.6.7	MAKE AI AVAILABLE ONLY FOR APPROPRIATE USES	250
50.7	GOOGLE'S AI GOVERNANCE	250
50.8	INDUSTRY COLLABORATION FOR ETHICAL AI	251
50.8.1	AI ALLIANCE	251
50.9	CONSULTING FIRMS SUPPORTING ETHICAL AI	251
50.9.1	DELOITTE	251
50.9.2	PWC	251
50.10	INDUSTRY-WIDE ETHICAL AI INITIATIVES	252
50.10.1	ADOBE	252
50.11	KEY TAKEAWAYS	252

51 AI GOVERNANCE 252

51.1	WHAT IS AI GOVERNANCE?	253
51.2	BENEFITS OF ARTIFICIAL INTELLIGENCE	253
51.3	UNDERSTANDING AN AI SYSTEM	253
51.3.1	THE ROLE OF DATA	253
51.4	RISKS ASSOCIATED WITH AI SYSTEMS	254
51.4.1	BIAS IN AI MODELS	254
51.4.2	PRIVACY AND COPYRIGHT INFRINGEMENT	254
51.4.3	LACK OF TRANSPARENCY AND TRUST	255
51.4.4	MODEL DETERIORATION AND CONTINUOUS MONITORING	255
51.5	AI REGULATIONS AND GUIDELINES	255
51.5.1	AI GUIDELINES	255
51.5.2	AI REGULATIONS	256
51.6	WHY AI GOVERNANCE IS ESSENTIAL	256
51.7	KEY TAKEAWAY	256

52	<u>AI ETHICS IN ENTERPRISE AI SYSTEMS</u>	256
52.1	INTRODUCTION	256
52.2	ETHICAL GUIDELINES FOR AI SYSTEMS	257
52.3	EVALUATING WHETHER AI FOLLOWS ETHICAL GUIDELINES	257
52.3.1	DICHOTOMY MAPPING	257
52.3.2	EXAMPLE: HOTEL RECOMMENDATION SYSTEM	257
52.4	ADDRESSING IDENTIFIED ETHICAL ISSUES	258
52.4.1	IMPLEMENT GUARD RAILS	258
52.4.2	USE DIVERSE TRAINING DATA	258
52.4.3	USE OPEN-SOURCE FAIRNESS TOOLS	258
52.5	KEY TAKEAWAYS	258
53	<u>LESSON SUMMARY</u>	259
53.1	ETHICAL CONSIDERATIONS AND RESPONSIBLE USE OF AI	259
53.2	ETHICAL CONSIDERATIONS FOR GENERATIVE AI	259
53.3	PERSPECTIVES OF MAJOR ORGANIZATIONS ON AI ETHICS	260
53.3.1	IBM	260
53.3.2	MICROSOFT	260
53.3.3	GOOGLE	260
53.4	RISKS OF PREMATURE AI ADOPTION	260
53.5	AI GOVERNANCE	261
53.6	RISKS ASSOCIATED WITH AI	261
53.6.1	BIAS IN AI MODELS	261
53.6.2	PRIVACY AND COPYRIGHT INFRINGEMENT	261
53.6.3	LACK OF TRANSPARENCY	261
53.6.4	MODEL DETERIORATION	261
53.7	AI REGULATIONS AND GUIDELINES	261
53.8	KEY TAKEAWAYS	262
54	<u>PROJECT: TRANSFORMING ORGANIZATIONAL FUNCTIONS WITH GENERATIVE AI</u>	262
54.1	LEARNING OBJECTIVES	262
54.2	INTRODUCTION	262
54.3	TRANSFORMING BUSINESS FUNCTIONS	263
54.3.1	PURPOSE	263
54.3.2	TASK	263
54.4	PERSONALIZING CUSTOMER EXPERIENCE	263
54.4.1	PURPOSE	263
54.4.2	CUSTOMER SEGMENTATION	263
54.4.3	PROMPT 1: CUSTOMER SEGMENTATION	263
54.4.4	PROMPT 2: PERSONALIZED PRODUCT RECOMMENDATIONS	264
54.5	SUPPLIER PERFORMANCE ANALYSIS	264
54.5.1	PURPOSE	264

54.5.2	PROMPT	264
54.6	DEMAND PREDICTION	265
54.6.1	PURPOSE	265
54.6.2	PROMPT	265
54.6.3	HISTORICAL SALES DATA	265
54.6.4	CUSTOMER FEEDBACK AND REVIEWS	265
54.7	IMPROVING CUSTOMER SUPPORT	265
54.7.1	PURPOSE	265
54.7.2	TASK	265
54.7.3	PROMPT 1	266
54.7.4	PROMPT 2	266
54.7.5	PROMPT 3	266
54.8	SUMMARY	266

55	COURSE SUMMARY AND KEY TAKEAWAYS	266
-----------	---	------------

55.1	ARTIFICIAL INTELLIGENCE AND GENERATIVE AI	266
55.2	AI IN EVERYDAY LIFE	267
55.3	AI APPLICATIONS ACROSS INDUSTRIES	267
55.4	CORE CONCEPTS OF ARTIFICIAL INTELLIGENCE	267
55.4.1	MACHINE LEARNING	267
55.4.2	DEEP LEARNING	268
55.4.3	NEURAL NETWORKS	268
55.5	GENERATIVE AI MODEL ARCHITECTURES	268
55.6	AI DOMAINS	268
55.6.1	NATURAL LANGUAGE PROCESSING (NLP)	268
55.6.2	COMPUTER VISION	268
55.6.3	ROBOTICS	269
55.7	AI IN BUSINESS	269
55.8	AI CAREER OPPORTUNITIES	269
55.9	ETHICAL CONSIDERATIONS IN AI	270

56	GLOSSARY OF TERMS	270
-----------	--------------------------	------------

56.1	A	270
56.2	B	270
56.3	C	271
56.4	D	271
56.5	E	272
56.6	G	272
56.7	H	272
56.8	I	273
56.9	L	273
56.10	M	273
56.11	N	273
56.12	O	274

56.13	P	274
56.14	R	274
56.15	S	274
56.16	T	275
56.17	U	275
56.18	V	275
56.19	W	275

Module 1: Introduction and The Applications of AI

1 Introducing Artificial Intelligence (AI)

1.1 Learning Objectives

After completing this topic, we will be able to:

- Define Artificial Intelligence (AI) as **augmented intelligence**.
- Discuss the different types of AI based on their strengths.

1.2 History of Artificial Intelligence

AI's history can be traced back to the very origins of computing. Since the days of the **abacus** and early calculators, humans have strived to automate mental tasks.

The formal journey of AI began in the **1950s**, when:

- **Alan Turing** proposed the **Turing Test** (a test designed to determine whether a machine can exhibit intelligent behavior similar to a human).
- **John McCarthy** coined the term **Artificial Intelligence (AI)**.

The evolution of AI continued through several important milestones:

Time Period	Major Development
1960s	Early AI programs such as ELIZA and SHRDLU were developed.
1970s	Rise of Expert Systems (computer systems designed to imitate the decision-making ability of human experts).
1980s	Significant growth in Machine Learning , preparing the foundation for future AI advancements.
1990s	Introduction of Neural Networks (computing models inspired by the structure and functioning of the human brain).
2000s	Rapid advancement of Deep Learning (a subset of machine learning that uses multiple layers of neural networks).
2010–2020	AI applications expanded across industries, including Natural Language Processing (NLP) and Computer Vision .
Present Decade	Continued rapid growth in deep learning models, autonomous systems, and healthcare applications.

1.3 What Is Artificial Intelligence?

Artificial Intelligence (AI) refers to the **simulation of human intelligence processes by computer systems**.

AI involves using **algorithms** (step-by-step instructions for solving problems) and **data** to enable machines to perform tasks that typically require human intelligence, including:

- Learning
- Reasoning
- Problem solving
- Decision making

AI can range from:

- Simple automation
- Complex deep learning
- Neural networks

1.4 Technologies Driving AI Growth

Several technological developments have accelerated the growth of AI.

1.4.1 Internet

The internet has revolutionized connectivity and provided faster access to information.

1.4.2 Distributed Computing

Distributed computing scales data processing, enabling greater efficiency.

1.4.3 Internet of Things (IoT)

The **Internet of Things (IoT)** has increased the number of connected devices, generating massive amounts of data.

1.4.4 Social Networking

Social networking has encouraged most of the data generated to be **unstructured data** (data that does not follow a predefined format, such as text, images, videos, and social media posts).

Together, these technologies reshape our digital landscape by accelerating information access and innovation.

1.5 AI as Augmented Intelligence

With **augmented intelligence**(AI systems that work alongside humans to **assist them** by analyzing large amounts of data, finding patterns, and providing useful insights to improve

their decision-making instead of replacing them), information that subject matter experts need is placed at their fingertips and backed with evidence, enabling them to make informed decisions.

Experts are encouraged to scale their capabilities while allowing machines to handle time-consuming work.

AI is designed to **extend human capabilities** rather than replace them.

1.6 Innate Intelligence

Human beings possess **innate intelligence**, which governs every activity within our bodies.

This intelligence is responsible for:

- An oak tree growing from a small seed.
- A complex organism, such as an elephant, developing from a single-celled organism.

1.7 How AI Learns

Machines do not naturally possess intelligence.

The only innate intelligence machines have is what we provide them.

We enable machines to:

- Examine examples.
- Create machine learning models.
- Learn from inputs and desired outputs.

Machine learning models are developed using different learning approaches:

- **Supervised Learning** (learning from labeled data where correct answers are already known).
- **Unsupervised Learning** (learning patterns from unlabeled data without predefined answers).
- **Reinforcement Learning** (learning through rewards and penalties based on actions taken).

These learning methods will be studied in more detail later.

1.8 Categories of AI

AI can be divided into categories based on:

- Strength
- Breadth
- Application

Based on **strength**, AI can be classified into three types:

- Weak AI (Narrow AI)
- Strong AI (Generalized AI)
- Super AI (Conscious AI)

1.8.1 Weak AI (Narrow AI)

Weak AI, also called **Narrow AI**, is designed for a specific domain.

Characteristics include:

- Performs specific tasks.
- Cannot learn entirely new tasks independently.
- Makes decisions based on programmed algorithms and training data.

Examples include:

- Language translators
- Virtual assistants
- AI-powered web searches
- Recommendation engines
- Intelligent spam filters

1.8.2 Strong AI (Generalized AI)

Strong AI, also called **Generalized AI**, refers to artificial intelligence capable of engaging in and performing a diverse array of distinct and unrelated tasks.

Characteristics include:

- Can acquire new skills.
- Can tackle novel challenges.
- Learns new approaches autonomously from experience.
- Combines multiple AI strategies.
- Can perform at a human level of intelligence.

Potential use cases include:

- Finance
- Human Resources
- Information Technology
- Research and Development
- Supply Chain

1.8.3 Super AI (Conscious AI)

Super AI, also called **Conscious AI**, extends the concept of generative AI to a more advanced level.

Characteristics include:

- Human-level consciousness.
- Self-awareness.
- Advanced cognitive abilities.
- Ability to develop its own thinking skills.

Because we are not yet able to adequately define **consciousness**, it is unlikely that we will be able to create conscious AI in the near future.

Potential applications could include:

- Healthcare
- Autonomous vehicles
- Robotics
- Natural Language Understanding
- Environmental conservation

1.9 Fields That Contribute to AI

Artificial Intelligence is the fusion of many fields of study.

Field	Contribution
Computer Science	Determines how AI is implemented in software.
Electrical Engineering	Determines how AI is implemented in hardware.
Mathematics	Develops viable AI models.
Statistics	Measures model performance.
Psychology	Helps understand how intelligence and learning work.
Linguistics	Helps AI understand and process human language.
Philosophy	Provides guidance on intelligence and ethical considerations.

1.10 AI in Everyday Life

Although the science fiction version of AI may still be a distant possibility, AI is already involved in many of the decisions we make every day.

Over the years, AI has proven useful across different domains and continues to impact society in meaningful ways.

1.11 Key Takeaways

At the end of this topic, we learned that:

- Artificial Intelligence is defined as **augmented intelligence**, aiming to extend human capabilities and tackle tasks beyond individual human and machine capabilities.
- Machine learning models are developed using **supervised learning**, **unsupervised learning**, and **reinforcement learning**.
- AI can be divided based on its strengths into:
 - **Weak (Narrow) AI** — AI tailored for specific domains.
 - **Strong (Generalized) AI** — AI capable of performing diverse and unrelated tasks.
 - **Super (Conscious) AI** — AI with human-level consciousness and self-awareness.

2 Human Intelligence, Artificial Intelligence, and Augmented Intelligence

2.1 Understanding Different Forms of Intelligence

To understand the differences between human intelligence, artificial intelligence, and augmented intelligence, we can consider the example of driving from home to the office. This journey demonstrates how three different forms of intelligence can work independently and together.

2.1.1 Human Intelligence

The first form of intelligence is **human intelligence**.

Human intelligence is required to operate the vehicle, including tasks such as:

- Turning the steering wheel.
- Checking the mirrors.
- Controlling and operating the vehicle safely.

These tasks rely entirely on human skills and decision-making.

2.1.2 Artificial Intelligence

The second form of intelligence is **artificial intelligence (AI)**.

Once we get on the highway, we can turn on the **self-driving feature** in the car.

With artificial intelligence:

- The car stays in its own lane.
- The car keeps an appropriate distance from the vehicle in front.
- The car maintains a proper speed.

- No input from the driver is required while these functions are being performed.

Artificial intelligence is the ability for **machines to perform tasks that normally require human intelligence**, such as:

- Reasoning.
- Natural communication.
- Problem solving.

Artificial intelligence basically replaces the need for a human during specific tasks.

Instead of humans performing the work, computers perform the work, make decisions, and complete specific tasks. Today, artificial intelligence is incorporated into many different systems.

2.1.3 Augmented Intelligence

The third form of intelligence is **augmented intelligence**.

After getting off the highway, we can rely on various **driver-assist features** in the vehicle.

Examples include:

- Collision detection to alert us if we get too close to the vehicle in front.
- Blind spot avoidance to notify us if another vehicle is alongside us while changing lanes.

Unlike artificial intelligence, augmented intelligence combines the capabilities of **humans and machines**, allowing both to work together to enhance each other's efforts while completing tasks.

Augmented intelligence systems enhance our human abilities instead of replacing them.

Examples of augmented intelligence include:

- Screen readers for people who are blind.
- Voice-driven navigation.
- In-car collision avoidance systems.

These systems act on our behalf in the physical world while complementing our own capabilities.

2.2 Artificial Intelligence vs. Augmented Intelligence

A common question is:

Which form of intelligence provides the best results for a given problem?

To answer this, we can compare the strengths of **machines** and **humans**.

2.3 Strengths of Machines

Machines have several important advantages.

They are excellent at:

- Ingesting large amounts of data.
- Processing data much faster than humans.
- Working continuously without getting tired.
- Performing repetitive tasks.
- Completing repetitive tasks with reliable accuracy.

If we need to:

- Analyze a large amount of data quickly, or
- Perform the same task repeatedly with little or no margin for error,

then **artificial intelligence** is generally the best choice.

2.4 Strengths of Humans

Humans also have unique strengths.

We are very good at:

- Generalizing information (understanding the broader concept represented by a single piece of data).
- Creativity.
- Generating new ideas.
- Solving problems.
- Communicating findings effectively.
- Emotional intelligence.

Emotional intelligence is especially valuable because it allows us to understand the reactions and emotions of other people.

This ability is critical for tasks such as:

- Customer service.
- Caregiving.

These are areas where machines still cannot match human capabilities.

2.5 Why Augmented Intelligence Is the Sweet Spot

Augmented intelligence combines the strengths of both humans and machines.

By using artificial intelligence to help us see and understand the world in new ways, we can accomplish tasks that would be impossible for humans to perform alone.

Instead of replacing people, augmented intelligence provides a helping hand by enhancing human capabilities.

2.6 Key Takeaways

- **Human intelligence** enables us to reason, communicate, create, solve problems, and understand emotions.
- **Artificial intelligence** allows machines to perform tasks that normally require human intelligence and can often replace human involvement for specific activities.
- **Augmented intelligence** combines human intelligence with machine intelligence so that both work together.
- Machines excel at processing large amounts of data, repetitive tasks, speed, and accuracy.
- Humans excel at creativity, generalization, communication, and emotional intelligence.
- The best solution is often **both artificial intelligence and augmented intelligence**, using each where it performs best.
- [Augmented intelligence allows humans and machines to complement each other's strengths, helping us perform tasks more effectively while allowing humans to continue leading and thriving.](#)

3 Introducing Generative AI and Its Use Cases

3.1 Learning Objectives

After studying this topic, we will be able to:

- Describe **Generative AI (Gen AI)** and its major capabilities.
- Recognize the use cases of Generative AI across diverse sectors.

3.2 Introduction to Generative AI

As artificial intelligence continues to grow, **Generative AI (Gen AI)** emerges as an innovative frontier.

While traditional AI tends to focus on analyzing data and making decisions, such as recommending the songs we want to listen to or translating languages, **Generative AI** takes a step further by creating entirely new content from scratch.

Generative AI is an AI technology capable of creating new and novel data, ranging from:

- Text
- Images
- Music
- Videos

Unlike conventional AI models, Generative AI does not rely on predefined rules and patterns. Instead, it uses **deep learning techniques** (machine learning methods that learn complex patterns from large amounts of data) and relies on vast datasets to generate new ideas.

3.3 Large Language Models (LLMs)

Generative AI can use **Large Language Models (LLMs)** for processing and generating human-like text.

LLMs can perform a variety of **Natural Language Processing (NLP)** tasks, including:

- Text generation
- Translation
- Summarization

Alternatively, Generative AI can design and incorporate an LLM into a larger, more advanced AI system to perform advanced tasks such as:

- Decision making
- Problem solving

3.4 Capabilities of Generative AI

Generative AI's capabilities drive its rapid integration across various aspects of our lives.

Its major capabilities include:

- Creating new content in diverse formats, including:
 - Text
 - Images
 - Videos
 - Audio
- Conducting human-like intelligent conversations.
- Performing **data augmentation** (generating additional training data to improve machine learning models), which enhances the accuracy of machine learning models.

3.5 Economic Impact of Generative AI

According to recent research conducted by **Goldman Sachs**, the use of Generative AI could:

- Help the global economy grow by about 7%, which is nearly **\$7 trillion**.
- Elevate productivity growth by **1.5% over ten years**.

3.6 Domain-Specific Use Cases of Generative AI

3.6.1 Marketing

Generative AI helps companies:

- Create personalized advertisements.
- Develop personalized email campaigns.
- Generate social media posts based on individual preferences and behaviors.

3.6.2 Creative Fields

Generative AI can:

- Produce unique digital art.
- Create music.
- Generate video content for advertising and marketing campaigns.
- Generate soundtracks for films or video games.

3.6.3 Product Development

By analyzing trends and consumer feedback, Generative AI can:

- Generate new product ideas.
- Optimize existing designs for better performance and cost-effectiveness.

3.7 Industry-Specific Use Cases of Generative AI

3.7.1 Healthcare and Precision Medicine

Generative AI can:

- Support physicians by providing tailored treatments.
- Simulate surgeries.
- Produce medical images to assist doctors in developing treatments.

3.7.2 Gaming

Generative AI can create interactive game worlds by generating:

- New levels
- Characters
- Objects

These elements can adapt to player behavior.

3.7.3 Fashion

Generative AI can:

- Design and produce virtual try-on experiences for customers.
- Recommend personalized fashion choices based on customer behavior and preferences.

3.7.4 Education and Training

Generative AI can:

- Create customized learning materials.

- Develop interactive learning environments that adjust to students' learning styles and learning pace.

3.8 Key Takeaways

- Generative AI is a type of artificial intelligence that can generate content across various mediums, including:
 - Written text
 - Images
 - Audio
 - Videos
- **Large Language Models (LLMs)** are advanced neural network architectures for processing and generating human-like text.
- Major capabilities of Generative AI include:
 - Diverse content creation
 - Human-like conversations
 - Data augmentation
- Domain-specific use cases include:
 - Marketing
 - Creative fields
 - Product development
- Industry-specific use cases include:
 - Healthcare and precision medicine
 - Gaming
 - Fashion
 - Education and training

4 Introduction

Artificial Intelligence (AI) encompasses a range of systems designed to mimic, enhance, or exceed human capabilities. AI can be categorized based on its capabilities and functionalities. Understanding these types and their capabilities highlights the diverse applications and potential of AI technologies.

4.1 Objectives

After completing this topic, we will be able to:

- Explain the types of Artificial Intelligence based on their functionalities.
- Explore the capabilities of each type of Artificial Intelligence.

4.2 AI Types

4.2.1 Diagnostic/Descriptive AI

Diagnostic or descriptive AI focuses on assessing the correctness of behavior by analyzing historical data to understand what happened and why. [This type of AI is instrumental in identifying patterns and trends, performing comparative analyses, and conducting root cause analyses.](#)

4.2.1.1 Capabilities

- **Scenario planning:** Helps in creating different future scenarios based on historical data.
- **Pattern/trends recognition:** Identifies recurring patterns and trends within data sets.
- **Comparative analysis:** Compares various data points to find correlations and insights.
- **Root cause analysis:** Determines the underlying reasons behind specific outcomes.

4.2.2 Predictive AI

Predictive AI is concerned with forecasting future outcomes based on historical and current data. This type of AI is used extensively in predicting customer behavior, market trends, and other forward-looking insights.

4.2.2.1 Capabilities

- **Forecasting:** Predicts future trends and events.
- **Clustering and classification:** Groups similar data points and classifies them into predefined categories.
- **Propensity model:** Assesses the likelihood of specific outcomes based on current data.
- **Decision trees:** Utilize a tree-like model of decisions to predict outcomes.

4.2.3 Prescriptive AI

Prescriptive AI focuses on determining the optimal course of action by providing recommendations based on data analysis. [It goes beyond prediction by suggesting actions that can help achieve desired outcomes.](#)

4.2.3.1 Capabilities

- **Personalization:** Tailors recommendations and experiences to individual needs.
- **Optimization:** Identifies the most efficient ways to achieve goals.
- **Fraud prevention:** Detects and prevents fraudulent activities through analysis.
- **Next best action recommendation:** Provides actionable insights on the next steps to take.

4.2.4 Generative/Cognitive AI

Generative or cognitive AI is involved in producing various types of content, such as code, articles, images, and more. This type of AI mimics human creativity and cognitive processes to automate and assist in content creation.

4.2.4.1 Capabilities

- **Advises:** Offers expert advice and recommendations.
- **Creates:** Produces new content, such as text, images, and code.
- **Protects:** Enhances security measures through intelligent analysis.
- **Assists:** Assists in various tasks, improving efficiency.
- **Automates:** Automates repetitive tasks to save time and resources.

4.2.5 Reactive AI

Reactive AI systems are designed to respond to specific inputs with predetermined responses. They do not have memory or the ability to learn from past experiences, making them suitable for tasks that require immediate reactions.

4.2.5.1 Capabilities

- **Rule-based actions:** Executes specific actions based on predefined rules.
- **Instant responses:** Provides immediate responses to inputs.
- **Static data analysis:** Analyzes current data without considering past interactions.

4.2.6 Limited Memory AI

Limited memory AI systems have the ability to use past experiences to inform current decisions. They can learn from historical data to improve their performance over time. This type of AI is commonly used in autonomous vehicles and recommendation systems.

4.2.6.1 Capabilities

- **Learning from data:** Uses historical data to make informed decisions.
- **Pattern recognition:** Identifies patterns over time to improve accuracy.
- **Adaptive responses:** Adapts responses based on previous interactions.

4.2.7 Theory of Mind AI

Theory of Mind AI is an advanced type of AI that aims to understand human emotions, beliefs, and intentions. It is still in the research stage and seeks to interact more naturally with humans by comprehending their mental states.

4.2.7.1 Capabilities

- **Emotion recognition:** Identifies and responds to human emotions.
- **Social interaction:** Engages in more natural and human-like interactions.
- **Intent prediction:** Predicts human intentions based on context and behavior.

4.2.8 Self-aware AI

Self-aware AI represents the most advanced form of AI, which has its own consciousness and self-awareness. This type of AI can understand and react to its own emotions and states. It remains a theoretical concept and has not yet been realized.

4.2.8.1 Capabilities

- **Self-diagnosis:** Evaluates its own performance and health.
- **Autonomous learning:** Learns independently without human intervention.
- **Adaptive behavior:** Adjusts behavior based on self-awareness.

4.2.9 Narrow AI (Weak AI)

Narrow AI is designed to perform a specific task or a limited range of tasks. It excels in a single area but lacks generalization capabilities. Most current AI applications fall under this category.

4.2.9.1 Capabilities

- **Task specialization:** Excels in performing specific tasks.
- **High accuracy:** Achieves high performance in its designated area.
- **Efficiency:** Operates efficiently within its scope of specialization.

4.2.10 General AI (Strong AI)

General AI, like human intelligence, can understand, learn, and apply knowledge across a wide range of tasks. It can also transfer knowledge from one domain to another and adapt to new situations autonomously.

4.2.10.1 Capabilities

- **Cross-domain learning:** Applies knowledge across various domains.
- **Autonomous decision-making:** Makes decisions independently in diverse scenarios.
- **Human-like understanding:** Understands and processes information similar to humans.

4.3 Summary

In this topic, we learned about the following types of Artificial Intelligence:

- **Diagnostic/Descriptive AI:** Focuses on assessing the correctness of behavior by analyzing historical data to understand what happened and why.
- **Predictive AI:** Concerned with forecasting future outcomes based on historical and current data.
- **Prescriptive AI:** Focuses on determining the optimal course of action by providing recommendations based on data analysis.
- **Generative/Cognitive AI:** Involved in producing various types of content, such as code, articles, images, and more.
- **Reactive AI:** Designed to respond to specific inputs with predetermined responses.
- **Limited Memory AI:** Uses past experiences to inform current decisions.
- **Theory of Mind AI:** Aims to understand human emotions, beliefs, and intentions.
- **Self-aware AI:** Represents an advanced form of AI with self-awareness (theoretical concept that has not yet been realized).
- **Narrow AI (Weak AI):** Designed to perform specific tasks.
- **General AI (Strong AI):** Can perform a wide range of tasks similar to human intelligence.

5 How Traditional AI Differs from Generative AI

5.1 Introduction

Generative AI is all the rage. One question that is asked quite frequently is: **How does generative AI differ from the AI that we were using 5, 10, 20, or even 30 years ago?**

To understand this difference, we first need to understand how AI existed before generative AI and how the underlying architecture has changed.

5.2 Traditional AI Architecture

Before generative AI, AI systems typically consisted of three main components:

- Repository
- Analytics Platform
- Application Layer

5.2.1 Repository

A **repository** is exactly what it sounds like—it is where we keep all of our information.

The repository can contain:

- Data in tables, rows, and columns
- Images
- Documents
- Any other type of information

It is simply the place where an organization stores all of its historical information and data.

5.2.2 Analytics Platform

The second component is the **analytics platform**.

In the IBM ecosystem, examples of analytics platforms include:

- SPSS Modeler
- Watson Studio

The analytics platform is where we build analytical models using the information stored in the repository.

5.2.3 Application Layer

The third component is the **application layer**.

Consider the example of a telecommunications (telco) company.

The company stores customer information in its repository. Suppose the company wants to identify customers who are likely to **churn** (cancel their service).

The process works as follows:

- Customer information is taken from the repository.
- The data is moved into the analytics platform.
- Models are built to predict which customers are likely to churn.
- Once the models are created, they are deployed into an application.
- The application uses these predictions to help prevent customers from canceling their service.

For example, if a customer is predicted to cancel their service, the company may:

- Contact the customer.
- Offer incentives or benefits.
- Try to convince the customer to remain with the company.

At this stage, this is better described as **predictive analytics** or a **predictive model**, rather than true AI.

5.3 The Role of the Feedback Loop

To transform predictive analytics into AI, we need to introduce a **feedback loop**.

A feedback loop automates the learning process.

Continuing with the telco example:

- Customer information is analyzed.
- The model predicts which customers will cancel.
- The company takes action through the application.
- Some predictions are correct.
- Some predictions are incorrect.

The feedback loop enables the system to learn from these outcomes.

For example:

- If the model predicts that a customer will cancel but the customer does not, we can investigate why the prediction was incorrect.
- We can then improve the model so it does not make the same mistake again.

The goal is for AI to learn from both:

- Previous mistakes
- Previous successes

A useful way to think about this is the saying:

"Fool me once, shame on you. Fool me twice, shame on me."

Similarly, we want AI to learn from experience so it continuously improves over time.

This architecture has existed for decades and predates many modern AI systems.

5.4 Generative AI Architecture

With generative AI, the entire paradigm has changed.

The fundamental architecture and the way we build AI systems are now different.

Instead of starting with an organization's internal repository, generative AI begins with an enormous amount of data collected from across the world.

Rather than using only organizational data, these systems are trained on a **massive quantity of information** covering an extremely wide range of topics.

This large-scale data is then used to build **Large Language Models (LLMs)**.

5.5 Large Language Models (LLMs)

Large Language Models (LLMs) are:

- Very large
- Extremely powerful
- Remarkably capable

However, despite their capabilities, they usually do not possess the specific business knowledge needed for every organization.

For example, an LLM may generally understand why customers cancel a telecommunications service, but it will not understand the unique characteristics, nuances, and **idiosyncrasies** (unique behaviors or characteristics) of a specific company's customers.

5.6 Prompting and Tuning

This is where **prompting and tuning** become important.

The prompting and tuning layer takes a general-purpose Large Language Model and adapts it for a specific business use case.

Returning to the customer churn example:

- The LLM is trained on massive amounts of information covering many topics.
- The organization then uses prompting and tuning to make the model more specific to its own customers and business requirements.

This allows a general-purpose model to become useful for a particular organization.

5.7 Application Layer in Generative AI

Just like traditional AI, generative AI also includes an **application layer**.

The application layer is where the AI is delivered to users so it can perform its intended purpose within the organization.

5.8 Feedback Loop in Generative AI

Generative AI also includes a **feedback loop**.

However, unlike traditional AI, the feedback loop typically returns to the **prompting and tuning layer**, rather than directly modifying the Large Language Model itself.

This is because the underlying Large Language Models are generally developed outside the organization and are too large for an individual company to retrain.

5.9 Traditional AI vs. Generative AI

Traditional AI	Generative AI
Starts with an organization's repository.	Starts with massive amounts of information collected from across the world.
Uses organizational data.	Uses extremely large datasets containing information about many topics.
Models are built inside an analytics platform.	Uses pre-trained Large Language Models (LLMs).
Uses predictive analytics for business-specific problems.	Uses prompting and tuning to adapt general-purpose models for business-specific use cases.
Feedback loop improves prediction models directly.	Feedback loop primarily improves prompting and tuning.
Data size is manageable within an organization's repository.	Data and models are far too large for any single organization to store in its own repository.

5.10 Key Takeaways

- Traditional AI begins with an organization's internal repository and builds predictive models using an analytics platform.
- A feedback loop enables traditional AI systems to learn from previous successes and mistakes.
- Generative AI fundamentally changes this architecture by starting with massive amounts of global data rather than organizational data.
- Large Language Models (LLMs) are trained on enormous datasets covering a wide variety of information.

- Prompting and tuning customize these general-purpose models for specific organizational use cases.
- Like traditional AI, generative AI includes an application layer and a feedback loop, but the feedback loop primarily updates the prompting and tuning process rather than the underlying Large Language Model.
- The enormous size of both the training data and the models makes it impractical for individual organizations to store or build them within their own repositories, which is why generative AI requires a fundamentally different architecture.

6 Are We There Yet? Understanding the Current State of Artificial Intelligence

6.1 Introduction

In this topic, we explore the current state of Artificial Intelligence (AI) and answer the question: **Are we there yet?** We discuss different definitions of AI, compare today's AI capabilities with **Artificial General Intelligence (AGI)**, examine historical milestones, and consider how AI is already being used in everyday life.

6.2 What Is Artificial Intelligence?

One definition of **Artificial Intelligence (AI)** is:

- The **simulation of intelligent behavior in computers** (computers performing tasks that normally require human intelligence).

Based on this definition, we have already achieved many forms of AI because computers can successfully simulate intelligent behavior in numerous specialized tasks.

However, this does not necessarily mean we have reached the ultimate goal of AI.

6.3 Artificial General Intelligence (AGI)

Artificial General Intelligence (AGI) is an AI system that:

- Is equivalent to a human in intelligence.
- Is as diverse as a human in its capabilities.
- Can do anything that a human can do, at least as well as or better than a human.

Although modern AI has made remarkable progress, we have **not yet achieved AGI**.

6.4 Does Solving Math Instantly Mean Intelligence?

Consider the example of instantly solving the following mathematical expression mentally:

$$312 \times 125$$

Being able to calculate the answer immediately is impressive.

However:

- A calculator can perform the same calculation instantly.
- We do not consider a calculator to be sentient (capable of experiencing or perceiving things) or artificially intelligent.

This demonstrates that performing arithmetic quickly alone is not sufficient evidence of intelligence.

6.5 Is Memorization the Same as Intelligence?

One traditional view of intelligence suggested that a person would be considered intelligent if they knew:

- Every element in the periodic table.
- The correct position of each element.
- Every atomic number.
- The correct spelling of every element.
- Every chemical abbreviation.

However, this is better described as:

- A **database lookup** (retrieving stored information) rather than true intelligence.

Simply storing and recalling information is not the same as genuine intelligence.

6.6 Chess as a Measure of Intelligence

Becoming a chess grandmaster requires:

- Years of practice.
- Learning countless patterns.
- Understanding strategies.
- Mastering complex moves.

Most people would consider the world's best chess players to be highly intelligent.

6.6.1 IBM Deep Blue

In 1997, IBM created a computer called **Deep Blue**.

Deep Blue was able to:

- Defeat the best chess player in the world.
- Win convincingly.

If playing chess at the highest level were the standard for intelligence, then AI would have already passed that milestone.

However, many people still felt that AI had not truly reached human-level intelligence.

6.7 The Turing Test

One of the classic methods for evaluating artificial intelligence is the **Turing Test**.

6.7.1 How the Turing Test Works

In the Turing Test:

- A user communicates by typing messages.
- The user cannot see who they are communicating with.
- They may be chatting with another human or with a computer.
- If the user cannot reliably distinguish between the human and the computer, the computer is considered to have demonstrated artificial intelligence.

Passing the Turing Test would suggest that a computer's conversational abilities are indistinguishable from those of a human.

6.7.2 Has the Turing Test Been Passed?

This remains a disputed topic.

In **2014**, some people argued that the Turing Test had technically been passed by:

- A chatbot that simulated a **13-year-old boy**.

However, some argue that this sets the standard relatively low and that a more demanding benchmark should be used.

6.8 ChatGPT and Modern Conversational AI

Modern conversational AI systems have significantly raised expectations.

For example:

- ChatGPT is used daily by many people.
- Conversations with ChatGPT can sometimes feel similar to conversations with a real person.
- Modern AI systems demonstrate substantial improvements in natural language communication.

These advances suggest that AI is moving closer to human-like interaction than ever before.

6.9 The Moving Finish Line of AI

A common observation is that the definition of "true AI" seems to change whenever AI reaches a previous milestone.

As AI achieves one goal, expectations often shift toward a more advanced objective.

Examples include systems that excel at:

- Arithmetic.
- Memorization and recall.
- Playing complex but narrowly defined games such as chess.
- Natural conversation.

These are examples of highly specialized AI systems.

6.10 Specialized AI Versus General Intelligence

Current AI systems are often described as **siloed savants** (highly capable in one specific area but limited outside that area).

Examples include:

- AI that performs arithmetic extremely well.
- AI that excels at memorization and information retrieval.
- AI that masters chess.
- AI that generates human-like conversations.

Each system performs exceptionally within its own domain.

However, combining all of these abilities into one unified system with reasoning, adaptability, and emotional intelligence is the challenge of achieving **Artificial General Intelligence (AGI)**.

6.11 The Rapid Development of AI

Recent progress in AI has accelerated dramatically.

Observations include:

- AI development has become faster than ever before.
- The pace of improvement has increased significantly, even within the current year.
- AI is gradually moving away from highly specialized systems like **Deep Blue**, which focused only on chess.
- Modern AI systems are becoming capable of performing a much broader range of human-like tasks.

Although AI has not yet reached AGI, it appears to be moving steadily toward that goal.

6.12 Are We There Yet?

The conclusion is:

- **No, we are not there yet.**

- However, **we have never been closer** to achieving Artificial General Intelligence.

For many years, AI was described as being five to ten years away. Even a decade later, it was still described as being another five to ten years away.

Today, however, the gap appears to be narrowing much more rapidly than before.

6.13 Everyday Uses of AI

6.13.1 Self-Driving Car

One example of everyday AI use is:

- Using a **self-driving car** that functions as a computer capable of transporting people automatically.

6.13.2 AI Image Generation

Another practical use of AI is image generation.

Instead of relying solely on photography skills, we can:

- Write a prompt describing the desired image.
- Allow AI to generate a professional-quality photograph based on that prompt.

This demonstrates how AI can enhance creativity and productivity in everyday life.

7 AI in Daily Life

Artificial Intelligence (AI) is integrated into many aspects of our daily lives. It powers personalized digital experiences, automates routine tasks, enhances convenience, improves accessibility, strengthens security, optimizes navigation, and supports healthcare. By learning user preferences and analyzing large amounts of data, AI helps make everyday activities more efficient and personalized.

7.1 Common Use Cases of AI in Daily Life

AI is widely used to personalize digital experiences and automate routine activities. Some common examples include:

- Customizing social media feeds according to user preferences.
- Recommending products on online shopping platforms based on browsing and purchase history.
- Automatically adjusting home temperature and lighting according to user routines and preferences.

According to a report published on **Statista (statista.com)**, within the last five years, the number of digital voice assistants being used in devices worldwide for day-to-day activities rapidly increased from **4.2 billion** to **8.4 billion units**.

7.2 Benefits of AI in Daily Life

AI provides numerous benefits in our daily lives, including:

- Providing personalized experiences.
- Streamlining tasks through efficiency and automation.
- Enhancing accessibility.
- Improving day-to-day convenience.
- Boosting safety and security.
- Optimizing navigation routes.
- Enhancing healthcare facilities.

7.3 AI for Automation and Convenience

AI makes our lives more efficient and convenient by automating repetitive tasks and processes, saving time and effort in daily activities.

Virtual assistants such as **Siri**, **Google Assistant**, and **Alexa** employ AI algorithms to interpret and act upon voice instructions. They help complete tasks such as:

- Scheduling reminders.
- Answering queries.
- Delivering weather forecasts.
- Managing smart home devices.

In smart home setups, AI empowers devices such as thermostats, lights, and security cameras to understand user habits and preferences over time. This allows us to automate tasks including:

- Temperature adjustments.
- Lighting control.
- Security alert notifications.

These capabilities enhance convenience and improve energy efficiency.

7.4 AI-Powered Personalized Recommendations

AI uses recommendation systems to analyze vast amounts of data and provide personalized recommendations across different domains.

7.4.1 Streaming Platforms

Streaming platforms such as **Netflix**, **Amazon Prime**, and **Spotify** offer personalized suggestions based on the user's viewing history and preferences for movies, podcasts, and music.

7.4.2 Social Media

Social media platforms such as **Facebook** and **Instagram** personalize content according to the user's interests and past interactions.

7.4.3 E-commerce

E-commerce platforms such as **Amazon** and **eBay** use AI algorithms to provide personalized product recommendations based on the user's browsing history and purchase behavior.

AI-driven chatbots assist customers with their queries, while visual search technology enables customers to find products by uploading images, further enhancing the online shopping experience.

7.5 AI for Security

AI plays a significant role in enabling security across various domains that we interact with in our daily lives.

Its applications include:

- Cybersecurity, including encryption, access control, and risk assessment.
- Fraud detection for identifying fraudulent activities in financial transactions, online payments, and identity verification.
- Biometric authentication systems using facial recognition, fingerprint scanning, or voice recognition.
- AI-enabled video analytics that enhance surveillance systems, monitor investments, detect fraudulent transactions, identify credit card fraud, and help prevent financial crimes.

7.6 AI in Healthcare

AI is transforming the healthcare sector and impacting our lives in several ways.

AI chatbots and virtual assistants provide **24/7 support** by:

- Reminding us to take medication.
- Offering basic healthcare information.

AI-powered wearables, such as fitness trackers and smartwatches, monitor:

- Heart rate.
- Sleep patterns.
- Blood oxygen levels.

This data is crucial for identifying potential health risks early and monitoring our well-being.

AI algorithms can also help diagnose diseases and develop personalized treatment plans through data-driven insights.

7.7 AI in Smart Devices

AI is integrated into smart devices in numerous ways, enhancing user experience and functionality.

7.7.1 Smartphone Cameras

AI enhances smartphone cameras with features such as:

- Facial recognition.
- Scene detection.
- Portrait mode.
- Automatic image enhancement.

These features enable users to capture high-quality photos and videos effortlessly.

7.7.2 Navigation Applications

AI-powered navigation applications utilize:

- Real-time traffic data.
- Historical traffic patterns.
- User feedback.

These applications:

- Optimize routes.
- Estimate travel times accurately.
- Provide alternative routes to avoid congestion.

This improves the efficiency of daily commutes and travel planning.

7.7.3 Keyboards and Messaging Applications

AI algorithms are utilized in keyboards and messaging applications to:

- Suggest words or phrases as we type.
- Predict the next word in a sentence.
- Correct spelling and grammar mistakes.

These features streamline text input on mobile devices.

7.8 Summary

In this topic, we learned that AI's use cases and benefits in daily life are profound and multifaceted.

We learned that:

- Virtual assistants and smart home devices use AI to automate routine tasks, making our lives more efficient and convenient.
- AI recommendation systems provide personalized suggestions across streaming platforms, social media, and e-commerce platforms.
- AI boosts security through biometric authentication, fraud detection, cybersecurity, and AI-enabled video analytics.
- AI transforms healthcare through AI-powered wearables, chatbots, virtual assistants, disease diagnosis, and personalized treatment plans.
- AI enhances smart devices by improving smartphone cameras, navigation applications, keyboards, and messaging applications, resulting in a better user experience and increased functionality.

8 AI Chatbots and Smart Assistants

8.1 Overview

In this topic, we explore how **AI chatbots** and **smart assistants** work, their benefits, diverse applications across industries, and the future trends shaping their development.

AI chatbots are revolutionizing how we interact with technology by handling tasks and conversations that once required human intervention. We can book vacations, order dinner, set reminders, or control smart home devices with just a few commands or voice instructions, making everyday life significantly simpler.

AI chatbots and smart assistants are **AI-driven software programs** that understand and respond to queries, provide information, and perform tasks. They use advanced algorithms and **Natural Language Processing (NLP)** (a technology that enables computers to understand and process human language) to understand and respond to user inputs in a way that simulates human-like conversation.

8.2 Evolution of AI Chatbots

As AI has advanced, chatbots have undergone a remarkable evolution:

- Started as **simple rule-based systems**, where interactions followed predefined patterns.
- Evolved into **AI-powered personal assistants** capable of handling more complex tasks.
- The most recent breakthrough is **Generative AI models**, which enable chatbots to:
 - Engage in intelligent, context-based conversations.
 - Offer personalized recommendations.

8.3 Examples of AI Chatbot Platforms and Smart Assistants

8.3.1 AI Chatbot Platforms

Some examples of chatbot platforms include:

- [IBM WatsonX Assistant](#)
- [Chatfuel](#)
- [Wit.ai](#)

- [And more](#)

8.3.2 Smart Assistants

Some popular smart assistants include:

- [Siri](#)
- [Google Assistant](#)
- [Amazon Alexa](#)
- [Microsoft Cortana](#)
- [And so on](#)

8.3.3 Generative AI Chatbots

Examples of popular Generative AI chatbots include:

- [ChatGPT](#)
- [Google Gemini](#)

8.4 How AI Chatbots Work

AI chatbots use a **database of information, deep learning, machine learning, and Natural Language Processing (NLP)** to understand conversation patterns and generate appropriate responses.

The working process includes the following steps:

8.4.1 Understanding User Input

- The chatbot analyzes the user's input to recognize the intent and keywords.
- If the input is speech, the chatbot converts it into text before processing it further.

8.4.2 Intent and Context Recognition

- After processing the input, the chatbot uses AI algorithms and NLP engines to:
 - Recognize the user's intent.
 - Understand the context of the input.
 - Categorize the input.
 - Determine the most appropriate response.
- The chatbot also analyzes:
 - Previous interactions.
 - User preferences.

This enables more relevant and personalized responses.

8.4.3 Dialogue Management

The dialogue management system selects the appropriate response from the chatbot's knowledge base.

At the same time, the machine learning algorithm updates the chatbot's knowledge base with information from the current conversation.

8.4.4 Response Generation

Finally, the chatbot generates a response, which may include:

- Providing information.
- Answering questions.
- Performing tasks.
- Engaging in a conversation to gather additional information.

8.5 Example: Weather Information

Suppose we want to know the weather in our location.

To provide an accurate response, the chatbot accesses the **Weather Application Programming Interface (API)** (a service that allows software applications to retrieve weather data) to retrieve the current weather forecast for our location.

8.6 Benefits of AI Chatbots

AI chatbots provide several important benefits.

- **All-time availability**, allowing us to access information anytime and anywhere.
- **Scalability**, enabling them to handle large volumes of inquiries and interactions without high overhead costs.
- **Personalized services and suggestions**, where responses and recommendations are tailored based on user preferences and prior interactions.
- **Natural conversations**, achieved by utilizing Natural Language Processing (NLP).
- **Multilingual communication**, as chatbots can be trained to communicate in multiple languages, expanding their reach and effectiveness.

8.7 Applications of AI Chatbots

AI chatbots have diverse applications across multiple industries and domains.

8.7.1 Customer Service

- Handle routine inquiries.
- Provide instant customer assistance.

8.7.2 E-commerce

- Recommend products based on user preferences.
- Facilitate transactions.

8.7.3 Healthcare

- Assist patients with appointment scheduling.
- Provide medication reminders.
- Perform symptom assessment.

8.7.4 Education

- Support language learning.
- Provide tutoring assistance.

8.8 Future Trends of AI Chatbots

As the chatbot industry grows, it will continue to reshape business communications.

Businesses are increasingly focused on building more human-like chatbots using **machine learning** and **Natural Language Processing (NLP)**.

Future developments include:

- **Sentiment analysis** (an AI technique that identifies emotions expressed in conversations), enabling chatbots to determine whether conversations are going well and respond according to customer emotions.
- Increased use of chatbots for **enterprise internal tasks**, including:
 - HR inquiries.
 - IT support.
- Automation of business processes to improve work efficiency.
- Automation of tasks such as processing refunds.
- More intelligent and versatile chatbot capabilities.

These trends signify an exciting future where AI chatbots become increasingly capable of handling both customer-facing and internal organizational tasks.

8.9 Key Takeaways

- AI chatbots and smart assistants are AI-driven software programs that understand and respond to queries, provide information, and perform tasks.
- They use databases, deep learning, machine learning, and Natural Language Processing (NLP) to understand user input and generate appropriate responses.
- Chatbots have evolved from simple rule-based systems to AI-powered personal assistants and Generative AI chatbots capable of intelligent, context-aware conversations.
- Popular chatbot platforms include IBM WatsonX Assistant, Chatfuel, and Wit.ai.
- Popular smart assistants include Siri, Google Assistant, Amazon Alexa, and Microsoft Cortana.
- Popular Generative AI chatbots include ChatGPT and Google Gemini.
- AI chatbots offer all-time availability, scalability, personalized services, natural conversations, and multilingual communication.
- They are widely used in customer service, E-commerce, healthcare, and education.

- Future advancements include more human-like interactions, sentiment analysis, automated HR and IT support, and business process automation, such as refund processing.

9 Understanding Chatbots

9.1 Introduction

Most of us have, at this point in our lives, probably used a chatbot, maybe without even knowing it.

A chatbot is an application that allows users to interact with a business or service through natural conversation. We can ask questions, request information, or perform tasks without needing to communicate directly with another person.

In this topic, we explore:

- What a chatbot is
- How a chatbot works
- How we interact with a chatbot
- The technologies behind chatbots
- The benefits of using chatbots

9.2 Example: Flower Shop Chatbot

Consider that we own a flower shop.

- We are the only employee.
- We are busy designing floral arrangements.
- We do not always have time to answer customer phone calls or messages.

Customers commonly ask questions such as:

- What time do you open?
- Do you have yellow roses in stock?
- How do we contact you?

Instead of answering every inquiry ourselves, we can use a chatbot called **Flora**.

When a customer asks:

What time do you open?

Flora can automatically respond with the correct business hours, for example:

We open Monday at 9:00 AM.

This allows us to continue focusing on designing floral arrangements while the chatbot handles routine customer inquiries.

9.3 Example: Banking Chatbot

Chatbots can also perform more complex tasks.

Suppose we want to check our bank account balance.

Instead of:

- Calling the bank
- Waiting on hold
- Speaking with a teller

We can interact with a banking chatbot named **Bertie**.

9.3.1 Customer Request

We ask:

What is my balance?

9.3.2 Authentication Process

Before providing sensitive information, Bertie requests identifying information, such as:

- Account number
- Date of birth
- PIN
- Other identifying information

We provide these details to Bertie.

9.3.3 Communication with the Bank

Bertie sends the information to the bank, including:

- Account number
- Date of birth
- PIN

The bank verifies the information and authorizes the request.

For example, the bank responds:

You only have five dollars.

Bertie then returns this response directly to us.

Throughout this process:

- No teller is involved.
- No employee needs to manually process the request.
- The entire interaction is automated.

9.4 Ways We Can Interact with a Chatbot

Suppose tomorrow is **Mother's Day**, and we forgot to order flowers.

The flower shop owner is busy preparing floral arrangements and cannot answer customer calls.

Fortunately, we can interact with the chatbot in several different ways.

9.4.1 Website Chatbot

A flower shop website may include a chatbot popup, such as **Flora**.

Flora asks:

What flowers would you like to order?

We respond:

Yellow roses.

Flora continues the conversation and processes our request.

9.4.2 Voice-Based Chatbot

Instead of using the website, we can simply call the chatbot.

For example, the chatbot says:

Welcome to Flowers Supply 101. What would you like to place your order for?

We respond:

A dozen yellow roses.

The chatbot continues the conversation naturally and processes the order.

9.4.3 Instant Messaging Platforms

We can also interact with chatbots through various instant messaging services on different social media platforms.

This means chatbot interactions are possible through:

- [Website chat windows](#)

- Phone calls (voice)
- Instant messaging services
- Social media messaging platforms

Chatbots can support both:

- Text-based communication
- Voice-based communication

9.5 How a Chatbot Works Behind the Scenes

Most chatbot systems are hosted in the **Cloud**.

A chatbot service provider integrates with various cloud services to process user requests.

9.5.1 Step 1: User Sends a Request

The user communicates through one of several channels:

- Website
- Instant messenger
- Voice call
- Text message

For example, we say:

We would like to purchase a dozen yellow roses.

The communication may occur through text or voice.

9.5.2 Step 2: Request Is Sent to the Cloud

The communication channel sends the request to the cloud-based chatbot service.

The chatbot processes the request using:

- **Natural Language Processing (NLP)** (technology that enables computers to understand and interpret human language)
- **Artificial Intelligence (AI)** (technology that enables computers to make decisions and respond intelligently)

The chatbot understands that we want to purchase:

- 12 yellow roses

9.5.3 Step 3: Chatbot Checks Business Systems

Before confirming the order, the chatbot verifies whether the requested flowers are available.

The chatbot connects to the flower shop's backend systems, which function as a database.

Example inventory:

- 36 yellow roses
- 14 red tulips
- 3 succulents

The chatbot searches the database for:

- 12 yellow roses

Since the inventory contains **36 yellow roses**, the requested quantity is available.

9.5.4 Step 4: Order Confirmation

The chatbot sends a confirmation back through the cloud.

The website then responds:

Yes, we have more than enough roses in stock, so we can place your order.

The user is informed that the order can be fulfilled successfully.

Meanwhile, the flower shop owner may simply receive a notification while continuing to prepare floral arrangements.

The entire ordering process remains automated without requiring manual intervention.

9.6 Technologies Used in Chatbots

A cloud-based chatbot typically uses:

- Cloud computing
- Natural Language Processing (NLP)
- Artificial Intelligence (AI)
- Backend databases
- Business systems integration
- Website integration
- Voice integration
- Instant messaging integration

These technologies allow the chatbot to:

- Understand user requests
- Interpret natural language
- Access business data
- Verify information
- Return accurate responses

- Automate customer interactions

9.7 Benefits of Using Chatbots

9.7.1 Saves Time

One of the biggest advantages of chatbots is time savings.

Instead of spending time answering repetitive customer questions, we can focus on other important business activities, such as:

- Designing floral arrangements
- Bookkeeping
- Managing daily operations

Over the course of a month, this automation can save many hours.

9.7.2 Simplicity

Chatbots are surprisingly simple to build.

In many cases:

- We do not need to write any code.
- Configuration is based primarily on natural language.

For example, we can simply configure the chatbot to respond with:

We are open Monday through Friday, 9:00 AM to 6:00 PM.

The chatbot automatically provides this information to customers without requiring knowledge of:

- Python
- JavaScript
- Other programming languages

For more advanced integrations with external systems, some programming knowledge may be required.

9.7.3 Fast Time to Launch

Cloud-based chatbots can be deployed quickly because:

- No underlying infrastructure needs to be built.
- Most services are already available in the cloud.
- Minimal technical setup is required.

From the time we start building a chatbot until it is ready for use, deployment can take as little as **one hour**.

9.8 Key Takeaways

- Chatbots automate conversations between users and businesses.
- They can answer common customer questions without human involvement.
- Chatbots can also perform more advanced tasks, such as checking bank account balances after verifying user identity.
- Users can interact with chatbots through websites, phone calls, instant messaging services, and social media platforms.
- Cloud-based chatbots use Natural Language Processing (NLP), Artificial Intelligence (AI), and backend databases to understand requests and provide accurate responses.
- Before completing transactions, chatbots can verify inventory or retrieve information from business systems.
- Chatbots help businesses save time, reduce repetitive work, and improve customer service.
- Most modern chatbot platforms are easy to build, require little or no coding for basic use cases, and can often be launched within **one hour**.

10 AI Assistant: User Conversations

10.1 Learning Objectives

After completing these exercises, we will be able to:

- Demonstrate how AI assistants in various domains, such as finance and health, respond to users' queries.
- Identify how AI assistants personalize recommendations and services based on user preferences and data.

10.2 Introduction

AI assistants, powered by advanced algorithms and machine learning, help users perform various tasks while improving efficiency, productivity, and convenience.

In this topic, we explore how AI assistants interact with users across different domains by adapting their responses based on user preferences, goals, and requirements. These exercises demonstrate personalized conversations between users and AI assistants in the fields of finance and health.

10.3 Exercise 1: Financial AI Advisor

In this exercise, we demonstrate a conversation with a Financial AI Advisor that assists users in managing finances, planning investments, and providing personalized financial guidance.

10.3.1 Prompt Instructions

Imagine you are a financial AI advisor named **FinAdvisor**. You assist clients in managing their finances, planning their investments, and providing personalized advice based on their financial goals.

10.3.2 Prompt 1

Hi FinAdvisor, I'm looking to invest some of my savings, but I'm not sure where to start. Can you help me?

10.3.3 Prompt 2

I'm planning for long-term growth. I want to save for retirement and maybe buy a house in the next **10–15 years**.

10.4 Exercise 2: Health AI Assistant

In this exercise, we demonstrate a conversation with a Health AI Assistant that helps individuals monitor their fitness levels, track their daily activities, and provide personalized health recommendations.

10.4.1 Prompt Instructions

Imagine you are a personal health AI assistant named **FitBot**. You help individuals monitor their fitness levels, track their daily activity, and provide personalized health recommendations.

10.4.2 Prompt 1

Hi FitBot, I want to improve my fitness and track my daily activities. Can you help me with that?

10.4.3 Prompt 2

I walk for an hour every morning and do yoga twice a week. My goal is to reduce stress and improve my flexibility.

10.5 Summary

In these exercises, we explored how AI assistants can simulate domain-specific conversations by adapting to user needs and preferences.

We observed how AI assistants:

- Provide personalized recommendations based on user preferences and goals.
- Streamline complex tasks through intelligent assistance.
- Improve user experiences in financial planning.
- Support health and fitness management through personalized guidance.

11 Applications of AI in Different Industries

11.1 Learning Objectives

After studying this topic, we will be able to:

- Explore the use cases and real-world examples of AI in different industries.
- Explain the impact and benefits of using AI in different industries.

11.2 AI Adoption Across Industries

AI has been evolving like never before, and its involvement in different industries has been increasing rapidly.

A recent global survey conducted by **Statista (statista.com)** on AI adoption in business estimates that:

- **23%** of company Chief Executive Officers (CEOs) have already adopted AI in their business operations.
- **43%** of CEOs plan to explore options for adopting AI in the future.
- Among Chief Medical Officers (CMOs), **32%** have adopted AI.
- **39%** of CMOs plan to explore AI adoption in the future.

These statistics demonstrate the growing interest and adoption of AI across industries.

11.3 Predictive Maintenance

Various industries face challenges such as unexpected machine failures in manufacturing plants and ensuring the reliability of critical medical equipment in the healthcare industry.

Deploying **AI-driven predictive maintenance systems** helps:

- Detect potential failures before they occur.
- Minimize downtime.
- Increase productivity.
- Reduce maintenance expenses.
- Enhance efficiency.
- Improve cost-effectiveness.

11.4 AI in Manufacturing

AI is a driving force behind a new era of efficiency and optimization in manufacturing.

11.4.1 AI-Driven Robotics and Automation

AI-driven robotics and automation transform production processes by handling repeated tasks and assembly. This frees human workers to focus on more complex and creative tasks while increasing:

- Speed
- Accuracy
- Scalability

Real-World Example

- **BMW** uses **cobots (collaborative robots that work alongside humans in a shared workspace)** in its manufacturing units to enhance efficiency and increase productivity.

11.4.2 Image Recognition for Quality Inspection

AI uses image recognition systems to inspect products for defects on assembly lines, ensuring that only top-quality products are delivered to customers.

Benefits include:

- Reduced waste.
- Improved product quality.
- Enhanced customer satisfaction.

11.4.3 Energy Optimization

AI optimizes energy consumption in manufacturing facilities by:

- Analyzing energy usage patterns.
- Identifying opportunities for energy savings.
- Supporting smart grids.
- Using AI-driven energy management systems.

Benefits include:

- Reduced energy costs.
- Improved sustainability.

11.4.4 Food and Beverage Manufacturing

The food and beverage manufacturing industry leverages AI-powered quality control systems to inspect:

- Defects.
- Contamination.
- Freshness in food products.

This ensures that only safe and high-quality food items are packed and distributed.

11.5 AI in Healthcare

AI is a game-changer in healthcare by offering innovative solutions and providing tangible benefits across multiple domains.

11.5.1 Medical Imaging Analysis

AI analyzes medical images such as:

- X-rays

- MRIs (Magnetic Resonance Imaging)
- CT scans (Computed Tomography scans)

This assists radiologists in:

- Detecting abnormalities.
- Making diagnoses.

11.5.2 Predictive Analytics

AI systems analyze:

- Electronic Health Records (**EHRs**)
- Patient data
- Historical trends

This helps to:

- Predict patient outcomes.
- Identify individuals at risk of developing specific conditions.
- Enable healthcare providers to anticipate and prevent diseases.
- Support more proactive and effective patient management.

11.5.3 Operational Efficiency

AI enhances operational efficiency by:

- Optimizing resource allocation.
- Streamlining scheduling.
- Managing supply chains more effectively.

11.5.4 Drug Discovery and Development

One of the most popular applications of AI in healthcare is drug discovery and development.

AI accelerates drug discovery by:

- Examining molecular structures.
- Predicting drug-target interactions.
- Simulating clinical trials.

Real-World Example

BenevolentAI leverages AI algorithms to identify potential drug candidates for various diseases, significantly reducing the time and cost involved in the drug development process.

This can lead to the discovery of novel treatments for diseases such as:

- Cancer

- [Alzheimer's disease](#)

11.6 AI in Finance

In the financial sector, AI serves as a strategic partner by revolutionizing banking, investment, and risk management.

11.6.1 AI-Powered Customer Support

AI-powered chatbots provide:

- Personalized assistance.
- Real-time support for banking transactions.

This enhances the overall customer experience.

11.6.2 Investment Analysis

AI algorithms help customers identify:

- Market trends.
- Investment opportunities.
- Financial risks.

This is achieved by processing large amounts of financial data.

11.6.3 Risk Management

AI enables risk management strategies to become more proactive and precise by:

- Predicting potential threats.
- Mitigating fraud.
- Optimizing portfolio performance.

11.6.4 Robo-Advisors

Robo-advisors use AI algorithms to provide:

- Automated investment suggestions.
- Portfolio management.
- Financial planning services.

11.6.5 Real-World Example: Bank of America – Erica

To provide customers with a more convenient way to manage their finances, **Bank of America** introduced **Erica**, an AI-powered virtual assistant.

[Erica](#) assists customers with tasks such as:

- Balance inquiries.
- Bill payments.
- Budgeting.
- Fraud alerts.
- Financial insights.

11.7 AI in Retail

AI has significantly impacted customer engagement, inventory management, marketing, and the overall retail experience.

11.7.1 Personalized Product Recommendations

AI-powered recommendation engines analyze customer preferences and browsing behavior to provide personalized product recommendations.

Real-World Example

Amazon (amazon.com) uses AI to suggest products based on:

- Customers' past purchases.
- Similar items viewed by other users.

11.7.2 Demand Forecasting

AI-driven demand forecasting algorithms analyze:

- Historical sales data.
- Market trends.
- External factors.

This enables retailers to:

- Predict future product demand.
- Optimize inventory levels.
- Reduce stockouts.
- Minimize excess inventory holding costs.

11.7.3 Personalized Marketing

AI enables retailers to create personalized marketing campaigns and promotions based on individual customer profiles and behavior.

Real-World Example

Salesforce Marketing Cloud uses machine learning algorithms to:

- Segment customers.
- Deliver personalized messages.

- Optimize campaign performance.

11.7.4 AI-Powered Cashierless Stores

Cashierless stores leverage:

- Computer vision.
- Machine learning algorithms.

These technologies enable frictionless retail shopping experiences.

Real-World Example

Amazon Go stores use AI technology to:

- Detect when customers take products from shelves.
- Automatically charge customers through their Amazon account upon exiting the store.

11.8 Future of AI

As AI continues to evolve, its potential for further innovation and growth is boundless. It promises transformative advancements that will reshape industries and revolutionize how we live and work.

11.9 Key Takeaways

In this topic, we learned that:

- AI-driven robotics and image recognition systems have introduced a new era of efficiency in the manufacturing industry.
- AI optimizes energy consumption and improves sustainability in manufacturing.
- AI offers innovative healthcare solutions through medical imaging analysis and predictive analytics.
- AI improves operational efficiency and accelerates drug discovery in healthcare.
- AI enhances customer service, investment analysis, and risk management in the financial sector.
- AI transforms retail through personalized customer engagement, demand forecasting, marketing automation, inventory optimization, and cashierless shopping experiences.

12 Generative AI Tools and Applications

12.1 Overview

After studying this topic, we will be able to:

- Recognize various applications of Generative AI across multiple domains.
- Explore the integration of Generative AI in industry-leading companies.

Generative AI has sparked a revolution across numerous fields by enabling machines to autonomously create new content. It can generate detailed images, fascinating videos, and compelling stories using advanced AI algorithms. This powerful technology has demonstrated immense potential across a wide range of industries.

12.2 Generative AI for Language

Language stands as the most crucial application area for Generative AI.

To create new content, Generative AI tools are built on underlying AI models, such as **Large Language Models (LLMs)**.

Initially, LLMs could only:

- Accept text as input.
- Generate text as output.

When **OpenAI** first released **ChatGPT**, it was built on the text-only LLM **GPT-3**.

With the development of **multimodal LLMs**, these models can now process multiple forms of data, including:

- Text
- Audio
- Images
- Videos

These multimodal models have significantly advanced Generative AI by enabling more accurate solutions to complex problems.

12.3 Large Language Models and Their Capabilities

12.3.1 OpenAI GPT Models

OpenAI's GPT models can now process both:

- Text
- Images

These models excel at solving complex problems with greater accuracy.

12.3.2 Google PaLM Models

Google introduced **PaLM Models**, which excel in textual and linguistic tasks.

Capabilities:

- Text input
- Text output

Limitation:

- Restricted to text-in and text-out workflows.

12.3.3 Google Gemini Models

Google introduced the **Gemini family of models**, which unlocked Generative AI's multimodal capabilities.

Gemini models can:

- Caption images.
- Answer questions about photographs.
- Describe videos.
- Discuss multimedia content.

Gemini is the foundational model powering the functionalities of the **Google Gemini** tool used today.

12.3.4 Other Major Large Language Models

Other LLMs are also transforming content creation and interaction, including:

- Amazon **Titan Models**
- Meta **Llama Models**
- Anthropic **Claude Models**

These models are revolutionizing how we create and interact with content.

12.4 Generative AI for Visual Arts and Design

Generative AI has significantly influenced the visual arts and design industry by empowering artists and designers with new creative tools and techniques.

12.4.1 Stable Diffusion

Stable Diffusion is among the most advanced technologies for **text-to-image generation**.

12.4.2 DALL·E

The **DALL·E** model can generate images that precisely match the input text description.

12.4.3 StyleGAN

StyleGAN is used for producing high-quality images of:

- Human faces
- Other objects

12.4.4 Super Resolution

Super Resolution is a model used for enhancing image resolution by increasing the number of pixels.

12.5 Generative AI for Voice and Music Generation

Along with text and image generation, Generative AI also facilitates voice and music generation.

12.5.1 Murph

Murph is a platform at the forefront of AI voice generation technology.

It excels in generating synthetic voices that closely replicate human speech nuances and tones.

12.5.2 OpenAI Whisper

OpenAI Whisper is an open-source model that enables:

- Transcription in multiple languages.
- Translation from those languages into English.

12.5.3 AI Music Generation

AI-powered music generators can:

- Produce a wide variety of musical genres.
- Create music ranging from classical compositions to contemporary beats.
- Tailor music to evoke specific moods, from cheerful melodies to melancholic tunes.

Musicians, producers, filmmakers, videographers, and businesses are all experimenting with Generative AI tools.

12.5.4 Jukedek

Jukedek uses Generative AI algorithms to compose original music tracks based on user input.

12.5.5 Amper Music

Amper Music also uses Generative AI algorithms to compose original music based on user requirements.

12.5.6 AIVA

AIVA enables users to generate new songs in **more than 250 different styles** within seconds.

Whether we are complete beginners or experienced professionals, these AI tools allow us to create our own songs.

12.6 Generative AI for Video Generation

Generative AI algorithms can generate videos that closely resemble reality.

These algorithms analyze existing human features and movements to create:

- Realistic characters.
- Lifelike backgrounds.

Generative AI models can:

- Generate visuals.
- Craft compelling stories.
- Produce highly engaging videos.

12.6.1 Google's Imagen Video

Google's Imagen Video is a machine learning model capable of generating high-definition videos.

12.6.2 OpenAI Sora

OpenAI Sora is a model that creates realistic and imaginative scenes from text instructions.

12.7 Adoption of Generative AI by Leading Companies

Many organizations are leveraging the capabilities of Generative AI.

According to a recent poll conducted by **Gartner**, **55% of organizations** (more than half) are either in the piloting stage or production mode with Generative AI.

12.7.1 Google

Google utilizes Generative AI in several products, including:

- **Google Photos** for image enhancements.
- **Google Duplex** for natural language understanding.
- **Google Magenta** for music generation.

12.7.2 Salesforce

Salesforce and **OpenAI** introduced a ChatGPT application named **Einstein** for the **Slack** platform.

The application leverages ChatGPT's AI technology.

12.7.3 Adobe

Adobe utilizes Generative AI through **Adobe Sensei**, the company's AI and machine learning platform.

Adobe Sensei powers features such as:

- Automated photo editing.
- Font recognition.

12.7.4 IBM

IBM introduced **watsonx**, an advanced AI and data platform that helps businesses create custom AI applications.

It provides tools for:

- Training AI models.
- Managing data.
- Following rules (governance).
- Working efficiently with other systems (integration).

12.8 Key Takeaways

- Generative AI has revolutionized multiple industries, leading to a wide variety of applications.
- Large Language Models (LLMs) such as **GPT**, **PaLM**, and **Gemini** are widely used for text generation.
- Popular tools such as **ChatGPT** and **Google Gemini** are built on these LLMs.
- Generative AI supports image generation through models such as **Stable Diffusion** and **DALL·E**.
- Voice and music generation have advanced through tools such as **Murph**, **AIVA**, **Jukedeck**, and **Amper Music**, enabling the creation of music across different genres and moods.
- Generative AI algorithms can create lifelike videos using models such as **Google's Imogen Video** and **OpenAI Sora**.
- Leading companies, including **Google**, **Salesforce**, **Adobe**, and **IBM**, have adopted and integrated Generative AI into their products and services.

13 Machine Learning Use Cases in Everyday Life

Machine learning (ML) is already a significant part of our everyday lives. Although everybody is talking about **Generative AI (Gen AI)**, it is important to understand that **Gen AI is a subset of the larger field of Machine Learning**.

Machine learning is a subfield of **Artificial Intelligence (AI)** in which machines learn from **datasets** and **past experiences** by recognizing patterns and generating predictions.

Machine learning is projected to become a **\$200 billion industry by 2029**, but it is already being widely used across many industries and everyday applications.

13.1 Natural Language Processing (NLP)

One of the most impactful areas of machine learning is **Natural Language Processing (NLP)**, which is the ability of machines to make sense of the unstructured nature of human language.

NLP enables computers to understand, interpret, and process text and spoken language, making many intelligent applications possible.

13.2 Customer Service

Machine learning powers many customer service solutions through **chatbots**, which act as virtual agents on many business and e-commerce websites.

Chatbots can:

- Handle text-based customer queries.
- Resolve many customer issues automatically.
- Route customers to the appropriate human customer service representative when necessary.

13.3 Voice Assistants and Speech Recognition

Machine learning also powers voice assistants such as:

- Siri
- Alexa

These systems work by:

- Converting spoken language into text (**Speech-to-Text**).
- Using NLP machine learning models to understand spoken commands.

The same capability is also used by services such as:

- Slack
- YouTube

These platforms use machine learning for **automatic transcription (autotranscription)** of spoken words in video content.

13.4 Mobile Applications and Smartphones

Many mobile applications rely on machine learning services.

Examples include:

- **Spotify**, which uses ML models to generate song recommendations.
- **LinkedIn**, which uses ML to provide employment suggestions.

Modern smartphones are also capable of performing certain machine learning tasks directly on the device.

Examples include:

- Computational photography for generating background blur in selfie photos.
- Facial recognition for unlocking smartphones.
- On-device image classification models that help us search our photo libraries.

For example, machine learning can quickly locate a specific image—such as a picture of a cat inside a dryer—from thousands of stored photos without requiring us to scroll manually.

(The dryer was not actually on.)

This is an example of a **needle in a haystack problem** (finding one specific item among a very large collection).

13.5 Financial Transactions and Fraud Detection

Machine learning plays a major role in detecting fraudulent financial transactions.

In the **United States alone**, there are approximately:

- **150 million credit card transactions every day.**
- This equals approximately **1,739 transactions every second.**

The vast majority of these transactions are legitimate, making fraud detection extremely challenging.

Financial institutions train machine learning and **deep learning** models using classification algorithms to:

- Recognize suspicious online transactions.
- Flag potentially fraudulent transactions for further investigation.

Processing **1,739 transactions every second** would be nearly impossible to perform manually, making machine learning essential.

Machine learning is also heavily used in stock market trading.

Between **60% and 73% of all stock market trading** is conducted by machine learning algorithms, and this percentage continues to increase every year.

13.6 Cybersecurity

Machine learning is frequently used in cybersecurity.

For example, **Reinforcement Learning** (a type of machine learning where models learn through trial and error by receiving rewards or penalties for their actions) is used to train models that can:

- Identify cyberattacks.
- Respond to cyber threats.
- Detect system intrusions.

13.7 Transportation

Machine learning plays an important role in modern transportation.

Examples include:

- **Google Maps**, which uses machine learning algorithms to analyze current traffic conditions and determine the fastest route.
- Ride-sharing platforms such as **Uber** and **Lyft**, which use machine learning to match riders with drivers.

13.8 Email Services

Machine learning is widely used in email systems.

Applications include:

- Classification of incoming email messages.
- Spam filtering.
- Autocomplete responses.

13.9 Healthcare

Healthcare is one of the most valuable areas where machine learning augments (supports and enhances) human capabilities.

One example is the evaluation of **mammograms**(X-rays of breast tissue).

According to the provided information:

- Doctors evaluating mammograms miss between **30% and 40% of cancers**.
- The rate of **false positives**(Doctors frequently suspect cancer where there isn't any, causing unnecessary stress, biopsies, and follow-up scans for patients) is even higher.

Machine learning helps by using **pattern recognition models** to classify tumors that are difficult for the human eye to detect.

This improves:

- The accuracy of interpreting radiology images.
- The reading time of radiologists.

As a result, radiologists can focus more attention on suspicious examinations that have been flagged by machine learning models.

Machine learning has also shown success in:

- Early lung cancer screening.
- Detecting bone fractures.

13.10 Marketing and Sales

According to **Forbes**, the department that uses AI and machine learning the most is the **Marketing and Sales department**.

Marketers use machine learning for:

- Lead generation.
- Data analytics.
- Search Engine Optimization (SEO).

They also build on top of existing machine learning models.

For example, recommendation algorithms used by **Netflix** analyze our tastes and interests to suggest TV shows and movies we are likely to watch next.

Marketing and sales teams can use these same recommendation models to create:

- Targeted marketing campaigns.
- Personalized advertisements tailored to those same tastes and interests.

13.11 Machine Learning vs. Artificial General Intelligence (AGI)

There is significant discussion today about the future of AI, particularly **Artificial General Intelligence (AGI)**.

AGI refers to AI that would one day match or surpass human intelligence.

However:

- AGI does **not currently exist**.
- It remains **theoretical**.

Machine learning, on the other hand, is a practical form of AI that already exists today and has become an integral part of our everyday lives.

14 Generative AI Tools in Action

14.1 Learning Objectives

After completing these exercises, we will be able to:

- Leverage common Generative AI tools that can help with daily tasks.
- Explore **Google Gemini**, **OpenAI ChatGPT**, and **Microsoft Copilot** for different everyday use cases.

14.2 Introduction

Generative AI models create new content by understanding context and patterns learned from large datasets containing text, code, images, and audio. These models power various AI tools that assist with creative writing, travel planning, learning new concepts, writing emails, summarizing content, and many other daily tasks.

In this topic, we explore how different Generative AI tools can be used effectively across multiple real-world scenarios.

14.3 Exploring Google Gemini

In this exercise, we explore **Google Gemini**, an AI-powered assistant capable of understanding text, audio, images, and video. We use Gemini for learning new skills and generating creative content through descriptive prompts.

14.3.1 Scenario: Learning a New Game

Prompt

I want to learn how to play chess, but I'm a complete beginner. Can you help me get started?

14.3.2 Scenario: Writing a Blog Post About the Revival of Board Games

Prompt 1

Develop a blog post outline on "The Board Game Boom" with content sections tailored to specific age groups.

Prompt 2

Can you provide a table or list that recommends specific board games for different age groups?

14.4 Experimenting with ChatGPT

In this exercise, we explore **ChatGPT** and how it can generate human-like conversations and assist with tasks such as summarization, brainstorming, content generation, and travel planning.

14.4.1 Scenario: Summarizing Text

Prompt

Summarize the article on "AI this Earth Day: Top opportunities to advance sustainability initiatives" based on the IBM blog post.

14.4.2 Scenario: Planning a Weekend Trip

Prompt 1

I need help planning a weekend getaway. Any suggestions for places to visit within a 3-hour drive from Delhi?

Prompt 2

Jim Corbett National Park sounds amazing. Any recommendations for a hotel or a place to stay there?

Prompt 3

That sounds perfect. Could you help me with an itinerary for a two-day trip?

14.5 Experimenting with Microsoft Copilot

In this exercise, we explore **Microsoft Copilot** and its capabilities for creative writing and image generation. We use Copilot to develop story ideas and generate illustrations from text prompts.

14.5.1 Scenario: Creating the Plot for a Fictional Story

Prompt 1

I want to create a short science fiction story for kids centered on the friendship between a young boy and a robot. Could you suggest some interesting plot ideas?

Prompt 2

With regard to the Robot Memories plot idea, could you provide a related illustration?

14.6 Summary

In these exercises, we explored the capabilities of three popular Generative AI tools:

- **Google Gemini**
- **OpenAI ChatGPT**
- **Microsoft Copilot**

We learned how these tools can assist with:

- Learning new skills.
- Creative writing.
- Content summarization.
- Travel planning.
- Brainstorming ideas.
- Generating text through effective prompting.
- Generating visual content through effective prompting.

15 Module Summary

At this point in the module, we know the following concepts and topics.

15.1 Artificial Intelligence (AI)

Artificial Intelligence (AI) is an augmented intelligence aiming to extend human capabilities and tackle tasks beyond human and machine capabilities.

Machine learning models are developed through:

- Supervised learning
- Unsupervised learning
- Reinforcement learning

AI can be divided based on its strengths:

- **Weak (Narrow) AI:** Specific domains.
- **Strong (Generalized) AI:** Diverse capabilities across unrelated tasks.
- **Super AI:** Hypothetical AI that surpasses human intelligence in all domains.

15.2 Generative AI

Generative AI is a type of artificial intelligence that can generate content across various media, such as:

- Written text
- Images
- Audio
- Videos

Large Language Models (LLMs) are advanced neural network architectures for processing and generating human-like text.

The capabilities of Generative AI include:

- Diverse content creation capabilities
- Human-like capabilities
- Data augmentation

15.2.1 Domain-Specific Use Cases of Generative AI

- Marketing
- Creative
- Product development

15.2.2 Industry-Specific Use Cases of Generative AI

- Healthcare
- Gaming
- Fashion
- Education and training

15.3 Benefits of AI

AI provides numerous benefits across daily life and technology.

- Virtual assistants and smart home devices use AI to automate routine tasks, making our lives more efficient and convenient.
- AI uses recommendation systems to provide personalized suggestions across various streaming, social media, and e-commerce platforms.
- AI boosts security and safeguards us against potential threats through biometric authentication and fraud detection.
- AI integration into smart devices enhances user experience and functionality.

15.4 Generative AI Use Cases

15.4.1 Text Generation

Various Large Language Models (LLMs), such as GPT, PaLM, and Gemini, are used for text generation.

Tools based on these models include:

- ChatGPT
- Google Gemini

15.4.2 Image Generation

Generative AI helps generate images using advanced models such as:

- Stable Diffusion
- DALL·E

15.4.3 Voice and Music Generation

Generative AI enhances voice and music generation by creating different genres and moods using tools such as:

- Murf
- AIVA

15.4.4 Video Generation

Generative AI algorithms can create lifelike videos using models such as:

- Google's Imagen Video
- OpenAI's Sora

15.5 AI Chatbots and Smart Assistants

AI chatbots and smart assistants are AI-driven software programs that understand and respond to queries, provide information, and perform tasks.

Chatbots have evolved from:

- Simple rule-based systems
- AI-powered personal assistants
- Generative AI chatbots capable of complete conversations

15.5.1 Benefits of AI Chatbots

- 24/7 availability
- Scalability
- Personalized services and suggestions
- Natural conversations
- Multilingual communication

15.5.2 Applications of AI Chatbots

AI chatbots are widely used across industries, including:

- Customer service
- E-commerce
- Healthcare
- Education

Their applications include:

- More human-like interactions
- Sentiment analysis
- Automated HR support tasks
- Automated IT support tasks
- Handling processes such as refunds

15.6 AI's Impact on Industries

AI has a transformative influence across various industries.

Industry	AI Applications
Manufacturing	AI-driven robotics, image recognition systems
Healthcare	Medical imaging analysis, predictive analytics, operational efficiency
Finance	Enhance customer service, investment analysis

Retail	Customer engagement, inventory management, marketing, cashier-less stores
---------------	---

Module 2: AI Concepts, Terminology and Application Domain

16 Cognitive Computing

16.1 Overview

Cognitive computing is a branch within Artificial Intelligence (AI) that mimics human thought processes to create intelligent machines. It aims to provide enhanced functionality, adaptability, and intelligence across various domains by creating systems that simulate human cognitive processes such as thinking, reasoning, and problem-solving.

Unlike traditional systems that simply execute commands, cognitive computing makes machines active partners. These systems can understand our needs, anticipate our questions, and proactively provide valuable insights.

For example:

- A bank can use a cognitive computing system to detect fraudulent transactions.
- A company can improve customer support by using a chatbot powered by cognitive computing.
- Cognitive computing can evaluate an individual's performance and provide personalized recommendations, making learning more adaptive than traditional predetermined learning paths.

16.2 Human Decision-Making Process

When we seek to understand something and make a decision, we generally go through four key steps:

- Observe visible phenomena and bodies of evidence.
- Draw on existing knowledge to interpret what we are seeing and generate a hypothesis about what it means.
- Evaluate which hypotheses are right or wrong.
- Decide by choosing the option that seems best and acting accordingly.

Just as humans become experts through observation, evaluation, and decision-making, cognitive systems use similar processes to reason about the information they process, but they can perform these tasks at massive speed and scale.

16.3 Core Elements of Cognitive Computing

Cognitive computing consists of three core elements:

- Perception
- Learning
- Reasoning

These elements work together to simulate human-like intelligence and decision-making.

16.3.1 Perception

Perception is the foundation of cognitive computing. It enables the system to interpret and understand its environment.

The perception process includes:

- Sensing, where structured or unstructured data is gathered from various sources.
- Interpreting the collected data to understand the surrounding environment.

16.3.2 Learning

Learning enables cognitive systems to improve continuously by analyzing data.

Machine learning algorithms are used to:

- Analyze structured and unstructured data.
- Extract meaningful information from the data.
- Identify patterns and trends.
- Understand complex relationships within the data.
- Continuously improve through repetitive learning processes.

16.3.3 Reasoning

Reasoning enables cognitive systems to:

- Evaluate available information.
- Make accurate predictions.
- Support intelligent decision-making based on data analysis.

16.4 How Cognitive Computing Works

The cognitive computing process follows these steps:

- Gather structured and unstructured data from multiple sources.

- Analyze the data using machine learning algorithms.
- Extract meaningful information.
- Identify patterns and trends.
- Understand complex relationships.
- Generate accurate predictions.
- Continuously improve through repetitive learning.

This continuous cycle demonstrates the evolving nature of cognitive computing.

16.5 Benefits of Cognitive Computing

Cognitive computing provides several important benefits:

- Elevated decision-making through the analysis of huge datasets.
- Improved efficiency by saving time and resources through task automation.
- More human-like and interactive communication between humans and machines using Natural Language Processing (NLP).
- Enhanced adaptability and intelligence across various domains.
- Personalized recommendations based on user performance and behavior.

16.6 Applications of Cognitive Computing

Cognitive computing has applications across many industries, including:

- Healthcare
- Finance
- Education
- Entertainment
- Customer Service
- Retail

16.7 Real-World Examples

Several organizations use cognitive computing technologies to improve their products, services, and operations.

16.7.1 IBM Watson

IBM Watson uses cognitive computing in:

- Healthcare
- Finance
- Retail
- Customer Service

16.7.2 Google

Google applies cognitive computing techniques in:

- Google Search
- Google Assistant
- Google Translate

16.7.3 Amazon

Amazon's Alexa virtual assistant uses cognitive computing to:

- Understand and respond to voice commands.
- Manage smart home devices.
- Provide personalized recommendations to users.

16.7.4 JPMorgan Chase and Wells Fargo

JPMorgan Chase and Wells Fargo use cognitive computing to:

- Detect fraud.
- Assess risk.
- Automate customer service.

16.8 Key Takeaways

- Cognitive computing is a branch of AI that mimics human cognitive processes such as thinking, reasoning, and problem-solving.
- It transforms machines from simple tools into active partners that understand needs, anticipate questions, and provide valuable insights.
- Human decision-making involves observation, hypothesis generation, evaluation, and decision-making, and cognitive systems follow similar processes at massive speed and scale.
- The three core elements of cognitive computing are:
 - **Perception** – Gathering and interpreting data from various sources.
 - **Learning** – Using machine learning algorithms to analyze data and extract meaningful information.
 - **Reasoning** – Making accurate predictions and intelligent decisions based on data analysis.
- Cognitive computing continuously improves by learning from patterns and relationships within data.
- Major benefits include improved decision-making, increased efficiency, automation of tasks, and more interactive human-machine communication through Natural Language Processing (NLP).
- Cognitive computing is widely used in healthcare, finance, education, entertainment, retail, and customer service.
- Organizations such as **IBM Watson**, **Google**, **Amazon (Alexa)**, **JPMorgan Chase**, and **Wells Fargo** use cognitive computing to enhance their products, services, and operations.

17 Terminologies and Related Concepts of Artificial Intelligence (AI)

17.1 Overview

In this topic, we explore the key terminologies and related concepts of Artificial Intelligence (AI). Understanding these concepts enhances our knowledge and enables us to leverage AI's full potential, drive innovation, and stay ahead in our field.

AI enables a world where machines can not only understand human language but also predict our needs, recognize our faces, and provide safety against threats.

17.1.1 Real-World Example

Autonomous vehicles (self-driving cars) rely heavily on AI technologies such as:

- Machine Learning
- Deep Learning
- Natural Language Processing (NLP)
- Computer Vision

These technologies enable vehicles to navigate and make real-time decisions. Understanding these key terms provides valuable insights into how these vehicles operate, along with their benefits and challenges.

17.2 Artificial Intelligence (AI)

Artificial Intelligence (AI) is a branch of computer science that focuses on creating systems capable of performing tasks that typically require human intelligence.

AI systems typically demonstrate behaviors associated with human intelligence, including:

- Planning
- Learning
- Reasoning
- Problem solving
- Knowledge representation
- Perception
- Motion
- Manipulation
- Social intelligence
- Creativity

17.2.1 Categories of Artificial Intelligence

AI is categorized into three types:

AI Category	Description
-------------	-------------

Narrow AI (Weak AI)	Performs specific tasks.
General AI (Strong AI)	Possesses human-like cognitive skills and is capable of learning and adapting across various tasks.
Super AI	Aims to surpass human intelligence. It is currently theoretical and has not yet been realized.

17.3 Machine Learning

Machine Learning (ML) is a subset of Artificial Intelligence that uses computer algorithms to analyze data and make intelligent decisions based on what the system has learned, without being explicitly programmed.

Machine Learning algorithms:

- Are trained with large datasets.
- Learn from examples.
- Do not follow rule-based algorithms.
- Enable machines to solve problems on their own.
- Make accurate predictions using the provided data.

17.4 Deep Learning

Deep Learning is an important concept and a specialized subset of Machine Learning.

It uses multi-layered neural networks, known as **Deep Neural Networks** (neural networks with many hidden layers that can learn highly complex patterns), to analyze complex data and simulate human decision-making.

Deep Learning algorithms can:

- Label information.
- Categorize information.
- Identify complex patterns in data.
- Learn continuously.
- Improve the quality and accuracy of results by evaluating the correctness of decisions.

17.5 Neural Networks

Neural Networks are a computational model inspired by the human brain's neural structure.

Neural Networks consist of interconnected nodes (neurons) and have three layers.

Layer	Function
Input Layer	Receives and processes raw data.
Hidden Layers	Perform complex computations and transform the data.
Output Layer	Converts the processed data into the output format and produces the final result.

17.6 Key Takeaways

At the end of this topic, we understand that:

- Artificial Intelligence (AI) simulates human intelligence in machines.
- AI is divided into three categories:
 - Narrow AI (Weak AI) for specific tasks.
 - General AI (Strong AI) with human-like cognitive abilities.
 - Super AI, which aims to surpass human intelligence and is currently theoretical.
- Machine Learning is a subset of AI that:
 - Uses algorithms to analyze data.
 - Makes decisions without explicit programming.
 - Enables autonomous problem solving.
- Deep Learning is a specialized subset of Machine Learning that uses multi-layered neural networks to analyze complex data.
- Neural Networks are computational models consisting of interconnected nodes (neurons) organized into three layers:
 - Input Layer
 - Hidden Layers
 - Output Layer

18 Machine Learning Fundamentals

18.1 Introduction

Machine learning is a subset of Artificial Intelligence (AI) that uses computer algorithms to analyze data and make intelligent decisions based on what it has learned.

Instead of following rules-based algorithms, machine learning builds models to classify data and make predictions from it.

18.2 Understanding Machine Learning with a Heart Failure Prediction Example

A practical example of machine learning is determining whether a heart can fail.

Suppose we are given the following dataset:

- Beats per minute
- Body Mass Index (BMI)
- Age
- Sex
- Result indicating whether the heart has failed or not

Using this data, machine learning can learn patterns and create a model that predicts whether a heart is likely to fail based on the input values.

18.3 Traditional Programming vs. Machine Learning

18.3.1 Traditional Programming

An algorithm is a mathematical technique.

In traditional programming:

- We take **data** and **rules**.
- We use them to develop an algorithm that produces an answer.
- The algorithm is typically based on predefined logic, such as **if-then-else** statements.
- When we submit inputs, the algorithm produces answers according to the fixed rules we created.
- The algorithm does **not** change unless we manually modify it.

For the heart failure example, a traditional algorithm would use values such as beats per minute and BMI to create fixed rules that determine whether the heart will fail.

18.3.2 Machine Learning

Machine learning follows a different approach.

Instead of taking **data and rules** to produce answers, machine learning takes **data and answers** to create the rules.

Rather than manually defining the logic, the machine learning model learns the rules automatically from the data.

The model:

- Determines the rules and decision logic.
- Learns the equivalent of the traditional **if-then-else** statements.
- Identifies the parameters that best explain the relationship between the inputs and the outputs.
- Can be continuously trained with new data.
- Improves over time and can be used to predict future values.

Instead of arbitrarily deciding that beats per minute plus BMI equals a certain result, we allow the model to determine the appropriate logic based on the data.

18.4 How Machine Learning Works

Machine learning relies on defining behavioral rules by examining and comparing large datasets to identify common patterns.

For example:

- We can provide a machine learning program with a large collection of bird images and label each image as **bird**.
- We can also provide images of cats and label them as **cat**.

- During training, the model learns the characteristics associated with each label.
- When presented with a new image, the model predicts whether it is a bird or a cat and returns its prediction with a certain level of confidence.

18.5 Types of Machine Learning

There are three main types of machine learning:

- Supervised Learning
- Unsupervised Learning
- Reinforcement Learning

18.6 Supervised Learning

Supervised learning trains an algorithm using **human-labeled data**.

The algorithm learns the relationship between the inputs and their corresponding labels.

Example:

- Provide many images labeled **bird**.
- Provide many images labeled **cat**.
- The algorithm learns from these labeled examples.
- When shown a new image, it predicts whether the image is a bird or a cat with some level of confidence.

The more labeled samples we provide to a supervised learning algorithm, the more precise it becomes at classifying new data.

18.7 Unsupervised Learning

Unsupervised learning relies on **unlabeled data**.

Instead of providing labels, we only provide the input data and allow the algorithm to discover patterns on its own.

The algorithm:

- Ingests unlabeled data.
- Draws inferences.
- Identifies hidden patterns.
- Groups similar data together.

This approach is particularly useful for **clustering** (grouping similar data points together based on their similarity while separating them from dissimilar data).

18.7.1 Network Traffic Example

Suppose we provide a machine learning algorithm with a continuous stream of network traffic.

The algorithm independently learns:

- The baseline (normal network activity).
- Outlier behavior.
- Potentially malicious network activity.

After clustering the data, different analytical techniques can be applied to further explore the discovered patterns.

18.8 Reinforcement Learning

Reinforcement learning provides the algorithm with:

- A set of rules.
- A desired goal.
- Allowed actions.
- Constraints.

The algorithm learns how to achieve the goal by trying different combinations of allowed actions.

During learning:

- Good decisions are rewarded.
- Poor decisions are punished.
- The algorithm attempts to maximize its total rewards while operating within the defined constraints.

18.8.1 Applications

Reinforcement learning can be used to teach machines to:

- Play chess.
- Navigate an obstacle course.

18.9 Machine Learning Model

A machine learning model is an algorithm used to find patterns in data without the programmer having to explicitly program those patterns.

Instead of manually defining every rule, the model automatically learns the relationships from the available data.

18.10 Key Takeaways

- Machine learning is a subset of Artificial Intelligence (AI) that uses computer algorithms to analyze data and make intelligent decisions.
- Unlike traditional programming, which relies on fixed rules, machine learning models learn from data and evolve over time.

- Traditional programming uses **data + rules → answers**, whereas machine learning uses **data + answers → rules (model)**.
- Machine learning models can be continuously trained to improve their predictive performance.
- Machine learning identifies patterns by analyzing large datasets.

The three primary types of machine learning are:

Machine Learning Type	Description	Common Applications
Supervised Learning	Trained using labeled data to classify or predict new data. The more labeled samples provided, the more accurate the model becomes.	Image classification, heart failure prediction
Unsupervised Learning	Finds patterns in unlabeled data. Useful for clustering similar data points and detecting anomalies.	Customer segmentation, network anomaly detection
Reinforcement Learning	Learns to achieve goals within defined rules and constraints by maximizing rewards and minimizing penalties.	Chess playing, obstacle course navigation, robotics

19 Machine Learning: Techniques and Training

19.1 Learning Objectives

After studying this topic, we will be able to:

- Explain the core concepts of machine learning, including:
 - Supervised learning
 - Unsupervised learning
 - Reinforcement learning
- Discuss the supervised learning tasks:
 - Regression
 - Classification
 - Neural Networks
- Describe the process of training a machine learning model.
- Explain the roles of the training, validation, and test datasets.

19.2 Machine Learning Categories

Machine learning is a broad field that can be divided into three categories:

- Supervised Learning
- Unsupervised Learning
- Reinforcement Learning

There are many different tasks that we can solve using these categories.

19.2.1 Supervised Learning

Supervised learning involves using datasets with predefined class labels to train models for predicting or classifying new data points.

This means the data we receive already includes labels indicating what each data point represents.

For example, a dataset may contain labels such as:

- Age
- Sex

These labeled datasets are used to train machine learning models.

19.2.2 Unsupervised Learning

With unsupervised learning, we do not have class labels. Instead, we must discover class labels from unstructured data.

This could involve deep learning, which looks at pictures in unstructured data, interprets them, processes them, and further groups them into clusters based on their features.

19.2.3 Reinforcement Learning

Reinforcement learning is another subset of machine learning.

It uses a reward function to:

- Penalize bad actions.
- Reward good actions.

The model learns by maximizing rewards while minimizing penalties.

19.3 Supervised Learning Tasks

Supervised learning can be divided into three categories:

- Regression
- Classification
- Neural Networks

19.3.1 Regression

Regression models are built by looking at the relationships between features $\mathbf{x}$ and the result $\mathbf{y}$, where $\mathbf{y}$ is a continuous variable.

Essentially, regression estimates continuous values.

19.3.2 Neural Networks

Neural networks refer to structures that imitate the structure of the human brain to:

- Process input data.
- Recognize patterns.
- Make decisions or predictions.

19.3.3 Classification

Classification focuses on discrete values.

We can assign discrete class labels $\mathbf{y}$ based on many input features $\mathbf{x}$.

For example, given a set of features $\mathbf{x}$, such as:

- Beats per minute
- Body Mass Index (BMI)
- Age
- Sex

The algorithm classifies the output $\mathbf{y}$ into two categories:

- True
- False

For example, it predicts whether the heart will fail or not.

Classification models can also categorize results into multiple groups.

For instance, we can predict whether a movie belongs to one of the following genres:

- Action
- Comedy
- Drama
- Horror

Some common classification algorithms include:

- Decision Trees
- Support Vector Machines (SVM)
- Logistic Regression
- Random Forests

19.4 Features and Data Points

Classification allows us to extract features from the data.

In the previous example, the features include:

- Beats per minute
- Age

Features are distinctive properties of input patterns that help determine the output categories or classes.

In a dataset:

- Each column is a feature.
- Each row is a data point.

Classification is the process of predicting the class of given data points.

The classifier uses training data to understand how the given input variables relate to the output class.

19.5 Model Training

Training refers to using a learning algorithm to determine and develop the parameters of a machine learning model.

19.5.1 Example: Spam Email Classification

Suppose we are training a model to predict whether an email is spam.

Output labels:

- **True** → Spam
- **False** → Not Spam

The training process is as follows:

- We first show the algorithm real-life emails labeled as **True**, indicating they are spam.
- Next, we show emails labeled as **False**, indicating they are not spam.
- We continue providing additional emails labeled as **True** or **False**.
- The algorithm continually adjusts its internal parameters until it learns to distinguish spam emails from non-spam emails.

19.6 Dataset Splitting

In machine learning, a dataset is usually divided into three parts.

19.6.1 Training Set

The training set is used to train the algorithm.

19.6.2 Validation Set

The validation set helps:

- Validate the results.
- Adjust the algorithm's parameters.

19.6.3 Test Set

The test set contains data that the model has not seen before.

It is used to assess the model's performance on unseen data.

19.7 Model Evaluation Metrics

After testing the model, we can measure its effectiveness using metrics such as:

- Accuracy
- Precision
- Recall

19.8 Key Takeaways

- Machine learning consists of three major categories:
 - Supervised Learning
 - Unsupervised Learning
 - Reinforcement Learning
- Supervised learning builds models using labeled data.
- Unsupervised learning discovers patterns in unlabeled data, often through clustering.
- Reinforcement learning employs a reward system to guide actions.
- Supervised learning is divided into:
 - Regression
 - Classification
 - Neural Networks
- Regression estimates continuous values by analyzing the relationship between input features ($\mathbf{x}$) and output ($\mathbf{y}$).
- Neural networks imitate the structure of the human brain to process data, recognize patterns, and make predictions.
- Classification assigns discrete class labels based on input features.
- Model training involves splitting the dataset into:
 - Training Set
 - Validation Set
 - Test Set
- The training set trains the algorithm.
- The validation set is used to fine-tune and validate the model.
- The testing set evaluates the model's performance on previously unseen data.

20 Deep Learning Fundamentals and Applications

20.1 Introduction

Deep learning is a specialized subset of machine learning, while machine learning itself is a subset of artificial intelligence (AI).

Deep learning layers algorithms to create a neural network, which is an artificial replication of the brain's structure and functionality. This enables AI systems to continuously learn on the job and improve the quality and accuracy of their results.

Unlike traditional machine learning, deep learning allows AI systems to learn from unstructured data, such as:

- Photos
- Videos
- Audio files

For example, deep learning enables the natural language understanding capabilities of AI systems, allowing them to work out the context and intent of what is being conveyed.

20.2 How Deep Learning Works

Deep learning algorithms do not directly map input to output. Instead, they rely on several layers of processing units.

The process works as follows:

- Each layer processes the input it receives.
- The processed output is passed to the next layer.
- The next layer performs additional processing and forwards the result to another layer.
- This layered processing continues until the final output is produced.

The presence of many processing layers is why it is called **deep learning**.

20.3 Developing and Training Deep Learning Models

When creating deep learning algorithms, developers and engineers configure:

- The number of layers in the neural network.
- The type of functions that connect the outputs of one layer to the inputs of the next layer.

After configuring the network, they train the model by providing it with a large number of annotated examples (data that includes labels describing the correct output).

For example:

- Thousands of images are provided to the deep learning algorithm.
- Each image is accompanied by a label that corresponds to the context of that image.
- The algorithm processes these examples through its layered neural network.
- During training, it adjusts the weights of the variables in each layer of the neural network.
- This enables the model to detect the common patterns that define images with similar labels.

20.4 Advantages of Deep Learning

Deep learning addresses one of the major limitations found in older generations of learning algorithms.

- The efficiency and performance of traditional machine learning algorithms eventually plateau (stop improving significantly) as datasets grow.
- Deep learning algorithms continue to improve as they are provided with more data.

This ability to scale with increasing amounts of data makes deep learning highly effective for solving complex problems.

20.5 Applications of Deep Learning

Deep learning has proven to be highly efficient in a wide range of applications, including:

- Image captioning
- Voice recognition
- Voice transcription
- Facial recognition
- Medical imaging
- Language translation
- Natural language understanding (understanding the context and intent of language)
- Driverless car technology

20.6 Key Takeaways

- Deep learning is a specialized subset of machine learning.
- It layers algorithms to create artificial neural networks that replicate the structure and functionality of the human brain.
- Deep learning enables AI systems to continuously learn and improve their performance.
- It can learn from unstructured data such as photos, videos, and audio files.
- Deep learning enhances AI's natural language understanding by grasping the context and intent of communication.
- Developers train deep learning models using large collections of annotated examples.
- During training, the model adjusts the weights of variables across multiple neural network layers to identify common patterns.
- Unlike traditional machine learning algorithms, deep learning continues to improve as more data becomes available.
- Deep learning is highly effective for image captioning, voice recognition, voice transcription, facial recognition, medical imaging, language translation, and driverless car technology.

21 Neural Networks

Neural networks are computational models influenced by the human brain's neural structure. They enable machines to learn in a manner similar to the human brain, making it possible to perform tasks such as recognizing faces and images or driving cars independently.

21.1 Overview of Neural Networks

An artificial neural network (ANN) consists of interconnected nodes known as **neurons**. Like the human brain's neurological network, these neurons take incoming data and learn to make decisions over time.

For example, by providing the network with images of cats and dogs, it can learn to recognize patterns and form connections. The more data the network is exposed to, the more effectively it learns.

21.2 Structure of a Neural Network

A neural network consists of:

- **One input layer**
- **One or more hidden layers**
- **One output layer**

21.2.1 Input Layer

The input layer receives the data.

For example, in an image recognition task, the input layer takes in the image's pixel values.

21.2.2 Hidden Layers

Hidden layers process the data.

Each hidden layer transforms the input data by applying an **activation function** (a mathematical function that enables the network to learn complex patterns).

As a network includes more hidden layers, it becomes **deeper**, which is why the term **deep learning** is used.

21.2.3 Output Layer

The output layer produces the final result of the network's processing.

21.3 How Neural Networks Learn

Neural networks learn through a process called **training**.

21.3.1 Forward Pass (Forward Propagation)

During the **forward pass** (also called **forward propagation**):

- Data passes through the layers of the neural network.
- The network computes an output.

- The computed output is compared to the actual answer.
- The difference between the predicted output and the actual answer is called the **error** or **loss**.
- This step shows how well the network's predictions match the actual results.

21.3.2 Backward Pass (Backpropagation)

During the **backward pass** (also called **backpropagation**):

- The calculated error is sent back through the network.
- The network adjusts its internal parameters, such as **weights** and **biases**, to reduce the error in future predictions.

The **forward pass**, **error calculation**, and **backward pass** are repeated many times with different sets of data until the neural network consistently makes accurate predictions.

21.4 Types of Neural Networks

There are various types of neural networks, including:

- Perceptron Neural Networks
- Feed-Forward Neural Networks
- Deep Feed-Forward Neural Networks
- Modular Neural Networks
- Convolutional Neural Networks (CNNs)
- Recurrent Neural Networks (RNNs)

21.4.1 Perceptron Neural Networks

Perceptron neural networks are the simplest type of artificial neural network.

They consist of only:

- An input layer
- An output layer

21.4.2 Feed-Forward Neural Networks

In a feed-forward neural network, information flows in only one direction, that is, **forward**. Typically, it contains only one hidden layer.

Each neuron in a layer:

- Receives input from neurons in the previous layer.
- Passes its output to neurons in the next layer.

21.4.3 Deep Feed-Forward Neural Networks

A deep feed-forward neural network is similar to a feed-forward neural network but contains **more than one hidden layer**.

21.4.4 Modular Neural Networks

A modular neural network combines **two or more neural networks** to arrive at the final output.

21.4.5 Convolutional Neural Networks (CNNs)

A **Convolutional Neural Network (CNN)** is particularly well suited for analyzing **visual data**.

The term **convolutional** refers to a mathematical operation in which one function is applied to another function, producing a mixture of the two functions.

In CNNs:

- The convolution process takes place through multiple layers.
- Each layer performs a convolution operation on the output from the previous layer.

21.4.6 Recurrent Neural Networks (RNNs)

In **Recurrent Neural Networks (RNNs)**, each neuron in the hidden layers receives an input with a specific delay in time.

This allows the RNN to consider the **context** of the input.

RNNs are useful when previous information needs to be accessed during current iterations.

RNNs have feedback loops that pass information across time steps ($t-1$, t , $t+1$). When an RNN processes data at the current step (t), it receives two inputs:

- The new current data point (x_t).
- The hidden state/memory from the previous time step (h_{t-1}).

For example, they are useful for predicting the **next word in a sentence**, as they consider the context and flow of a conversation.

21.5 Key Points

- Neural networks are the building blocks of AI systems inspired by the structure of the human brain.
- Neural networks consist of an **input layer**, **one or more hidden layers**, and an **output layer**.
- The input layer receives data, hidden layers process the data using activation functions, and the output layer produces the final result.
- Neural networks learn through a training process involving **forward propagation**, **error (loss) calculation**, and **backpropagation**.

- During backpropagation, the network adjusts internal parameters such as **weights** and **biases** to improve future predictions.
- As the number of hidden layers increases, the network becomes deeper, leading to **deep learning**.
- Common types of neural networks include:
 - Perceptron Neural Networks
 - Feed-Forward Neural Networks
 - Deep Feed-Forward Neural Networks
 - Modular Neural Networks
 - Convolutional Neural Networks (CNNs)
 - Recurrent Neural Networks (RNNs)

22 Machine Learning vs. Deep Learning

22.1 Introduction

To understand the difference between **Machine Learning (ML)** and **Deep Learning (DL)**, we can use a simple pizza-ordering example. Deep Learning is a **subset of Machine Learning**, and both belong to the broader field of **Artificial Intelligence (AI)**.

The hierarchy is as follows:

- Artificial Intelligence (AI)
 - Machine Learning (ML)
 - Neural Networks (NN)
 - Deep Learning (DL)

Neural Networks form the **backbone of Deep Learning algorithms**.

22.2 Machine Learning

Machine Learning algorithms leverage **structured, labeled data** to make predictions.

22.2.1 Pizza Ordering Example

Suppose we want to build a model to determine whether we should order pizza for dinner.

The model uses three input factors.

Input	Question	Value
X1	Will it save time by ordering out?	1 (Yes)
X2	Will we lose weight by ordering pizza?	0 (No)
X3	Will it save us money?	1 (Yes, because we have a free pizza coupon.)

For simplicity, the inputs are represented as **binary values (1 and 0)**.

Note: Although binary values are used in this example, neurons in a neural network can represent values ranging from **negative infinity to positive infinity**.

22.2.2 Assigning Weights

Weights determine the importance of each input. Larger weights make an input contribute more significantly to the final output.

Weight	Associated Input	Value	Reason
W1	X1	5	Saving time is considered most important.
W2	X2	3	There is some interest in maintaining fitness.
W3	X3	2	The dinner cost is not a major concern.

The model uses a **threshold value of 5**.

22.2.3 Calculating the Output

The predicted output ($\hat{y}$) is calculated using the weighted inputs and the threshold.

$$\begin{aligned}\hat{y} &= (1 \times 5) + (0 \times 3) + (1 \times 2) - 5 \\ &= 5 + 0 + 2 - 5 \\ &= +2\end{aligned}$$

Since the output is a **positive number (+2)**, the model predicts:

Decision: Order pizza (Pizza Night).

22.3 What Makes Deep Learning Different?

A neural network is considered a **Deep Neural Network (DNN)** when it contains **more than three layers**, including the **input layer** and the **output layer**.

A deep neural network consists of:

- Input layer
- Multiple hidden layers
- Output layer

The presence of multiple hidden layers enables Deep Learning models to learn more complex patterns.

22.4 Human Intervention in Machine Learning

Classical Machine Learning relies more heavily on **human intervention**.

Human experts determine which features are important for distinguishing between different types of data.

For example, suppose we have images of:

- Pizza
- Burger
- Taco

A human expert labels these images before they are processed by the neural network. The expert also determines distinguishing characteristics of each food item.

For example:

- The type of bread may be a distinguishing feature between pizza, burgers, and tacos.

This approach is known as **Supervised Learning** because it incorporates **human intervention (human supervision)** during the learning process.

22.5 Deep Learning and Unsupervised Learning

Deep Learning does not necessarily require a labeled dataset.

It can ingest **unstructured data** (data that does not have predefined labels or organization), such as:

- Text
- Images

Instead of relying on human experts, a Deep Learning model can automatically determine the features that distinguish pizza, burgers, and tacos from one another.

By observing patterns in the data, the model can cluster similar inputs together.

These algorithms discover hidden patterns and natural groupings in data **without requiring human intervention**.

This approach is known as **Unsupervised Learning**.

22.6 Feedforward Neural Networks

Most Deep Neural Networks are **Feedforward Networks**.

In a feedforward network:

- Information moves in one direction.
- Data flows from the input layer to the output layer.

22.7 Backpropagation

Neural networks can also be trained using **Backpropagation** (the process of sending information from the output layer back toward the input layer to calculate errors and improve the model).

Unlike feedforward processing, backpropagation moves in the opposite direction:

- Output layer → Input layer

Backpropagation allows us to:

- Calculate the error associated with each neuron.
- Attribute(give) the error to individual neurons.
- Adjust the network weights.
- Fit (train) the algorithm more accurately.

22.8 Key Differences Between Machine Learning and Deep Learning

Machine Learning	Deep Learning
A subset of Artificial Intelligence.	A subset of Machine Learning.
Uses structured, labeled data.	Can work with unstructured data such as text and images.
Relies more on human intervention to identify and label features.	Can automatically discover features from raw data.
Often uses supervised learning.	Can perform unsupervised learning without labeled data.
Typically uses simpler neural network structures.	Uses deep neural networks with more than three layers.

22.9 Summary

Machine Learning and Deep Learning belong to the same field of study and are both subfields of Artificial Intelligence.

Neural networks provide the foundation for both approaches.

The primary distinctions are:

- Deep Learning uses **neural networks with more than three layers**.
- Machine Learning generally requires **human intervention** to label data and identify important features.
- Deep Learning can automatically learn features from raw, unstructured data with little or no human intervention.

23 Generative AI Models

23.1 Overview

After studying this topic, we will be able to:

- Explain the concept of **Generative AI Models**.
- List the various types of generative AI models.
- Understand how different generative AI model architectures work.
- Differentiate between **unimodal** and **multimodal** models.

Generative AI models are making it possible to mimic human creativity for generating text, art, music, and video.

Generative AI models are a class of AI systems that learn from large datasets by recognizing patterns and trends. By learning from the data, they use **machine learning** and **deep learning** algorithms to create new content.

The design of a generative AI model changes depending on what it is designed to do and how it will be used.

23.2 Types of Generative AI Models

Some common types of generative AI models are:

- Variational Autoencoders (VAEs)
- Generative Adversarial Networks (GANs)
- Autoregressive Models
- Transformers

23.3 Variational Autoencoders (VAEs)

VAEs work by transforming input data through **encoding** and **decoding**.

They have three main parts:

- Encoder Network
- Latent Space
- Decoder Network

The encoder takes the input data and turns it into a simpler form called the **latent space representation** (a compressed representation that captures the most important characteristics of the data).

The latent space holds the key features of the data.

The decoder then uses this latent space representation to create new outputs.

23.3.1 Applications of VAEs

VAEs are used for:

- Image generation
- Anomaly detection
- Other generative tasks

23.3.2 Example

The **Fashion MNIST VAE model** is used to generate and reconstruct images from the **Fashion MNIST** dataset, which includes different clothing items such as:

- Shirts
- Shoes
- Bags

The model creates new, realistic images of clothing items based on the patterns it has learned from the data.

23.4 Generative Adversarial Networks (GANs)

GANs involve two neural networks:

- Generator
- Discriminator

The **generator** creates new data samples.

The **discriminator** checks whether the generated data is real or fake.

Both networks train together.

- The generator tries to create data that looks real.
- The discriminator tries to distinguish between real and fake data.

This process continues until the generator becomes so good at producing realistic data that the discriminator can no longer tell the difference.

23.4.1 Applications of GANs

GANs are used for:

- Image synthesis
- Style transfer
- Data augmentation
- Other image generation tasks

23.4.2 Example

GANs can be used in the fashion industry.

A famous example is [NVIDIA's StyleGAN](#), which can generate high-quality and realistic images of:

- [Faces](#)
- [Animals](#)
- [Landscapes](#)
- [More](#)

23.5 Autoregressive Models

Autoregressive models create data sequentially, considering the context of earlier generated elements.

These models predict the next element in the sequence based on the previous one.

They can generate sequences of data such as:

- Text
- Music

23.5.1 Examples

A music composition tool powered by an autoregressive model can generate new melodies.

[WaveNet](#) is an example of an autoregressive model that generates raw audio waveforms, producing high-quality, natural-sounding speech.

23.6 Transformers

Transformers are generally used in **Natural Language Processing (NLP)** tasks.

They consist of:

- Encoder layers
- Decoder layers

These layers enable the model to:

- Generate text sequences effectively.
- Perform cross-language translations.

23.6.1 Examples

A multilingual chatbot uses transformers to understand and respond to queries in multiple languages.

Large Language Models (LLMs) such as:

- [OpenAI's Generative Pre-trained Transformer \(GPT\) family of models](#)
- [Google Gemini models](#)

are examples of transformer-based models that can generate different creative texts.

23.7 Categories of Generative AI Models

Generative AI models can usually be categorized into:

- [Unimodal Models](#)
- [Multimodal Models](#)

23.7.1 Unimodal Models

Unimodal models process inputs and generate outputs within the same modality.

Modality refers to the type of data being processed.

A unimodal model handles only one type of data.

23.7.1.1 Example

GPT-3 is a unimodal model that:

- Takes text as input.
- Generates text as output.

It can:

- Complete sentences.
- Generate stories.
- Answer questions based on textual prompts.

23.7.2 Multimodal Models

Multimodal models handle inputs from one modality and produce outputs in a different modality.

They process multiple types of data.

23.7.2.1 Example: DALL·E

DALL·E is a multimodal model that generates images from textual descriptions.

For example, given the prompt:

"An elephant playing with a ball"

DALL·E can generate a corresponding image.

23.7.2.2 Example: Meta's ImageBind

Meta's ImageBind is an open-source multimodal AI model that processes diverse data types, including:

- Text
- Audio
- Visual data
- Movement

Its unique ability to combine information across these modalities allows it to create art from mixed inputs.

For example, it can merge the sound of a flowing river with a cityscape visual, showcasing its versatility.

23.8 Future of Generative AI Models

As we move forward, these models will continue pushing the boundaries of creativity, revolutionizing industries, and enhancing human experiences.

23.9 Key Takeaways

- Generative AI models are a class of artificial intelligence systems that use machine learning and deep learning algorithms to create new content by learning patterns from large datasets.
- Common generative AI model architectures include:
 - Variational Autoencoders (VAEs)
 - Generative Adversarial Networks (GANs)
 - Autoregressive Models
 - Transformers
- VAEs use:
 - An encoder to convert input data into a latent space representation.
 - A latent space to capture the key characteristics of the data.
 - A decoder to generate new outputs from the latent space representation.
- GANs consist of:
 - A generator that creates new data samples.
 - A discriminator that verifies whether the generated data is real or fake.
- Autoregressive models generate data sequentially by considering the context of previously generated elements.
- Transformers effectively generate text sequences and perform cross-language translations.
- Generative AI models are categorized as:
 - **Unimodal models**, which process inputs and generate outputs within the same modality.
 - **Multimodal models**, which handle inputs from one modality and produce outputs in a different modality while processing multiple types of data.

24 Foundation Models and Large Language Models (LLMs)

24.1 Introduction

Over the past couple of months, **Large Language Models (LLMs)**, such as **ChatGPT**, have taken the world by storm. Whether it is writing poetry or helping to plan an upcoming vacation, we are seeing significant improvements in AI performance and its potential to drive enterprise value.

24.2 Foundation Models

Large Language Models are part of a broader class of AI models known as **Foundation Models**.

The term **Foundation Models** was first introduced by a team from **Stanford** when they observed that the field of AI was shifting toward a new paradigm.

Previously, AI applications were developed by training many separate AI models, where each model:

- Was trained on task-specific data.
- Performed only one specific task.
- Required its own dedicated training process.

The new paradigm introduces a single foundational capability—a **Foundation Model**—that can support all of these applications and many additional use cases.

The key idea is that the same model can be transferred to many different tasks instead of training a separate model for every application.

24.3 Why Foundation Models Can Perform Multiple Tasks

Foundation Models can transfer their knowledge across multiple tasks because they are trained:

- On an enormous amount of unstructured data.
- Using an **unsupervised learning** approach (learning from data without manually labeled examples).

For language models, this means feeding the model **terabytes of text data**.

For example:

- Beginning of a sentence:
 - *No use crying over spilled...*
- Expected next word:
 - *Milk*

During training, the model learns to predict the next word based on the words that came before it.

This ability to generate the next word is the core capability of Foundation Models.

24.4 Foundation Models and Generative AI

Foundation Models belong to the field of **Generative AI** because they generate new content.

In language models, the generated content is typically the **next word in a sentence**.

Although these models are trained primarily for text generation, they can also be adapted to perform many traditional **Natural Language Processing (NLP)** tasks.

These include:

- Text classification
- [Named Entity Recognition \(NER\)](#) (identifying names of people, organizations, locations, etc. within text)
- Other traditional NLP tasks

24.5 Tuning a Foundation Model

[A Foundation Model can be adapted to perform a specific NLP task by introducing a relatively small amount of labeled data.](#)

This process is called **Tuning**.

[During tuning:](#)

- [A small labeled dataset is added.](#)
- [The model's parameters are updated.](#)
- [The model becomes specialized for a specific NLP task.](#)

[Instead of training a completely new model, we build upon the knowledge already learned during pre-training.](#)

24.6 Prompting and Prompt Engineering

Even when little or no labeled data is available, Foundation Models can still perform effectively.

This is achieved through **Prompting** or **Prompt Engineering**.

Prompt Engineering involves providing carefully designed instructions so the model performs the desired task.

24.6.1 Example: Sentiment Classification

Suppose we provide the model with a sentence and ask:

Does this sentence have a **positive** sentiment or a **negative** sentiment?

The model attempts to complete the text naturally.

Its generated answer becomes the solution to the classification problem.

Possible outputs include:

- Positive
- Negative

These models perform surprisingly well even in domains where only a very small amount of labeled data is available.

24.7 Advantages of Foundation Models

24.7.1 High Performance

The greatest advantage of Foundation Models is their performance.

Because they have been trained on **terabytes of data**, they already possess a broad understanding of language before being adapted to a specific task.

As a result, they can significantly outperform models trained using only small task-specific datasets.

24.7.2 Improved Productivity

Foundation Models also provide substantial productivity gains.

Through **Prompting** or **Tuning**:

- Much less labeled data is required.
- Task-specific models can be created much faster.
- There is no need to build models entirely from scratch.

This is possible because the model leverages everything it learned from its large-scale pre-training.

24.8 Disadvantages of Foundation Models

24.8.1 High Compute Costs

Training Foundation Models is computationally expensive.

Because they are trained on enormous datasets:

- Training requires significant computing resources.
- Smaller organizations often cannot afford to train their own Foundation Models.

Large Foundation Models may contain **billions of parameters**, making inference (running the model to generate predictions) expensive.

Running inference may require:

- Multiple GPUs simultaneously.
- Large amounts of computing infrastructure.

Consequently, Foundation Models can be much more costly than traditional AI approaches.

24.8.2 Trustworthiness Challenges

Although large-scale data is one of the biggest strengths of Foundation Models, it is also one of their biggest weaknesses.

Many language models are trained using text scraped from the Internet.

This creates several challenges:

- The training data may contain bias.
- The data may include hate speech.
- The data may contain toxic or harmful content.
- It is practically impossible for human annotators to review every training example.

In many open-source Foundation Models, the exact datasets used during training are not fully known.

This lack of transparency creates significant trustworthiness concerns.

24.9 IBM's Research and Innovations

IBM recognizes the enormous potential of Foundation Models.

Researchers at **IBM Research** are working to improve:

- Model efficiency.
- Trustworthiness.
- Reliability.

These improvements aim to make Foundation Models more suitable for enterprise and business environments.

24.10 Foundation Models Beyond Language

Foundation Models are not limited to language processing.

They are being applied across many different domains.

24.10.1 Vision Models

Foundation Models can generate images from text descriptions.

Example:

- **DALL·E 2**

DALL·E 2 accepts text as input and generates custom images.

24.10.2 Code Models

Foundation Models are also used for software development.

Example:

- **GitHub Copilot**

Copilot helps developers by automatically completing code while it is being written.

24.10.3 IBM Products Using Foundation Models

IBM is integrating Foundation Models into several enterprise products.

Examples include:

- **Watson Assistant**
- **Watson Discovery**
- **Maximo Visual Inspection**

IBM is also collaborating with **Red Hat** on code-generation models through **Project Wisdom**.

24.11 Foundation Models for Scientific Research

IBM is applying Foundation Models to scientific research in multiple domains.

24.11.1 Chemistry

IBM developed **MoLFormer**, a Foundation Model designed to promote molecule discovery for targeted therapeutics (medical treatments designed to target specific diseases or biological processes).

24.11.2 Climate Change

IBM is developing **Earth Science Foundation Models** using geospatial data (data associated with geographic locations).

These models are intended to improve:

- Climate research.
- Earth science analysis.
- Environmental understanding.

24.12 Summary

Foundation Models represent a major shift in artificial intelligence.

Unlike traditional AI systems that require separate models for individual tasks, Foundation Models provide a single transferable model capable of performing many different applications.

Key characteristics include:

- Training on terabytes of unstructured data.
- Learning through unsupervised training.
- Generating new content, making them part of Generative AI.
- Supporting traditional NLP tasks through Tuning and Prompt Engineering.
- Delivering higher performance and increased productivity.
- Requiring significant computational resources.
- Raising important trustworthiness and transparency challenges.
- Extending beyond language into vision, coding, chemistry, climate science, and many other domains.

IBM continues to advance Foundation Models by improving their efficiency, reliability, and trustworthiness while expanding their use across diverse business and scientific applications.

25 Artificial Intelligence, Machine Learning, Deep Learning, Foundation Models, Large Language Models, and Generative AI

Artificial Intelligence (AI) has become a widely discussed topic, and many related terms such as **Machine Learning (ML)**, **Deep Learning (DL)**, **Foundation Models**, **Large Language Models (LLMs)**, and **Generative AI** are often used together. Understanding how these terms relate to one another helps us build a clear picture of the modern AI landscape.

25.1 Artificial Intelligence (AI)

Artificial Intelligence (AI) refers to the **simulation of human intelligence in machines, enabling them to perform tasks that typically require human thinking.**

AI, in its various forms and paradigms, has existed for decades.

One early example is **Eliza**, a chatbot developed in the **mid-1960s**, which could mimic human-like conversation to a certain extent.

25.2 Machine Learning (ML)

Machine Learning is a **subfield of Artificial Intelligence**.

It focuses on developing algorithms that allow computers to **learn from data and make decisions based on that data**.

Rather than being explicitly programmed to perform a specific task, these algorithms use **statistical techniques** to:

- Learn patterns from data.
- Make predictions.
- Make decisions without human intervention.

Machine Learning is a broad field that encompasses a wide range of techniques and approaches, ranging from traditional statistical methods to complex neural networks.

25.2.1 Core Categories of Machine Learning

The major categories of Machine Learning include:

- **Supervised Learning**
 - Models are trained using **labeled data**.
- **Unsupervised Learning**
 - Models identify patterns in data **without predefined labels**.
- **Reinforcement Learning**
 - Models learn by interacting with an environment and receiving feedback based on their actions.

25.3 Deep Learning (DL)

Deep Learning is a **subset of Machine Learning**.

It specifically focuses on **artificial neural networks with multiple layers**.

The multiple layers in these neural networks are what give Deep Learning its name.

Deep Learning excels at handling:

- Large amounts of unstructured data.
- Images.
- Natural language.
- Complex patterns and intricate structures within data.

While traditional Machine Learning techniques are often efficient for linear separations or simpler patterns, Deep Learning performs exceptionally well when dealing with highly complex data.

25.3.1 Important Note

Not all Machine Learning is Deep Learning.

Traditional Machine Learning methods continue to play a pivotal role in many real-world applications.

Common traditional Machine Learning techniques include:

- Linear Regression
- Decision Trees
- Support Vector Machines (SVMs)
- Clustering Algorithms

These techniques have been widely used for a long time.

In many scenarios, Deep Learning may be unnecessary or may not be the most suitable approach.

25.4 Foundation Models

Foundation Models primarily belong to the field of **Deep Learning**.

The term **Foundation Model** was popularized in **2021** by researchers at the **Stanford Institute**.

Foundation Models are:

- Large-scale neural networks.
- Trained on vast amounts of data.
- Designed to serve as a base (foundation) for many different applications.

Instead of training a completely new model for every individual task, we can:

- Start with a pre-trained Foundation Model.
- Fine-tune it for a specific application.

This approach saves significant time and computational resources.

Foundation Models are trained on diverse datasets that capture a broad range of knowledge.

They can be adapted for many tasks, including:

- Language translation
- Content generation
- Image recognition

Foundation Models represent a shift toward more:

- Generalized AI solutions
- Adaptable AI systems
- Scalable AI systems

25.5 Large Language Models (LLMs)

Large Language Models (LLMs) are a **specific type of Foundation Model**.

They are designed specifically for **processing and generating human-like text**.

25.5.1 Understanding the Term "LLM"

Large

The word **Large** refers to the enormous scale of these models.

LLMs contain a vast number of parameters, often **billions or even more**.

This massive scale contributes to their nuanced understanding and capabilities.

Language

The word **Language** indicates that these models are designed to understand and interact using human languages.

Because they are trained on massive datasets, LLMs can understand:

- Grammar
- Context
- Idioms
- Cultural references

Model

The word **Model** refers to the computational model itself.

At their core, LLMs consist of algorithms and parameters working together to process input and generate output.

25.5.2 Common Language Tasks Performed by LLMs

LLMs can perform a wide range of language-related tasks, including:

- Answering questions
- Language translation
- Creative writing

25.6 Other Types of Foundation Models

Large Language Models are only one category of Foundation Models.

Other important Foundation Models include:

- **Vision Models**
 - Designed to interpret ("see") and generate images.
- **Scientific Models**
 - Used in scientific fields such as biology.
 - Example: Predicting how proteins fold into three-dimensional (3D) shapes.
- **Audio Models**
 - Generate human-sounding speech.
 - Compose music or audio content.

25.7 Generative AI

Generative AI is another important AI-related term.

Generative AI refers to models and algorithms that are **specifically designed to generate new content**.

Foundation Models provide the underlying knowledge, structure, and understanding.

Generative AI harnesses that knowledge to create something new.

It represents the creative expression that emerges from the vast knowledge base learned by Foundation Models.

Generative AI can generate different types of content, including:

- Text
- Images
- Speech
- Music
- Other creative outputs

25.8 Relationship Between the AI Terms

The relationship between these concepts can be summarized as follows:

- **Artificial Intelligence (AI)** is the broadest field.
- **Machine Learning (ML)** is a subfield of AI.
- **Deep Learning (DL)** is a subset of Machine Learning.
- **Foundation Models** are large-scale Deep Learning models.
- **Large Language Models (LLMs)** are a specific type of Foundation Model focused on language.
- **Generative AI** uses Foundation Models to generate new content rather than simply analyzing or predicting information.

26 Natural Language Processing, Speech Technology, and Computer Vision

26.1 Overview

Natural Language Processing (NLP), Speech Technology, and Computer Vision are major fields of Artificial Intelligence (AI) that enable computers to understand human language, process speech, and interpret visual information. Together, these technologies make human-computer interaction more natural, efficient, and intelligent.

26.2 Natural Language Processing (NLP)

Natural Language Processing (NLP) is a subset of Artificial Intelligence (AI) that enables computers to comprehend, interpret, and produce human language.

Humans naturally communicate using language. Although we can use computers and smartphones to send text and voice messages, computers do not innately know how to understand or process natural language. NLP provides the capability to process and understand human communication.

26.2.1 NLP Market Growth

According to a global survey by Fortune Business Insights:

- Current market size: **USD \$29.71 billion**
- Expected market size within the next eight years: **USD \$158.04 billion**
- Compound Annual Growth Rate (CAGR): **23.2%**

26.2.2 How NLP Works

NLP uses Machine Learning (ML) and Deep Learning (DL) algorithms to determine the semantic meaning (meaning of words within context) of language.

It accomplishes this by:

- Deconstructing sentences grammatically.
- Analyzing relationships between words.
- Understanding sentence structure.
- Interpreting the context in which words are used.

For example, based on the context of a conversation, NLP can determine whether the word "**cloud**" refers to:

- Cloud computing, or
- The mass of condensed water vapor floating in the sky.

NLP systems also understand user intent and emotion, including whether a user is asking a question because of:

- Frustration
- Confusion
- Irritation

By understanding the real intent behind human language, NLP systems draw inferences using a broad array of linguistic models and algorithms.

NLP is further divided into several subcategories related to audio and visual tasks.

26.3 Speech Technology

For computers to communicate naturally with humans, they must be able to:

- Convert speech into text.
- Convert text into speech.

These capabilities make communication more natural and easier to process.

26.3.1 Speech-to-Text (STT)

Speech-to-Text (STT) technology converts spoken words into written text using neural networks.

26.3.2 How STT Works

The neural network:

- Analyzes many voice samples along with their corresponding text versions.
- Identifies pronunciation patterns.
- Uses the learned patterns to convert new voice recordings into accurate text.

26.3.3 Applications of STT

Speech-to-Text enables:

- Real-time transcription
- Voice commands
- Dictation
- Transcription services
- Voice search

Examples include:

- **YouTube** uses STT to provide automatic closed captioning.
- **Siri** uses STT to process user commands.
- **Google Assistant** uses STT to process spoken requests.
- **Google Voice Search** uses STT to answer spoken queries.

26.4 Text-to-Speech (TTS)

Text-to-Speech (TTS), also known as **speech synthesis**, converts written text into spoken words.

26.4.1 How TTS Works

Neural networks enable TTS through the following process:

- One neural network learns a person's voice by analyzing numerous voice samples.
- A second neural network generates new speech.
- The generated speech is compared with the original voice.
- If the generated voice does not sufficiently match the original, the second neural network adjusts the generated audio.
- This process repeats until the generated speech sounds natural and closely matches the original voice.

26.5 Integration of STT and TTS

Together, Speech-to-Text and Text-to-Speech enable seamless human-machine interaction through natural language.

26.5.1 Applications

26.5.1.1 Translation Services

Applications such as **Google Translate**:

- Listen to spoken language using STT.
- Translate the spoken language.
- Speak the translated output using TTS.

26.5.1.2 Smart Home Devices

Smart home devices:

- Use STT to understand spoken commands.
- Use TTS to provide spoken feedback.

26.5.1.3 Customer Support

Customer support systems use these technologies in the following sequence:

- STT transcribes customer queries.
- NLP understands and interprets the queries.
- NLP generates responses.
- TTS delivers the responses as speech.

26.5.1.4 Accessibility

For accessibility applications:

- STT converts speech into text in real time.
- The text is displayed on a screen or mobile device.
- NLP analyzes and interprets the content.
- TTS converts the processed information into comprehensible speech.

26.6 Computer Vision

Computer Vision is a field of Artificial Intelligence that enables machines to interpret and understand visual data.

It analyzes images and videos to derive meaningful conclusions and make informed decisions.

Computer Vision bridges the digital and physical worlds.

One common example is **facial recognition**, where a phone analyzes a user's facial image and compares it with a pre-existing image before granting access.

Another example is **self-driving cars**, which use computer vision to understand and interpret their surroundings.

26.7 Neural Networks in Computer Vision

Neural networks play a crucial role in advancing computer vision applications, including:

- Image classification
- Object detection
- Image segmentation

26.7.1 Image Classification

Image classification involves categorizing images into predefined classes.

Examples include:

- Sorting products on e-commerce platforms.
- Detecting disorders in medical images.

26.7.2 Object Detection

Object detection not only identifies objects but also determines their locations within an image.

Common algorithms include:

- YOLO (You Only Look Once)
- Faster R-CNN

Applications include:

- [Surveillance\(Satellite & Drones Technology\)](#)
- [Autonomous vehicles](#)

26.7.3 Image Segmentation

Image segmentation divides images into meaningful segments.

It provides detailed labeling for every pixel, allowing systems to distinguish between different object types or categories within an image.

26.8 Industry Applications of Computer Vision

26.8.1 Retail

Companies including **Amazon**, **Walmart**, and **Alibaba** deploy computer vision for:

- [Inventory management](#)
- [Personalized shopping experiences](#)
- [Operational optimization](#)
- [Enhanced customer satisfaction](#)

26.8.2 Manufacturing

Organizations such as **Toyota**, **Siemens**, and **Bosch** integrate computer vision into production lines for:

- [Quality control](#)
- [Automation](#)
- [Product consistency](#)
- [Improved efficiency](#)

26.8.3 Agriculture

Companies including **John Deere** and **Monsanto** leverage computer vision for precision farming by enabling farmers to:

- [Monitor crop health.](#)
- [Optimize agricultural fields.](#)

26.9 Key Takeaways

- Natural Language Processing (NLP) enables computers to interpret, understand, and produce human language.
- NLP uses Machine Learning and Deep Learning algorithms to determine semantic meaning by deconstructing sentences grammatically, relationally, and structurally.

- Speech-to-Text (STT) converts spoken language into written text using neural networks and supports applications such as voice commands, dictation, transcription, and voice search.
- Text-to-Speech (TTS), also known as speech synthesis, converts written text into natural-sounding spoken language.
- STT and TTS can be integrated with NLP to enable seamless human-machine interaction.
- Computer Vision enables machines to analyze images and videos, derive meaningful insights, and make informed decisions.
- Major computer vision tasks include image classification, object detection, and image segmentation.
- Computer vision is widely used in retail, manufacturing, agriculture, surveillance, autonomous vehicles, healthcare, and facial recognition systems.

27 What is Natural Language Processing (NLP)?

Natural Language Processing (NLP) is the process of enabling computers to understand, interpret, and generate human language. As humans, we naturally listen to words and sentences and form comprehension from them. When we ask a computer to perform the same task, it is called **Natural Language Processing (NLP)**.

27.1 Unstructured Text vs. Structured Data

NLP begins with **unstructured text**.

Unstructured text is the natural way humans communicate through spoken or written language.

Example of unstructured text:

Add eggs and milk to my shopping list.

Humans immediately understand the meaning of this sentence, but to a computer, it is simply unstructured text.

To enable a computer to process this information, we need to convert it into a **structured representation**.

A structured representation might look conceptually like this:

- Shopping List
 - Item: Eggs
 - Item: Milk

This structured format allows a computer to understand and process the information.

27.2 The Role of NLP

The primary job of Natural Language Processing is to translate between **unstructured text** and **structured data**.

There are two directions in this translation:

- **Natural Language Understanding (NLU):** Converts unstructured language into structured data.
- **Natural Language Generation (NLG):** Converts structured data into natural human language.

In this topic, we primarily focus on **Natural Language Understanding (NLU)**, which converts human language into a structured format that computers can understand.

27.3 Applications of NLP

Natural Language Processing is used in many real-world AI applications.

27.3.1 Machine Translation

Machine translation converts text from one language to another.

Accurate translation requires understanding the **context** and **overall structure** of a sentence rather than translating each word.

A well-known example demonstrates why context is important:

Original English sentence:

"The spirit is willing, but the flesh is weak."

After translating from English to Russian and then back to English, the result may become:

"The vodka is good, but the meat is rotten."

This completely changes the intended meaning, demonstrating why NLP must understand context rather than simply translating words individually.

27.3.2 Virtual Assistants and Chatbots

Virtual assistants such as **Siri** and **Alexa** use NLP to understand spoken commands and determine the appropriate action to perform.

Chatbots use NLP in written conversations by interpreting user input and traversing a decision tree to determine the appropriate response or action.

27.3.3 Sentiment Analysis

Sentiment analysis uses NLP to determine the emotional tone expressed in text.

For example, NLP can analyze:

- Email messages

- Product reviews
- Customer feedback

It can determine:

- Whether the sentiment is positive or negative.
- Whether the statement is genuine or sarcastic.

27.3.4 Spam Detection

NLP is widely used for identifying spam emails.

It analyzes the content of an email and looks for indicators such as:

- Overused words
- Poor grammar
- Inappropriate claims of urgency

These characteristics can indicate that an email is spam rather than a legitimate message.

27.4 How NLP Works

NLP is **not a single algorithm**.

Instead, it is a **collection (bag) of tools**, where different techniques are combined depending on the application.

The input to NLP is unstructured text, which may be:

- Written text
- Spoken language converted into text using a **speech-to-text algorithm**

The following tools are commonly applied to transform unstructured text into structured data.

27.5 Tokenization

The first stage of NLP is **Tokenization**.

Tokenization breaks a string of text into smaller units called **tokens**.

Example:

Add eggs and milk to my shopping list.

This sentence contains **eight words**, which become **eight tokens**.

Once tokenized, NLP processes one token at a time.

27.6 Stemming

Stemming reduces a word to its **word stem** by removing prefixes, suffixes, and normalizing tense.

Example:

Word	Stem
Running	Run
Runs	Run
Ran	Run

Stemming helps treat different forms of the same word as a single concept.

However, stemming does not always produce meaningful results.

Example:

- Universal
- University

These words should not both stem to **Universe**.

27.7 Lemmatization

For situations where stemming is insufficient, NLP uses **Lemmatization**.

Lemmatization determines the root form (called the **lemma**) of a word by using its dictionary definition and meaning.

Example:

Word	Stem	Lemma
Better	Bet	Good

This example illustrates why choosing between stemming and lemmatization is important, as lemmatization preserves the actual meaning of words.

27.8 Part-of-Speech (POS) Tagging

Part-of-Speech (POS) Tagging identifies the grammatical role of a token based on its context within a sentence.

Consider the word **make**.

Example 1:

- I'm going to make dinner.

- Here, **make** is a **verb**.

Example 2:

- What **make** is your laptop?
- Here, **make** is a **noun**.

The grammatical role depends entirely on the context in which the word appears.

27.9 Named Entity Recognition (NER)

Named Entity Recognition (NER) identifies whether a token represents a specific entity.

Examples:

Token	Recognized Entity
Arizona	U.S. State
Ralph	Person's Name

NER enables NLP systems to recognize real-world entities such as:

- People
- Locations
- Organizations
- Dates
- Products
- Other named entities

27.10 NLP Processing Pipeline

The overall NLP process can be summarized as follows:

1. Receive unstructured text.
2. Convert spoken language to text (if necessary) using a speech-to-text algorithm.
3. Perform tokenization.
4. Apply stemming or lemmatization.
5. Perform Part-of-Speech (POS) Tagging.
6. Perform Named Entity Recognition (NER).
7. Convert the processed information into structured data.
8. Use the structured data in AI applications.

27.11 Key Takeaways

- Natural Language Processing (NLP) enables computers to understand and generate human language.
- NLP translates between **unstructured text** and **structured data**.
- **Natural Language Understanding (NLU)** converts unstructured text into structured data.

- **Natural Language Generation (NLG)** converts structured data into natural language.
- Major NLP applications include:
 - Machine translation
 - Virtual assistants
 - Chatbots
 - Sentiment analysis
 - Spam detection
- NLP is a collection of techniques rather than a single algorithm.
- The major NLP techniques include:
 - Tokenization
 - Stemming
 - Lemmatization
 - Part-of-Speech (POS) Tagging
 - Named Entity Recognition (NER)
- After converting unstructured language into structured data, the resulting information can be used by a wide variety of AI applications.

28 Self-Driving Cars and the Role of Computer Vision

28.1 Overview

We have been working on self-driving cars for the last few years. This domain has experienced tremendous growth in interest since the early self-driving competitions around **2005**.

Our work focuses on developing our own self-driving vehicle that was able to drive on **public roads in the Waterloo region** last August. This represents an important milestone in advancing autonomous driving technology.

28.2 3D Object Detection in Self-Driving Cars

One of our key research domains is **3D object detection**.

This remains a challenging task for algorithms to perform automatically. The goal is to identify every:

- Vehicle
- Pedestrian(walking human)
- Road sign

present in a driving environment so that the vehicle can make the correct decisions about how it should move and interact safely with surrounding objects.

To achieve this, we extensively work on combining information from multiple sensors, including:

- Laser data
- Vision (camera) data
- Radar data

These different sources of information are fused together to create a complete view of the world surrounding the vehicle.

28.3 Why Computer Vision Is Closely Associated with Self-Driving Cars

When we think of computer vision, we usually think immediately of self-driving cars because driving requires constant visual attention.

For example:

- We cannot safely look at a smartphone while simultaneously paying attention to the road.
- It is often difficult to predict what people on the street are going to do, such as crossing the road on a bicycle, skateboard, or other means of transportation.

Having cameras and sensors that continuously monitor the environment helps detect these situations and can prevent accidents before they potentially occur.

28.4 Limitations of Human Vision

One of the major limitations of human vision is **visual attention** (the ability to focus on only a limited number of objects at one time).

For example, we may be looking at one person, while something else—such as a delicious slice of pizza behind them—is also visible in our field of view.

Although both objects are visible, we can only actively pay attention to one or a limited number of things at a time. We cannot attend to everything in our visual field simultaneously.

28.5 How Computer Vision Overcomes Human Limitations

Unlike humans, cameras and computer vision systems can continuously observe the entire surrounding environment.

This allows them to:

- Monitor the complete visual scene.
- Detect important objects that humans may overlook.
- Help ensure that attention is maintained across the entire environment without requiring us to constantly look in every direction.

This is one of the major advantages of computer vision in self-driving cars.

28.6 The Future of Self-Driving Cars

Self-driving cars are widely considered to be the future of transportation.

There is tremendous interest in autonomous vehicles because of their potential to fundamentally change how society works and operates.

One exciting possibility is being able to:

- Ride in a self-driving car.
- Read during the commute.
- Talk on the phone while traveling to work.

Instead of actively driving through traffic, such as the heavy traffic in **Toronto**, passengers could use their travel time more productively.

28.7 Remaining Challenges

Although self-driving cars represent an exciting step forward, there is still significant work to be done.

Many interesting challenges remain before autonomous vehicles become:

- Robust
- Safe
- Capable of driving themselves **100% of the time** on public roads without human intervention.

29 AI and Cloud Computing, Edge Computing, and IoT

29.1 Overview

In this topic, we learn the basics of the **Internet of Things (IoT)**, **Cloud Computing**, and **Edge Computing**, and understand how they work together with **Artificial Intelligence (AI)** to enable smart, real-time applications.

The convergence of AI, cloud computing, edge computing, and IoT transforms raw data into meaningful solutions, helping build a smarter and more connected future.

29.2 Internet of Things (IoT)

The **Internet of Things (IoT)** refers to a network of physical devices connected to the Internet that collect and share data for processing and analysis.

IoT devices can include:

- Fitness trackers
- Smartwatches that monitor health and activity levels
- Washing machines that can be operated remotely through smartphone apps
- Smart thermostats that learn temperature preferences and automatically adjust heating or cooling
- Sensors
- Cameras

- Other devices that generate data

29.2.1 How IoT Devices Work

IoT devices follow a simple workflow:

- Collect data from their surroundings.
- Send the collected data through the Internet.
- Store the data in the cloud.
- Process and analyze the data to generate useful insights.

29.3 Cloud Computing

Cloud computing allows us to store and use data and services over the Internet instead of keeping everything on our own computer.

The main goal of cloud computing is to allow users to access and use powerful computing resources from remote data centers.

In simple terms, cloud computing means using someone else's computers and software over the Internet instead of relying on our own device and computing resources.

29.3.1 Example of Cloud Computing

Gmail is an example of a cloud service.

- Emails are stored in the cloud.
- We can access our emails from anywhere through the Internet.

29.4 AI in Cloud Computing

Cloud computing provides the storage and computing power needed to analyze data, while AI algorithms provide intelligent capabilities to:

- Find more detailed information in data.
- Automate tasks.
- Make decisions based on data.

Together, AI and cloud computing help us use data to its fullest.

29.4.1 Example: AI in Cloud-Based Email Services

Cloud-based email services use AI to:

- Analyze incoming emails.
- Distinguish between spam and legitimate messages.
- Use patterns and user behavior to automatically move spam emails to the junk folder.
- Continuously improve spam detection as AI learns from user actions, ensuring users receive fewer unwanted emails.

29.5 Edge Computing

Not all data collected by IoT devices needs to be sent to the cloud.

Sometimes, **Edge Computing** is used, where devices process data and make decisions locally before sending data to the cloud. This enables faster decision-making.

Edge computing refers to the practice of processing data closer to the source of generation rather than relying on centralized data centers.

29.5.1 Example: Smart Thermostat

Traditionally, a thermostat might rely solely on the cloud to manage heating and cooling.

With edge computing, the thermostat can:

- Instantly sense the temperature.
- Turn on the heater immediately if it becomes too cold.
- Make decisions locally without waiting for cloud processing.

29.6 Edge AI

Edge AI is a type of AI that runs directly on the device instead of relying on the cloud.

It can be thought of as a mini brain built into everyday devices.

29.6.1 Example of Edge AI

A security camera with facial recognition can:

- Detect unauthorized personnel.
- Trigger alerts immediately without depending on cloud processing.

29.7 How AI, IoT, Cloud Computing, and Edge Computing Work Together

These technologies complement one another to deliver intelligent and real-time solutions.

29.7.1 Example: Fitness Tracker

A fitness tracker is an IoT device that collects and shares data such as:

- Heart rate
- Steps
- Activity level

With edge computing, the fitness tracker can process data locally without needing an Internet connection.

It can:

- Count steps.
- Monitor heart rate.
- Analyze activity in real time.
- Alert us immediately if our heart rate becomes too high during exercise.
- Suggest slowing down when necessary.

By combining edge computing with AI, the fitness tracker can:

- Recognize whether we are running or walking.
- Analyze activity levels.
- Provide instant coaching instructions.
- Recommend changing workout exercises based on heart rate patterns.

29.7.2 Role of Cloud Computing

When the fitness tracker connects to a smartphone through **Bluetooth**, cloud computing becomes available.

The workout data and health metrics are sent to the cloud, where they can:

- Be stored for a long time.
- Be analyzed.
- Be viewed through the smartphone application.

29.7.3 Role of Cloud-Based AI

Cloud-based AI can provide advanced features such as:

- Analyzing sleep patterns.
- Providing recommendations for better sleep.
- Analyzing exercise data.
- Creating personalized workout plans based on fitness goals.

29.8 Real-World Applications

The integration of AI, IoT, cloud computing, and edge computing enables many intelligent applications, including:

- [AI-powered traffic lights](#)
- [Smart public transportation](#)
- [Smart agriculture](#)
- [AI-based smart buildings](#)

These technologies work together to make everyday life more efficient.

29.9 Key Takeaways

- IoT devices are a network of physical devices connected to the Internet that collect and share data for processing and analysis.
- Cloud computing allows us to store and use data and services over the Internet.
- Edge computing processes data closer to where it is generated instead of relying on centralized data centers.
- Edge AI enables AI models to run directly on devices for faster local decision-making.
- The combination of AI, IoT, cloud computing, and edge computing enables smart, intelligent, and real-time applications.
- Real-world applications include AI-powered traffic lights, smart public transportation, smart agriculture, and smart buildings.

30 Module Summary

At this point, we know the following concepts and technologies related to Artificial Intelligence (AI), Machine Learning (ML), Deep Learning, Generative AI, and supporting technologies.

30.1 Machine Learning

Machine learning is a subset of Artificial Intelligence (AI) that uses algorithms to analyze data, make decisions without explicit programming, and enable autonomous problem-solving.

30.1.1 Types of Machine Learning

There are three main types of machine learning:

30.1.1.1 Supervised Learning

Supervised learning is trained on labeled data to classify or predict new data. Its performance becomes more precise as more training data becomes available.

Categories of supervised learning include:

- **Regression:** Estimates continuous values.
- **Neural Networks:** Imitate the structure of the human brain.
- **Classification:** Focuses on discrete values.

30.1.1.2 Unsupervised Learning

Unsupervised learning finds patterns in unlabeled data.

It is useful for:

- Clustering similar data points.
- Detecting anomalies.

30.1.1.3 Reinforcement Learning

Reinforcement learning achieves goals within a defined set of rules and constraints by maximizing rewards.

It is useful for tasks such as:

- Playing chess.
- Navigation.

30.2 Model Training Process

Training a machine learning model involves splitting a dataset into three sets:

Dataset	Purpose
Training Set	Trains the algorithm.
Validation Set	Fine-tunes and validates the model.
Test Set	Evaluates the model's performance.

30.3 Deep Learning

Deep learning uses neural networks with multiple layers to analyze complex data. It enables continuous improvement and learning.

It enhances AI's natural language understanding by grasping context and intent.

Deep learning excels in various tasks, including:

- Image captioning.
- Voice recognition.
- Facial recognition.
- Medical imaging.
- Language translation.
- Driverless cars.

30.4 Neural Networks

A neural network is a computational model consisting of interconnected nodes organized into three layers:

- **Input Layer**
- **One or More Hidden Layers**
- **Output Layer**

30.4.1 Types of Neural Networks

Various types of neural networks include:

- Perceptron.

- Feed-forward Neural Network.
- Deep Feed-forward Neural Network.
- Modular Neural Network.
- Convolutional Neural Network (CNN).
- Recurrent Neural Network (RNN).

30.5 Generative AI Model Architectures

Generative AI uses several model architectures, including:

- Variational Autoencoders (VAEs)
- Generative Adversarial Networks (GANs)
- Autoregressive Models
- Transformers

30.5.1 Variational Autoencoders (VAEs)

A Variational Autoencoder (VAE) consists of:

- **Encoder:** Converts input data into a latent space representation.
- **Latent Space:** Captures the essential characteristics of the data.
- **Decoder:** Generates outputs based on the latent space representation.

30.5.2 Generative Adversarial Networks (GANs)

A Generative Adversarial Network (GAN) consists of:

- **Generator:** Creates new data samples.
- **Discriminator:** Verifies whether the generated data is realistic.

30.5.3 Autoregressive Models

Autoregressive models:

- Create data sequentially.
- Consider the surrounding context while generating data.

30.5.4 Transformers

Transformers:

- Generate text sequences.
- Perform cross-language translations effectively.

30.6 Types of AI Models

AI models can be categorized into:

30.6.1 Unimodal Models

Unimodal models process inputs and generate outputs within the same modality (the same type of data, such as text-to-text or image-to-image).

30.6.2 Multimodal Models

Multimodal models handle inputs from one modality and produce outputs in a different modality (different types of data, such as text-to-image).

30.7 Cognitive Computing

Cognitive computing technology mimics human cognitive processes such as:

- Thinking.
- Reasoning.
- Problem-solving.

30.8 Natural Language Processing (NLP)

Natural Language Processing (NLP) helps computers interpret and produce human language.

It uses machine learning and deep learning algorithms to understand a word's semantic meaning by deconstructing sentences:

- Grammatically.
- Relationally.
- Structurally.

30.9 Speech Technologies

30.9.1 Speech-to-Text (STT)

Speech-to-Text (STT) technology transforms spoken words into written text.

30.9.2 Text-to-Speech (TTS)

Text-to-Speech (TTS) technology converts written text into spoken words.

30.10 Computer Vision

Computer vision enables machines to understand visual data by:

- Analyzing images or videos.
- Drawing meaningful insights.
- Making informed decisions.

30.11 Internet of Things (IoT)

Internet of Things (IoT) devices are a network of physical devices connected to the internet that collect and share data for processing and analysis.

30.12 Cloud Computing

Cloud computing allows us to store and use data and services over the internet.

30.13 Edge Computing

Edge computing refers to the practice of processing data closer to the source where it is generated, rather than relying on a centralized data center.

30.14 Integration of AI, IoT, Cloud Computing, and Edge Computing

The intersection of AI, IoT, cloud computing, and edge computing enables smart and real-time applications.

30.14.1 Real-World Applications

Real-world applications of AI, cloud computing, edge computing, and IoT include:

- AI-powered traffic lights.
- Smart public transportation.
- Smart agriculture.
- Smart buildings.

Module 3: Business and Career Transformation through AI

31 AI Agents

31.1 Introduction

Understanding AI agents is crucial because they enable businesses to use automation and data-driven decision-making, improving efficiency, innovation, and competitiveness across various fields.

The global AI market is expected to increase to **\$1.8 trillion** within the next **seven years**, with a **compound annual growth rate (CAGR)** of **36.6%**.

According to **Gartner**, using AI agents in customer service will save companies around **\$80 billion** by reducing labor costs.

31.2 What Are AI Agents?

AI agents are software programs that:

- Interact with their environment.
- Collect and process data.
- Perform tasks on their own to achieve goals set by humans.
- Make decisions.
- Solve problems.
- Adapt to new information without needing constant human assistance.

31.3 How AI Agents Work

A **self-driving car** is an excellent example of how AI agents operate. The process consists of several phases.

31.3.1 Perception

AI agents use sensors such as:

- Cameras
- Radar

These sensors gather data about the environment by identifying:

- Vehicles
- Pedestrians
- Road signs

31.3.2 Understanding

Algorithms process the collected data to identify:

- Objects
- Their speed
- Their movement

31.3.3 Decision-Making

The AI agent uses a knowledge base and logic to determine actions such as:

- When to accelerate
- When to brake

- When to turn

31.3.4 Action

Actuators execute the decisions by controlling:

- Steering
- Brakes

This enables the vehicle to navigate safely.

31.3.5 Learning

The AI agent continually improves over time using **Machine Learning**.

It analyzes previous driving experiences to enhance future performance. For a self-driving car, this means continuously learning from past driving experiences.

31.4 Characteristics of AI Agents

AI agents possess several important characteristics:

- [Social ability](#)
- [Autonomy](#)
- [Reactiveness](#)
- [Proactiveness](#)

31.4.1 Social Ability

Social ability refers to an agent's capacity to communicate and collaborate with other agents and entities in its environment.

31.4.2 Autonomy

Autonomy is the agent's ability to operate independently and make decisions without external assistance.

31.4.3 Reactiveness

Reactiveness allows agents to respond immediately to changes in their environment.

31.4.4 Proactiveness

Proactiveness enables agents to take initiative and make decisions to achieve their objectives.

31.5 Applications of AI Agent Characteristics

31.5.1 Healthcare

The **social ability** feature is extensively used in healthcare.

Examples include healthcare chatbots that:

- Give medical advice.
- Check symptoms.
- Schedule appointments.
- Send reminders.

31.5.2 Self-Driving Cars

Self-driving cars use the **autonomy** feature to:

- Navigate roads.
- Make driving decisions.
- Operate without human assistance.

31.5.3 Smart Thermostats

Reactiveness is employed in smart thermostats, which adjust home temperatures based on real-time data.

31.5.4 Predictive Maintenance

In industries, predictive maintenance systems use the **reactiveness** feature to:

- Analyze equipment data.
- Schedule maintenance before breakdowns occur.

31.6 Multi-Agent Systems

Managing and coordinating many independent systems has become increasingly important in today's rapidly changing technology landscape.

Examples include:

- Managing resources spread across large geographical areas.
- Operating smart transportation systems.
- Creating complex simulations.

This is where **Multi-Agent Systems (MAS)** become valuable.

In multi-agent systems, multiple agents work together to achieve both individual and collective goals.

Multi-agent systems enable:

- Distributed problem-solving.
- Cooperative decision-making.
- Emergent behaviors from interactions among agents (complex behaviors that naturally arise when multiple agents interact, even though no single agent was explicitly programmed to produce them).

31.7 Applications of Multi-Agent Systems

31.7.1 Online Marketplaces

Buyer and seller agents:

- Negotiate prices.
- Manage transactions for users.
- Facilitate efficient and automated commerce.

31.7.2 Robotic Coordination

Teams of robots collaborate on complex tasks such as:

- Warehouse logistics.
- Search and rescue missions.

31.7.3 Traffic Management

Agents representing autonomous vehicles communicate to:

- Optimize traffic flow.
- Reduce congestion.

31.8 AI Agents in Real-World Products and Services

31.8.1 Google

Google integrates AI agents into several products.

31.8.1.1 YouTube

AI agents:

- Recommend content.
- Moderate content.
- Analyze user behavior to suggest videos.
- Detect and remove inappropriate content.
- Enforce community guidelines.
- Manage copyright issues.

31.8.1.2 Gmail

AI agents power features such as:

- Smart Compose
- Smart Reply

They also:

- Organize emails.
- Filter spam.
- Identify important messages.

31.8.1.3 Google Maps

Google Maps uses AI agents for real-time navigation by gathering real-time data from multiple sources to:

- Recommend accurate routes.
- Estimate travel times.
- Suggest alternative routes to avoid traffic congestion.

31.8.2 Amazon

Amazon employs AI agents across multiple products and services.

31.8.2.1 Alexa

Amazon Alexa uses AI for:

- Voice recognition.
- Understanding natural language.
- Controlling smart home devices.

31.8.2.2 Amazon.com

The Amazon.com e-commerce platform uses AI to:

- Recommend products.
- Improve search results.
- Manage inventory.

31.8.2.3 Amazon Web Services (AWS)

Amazon Web Services (AWS) provides AI services for developers, including:

- Machine Learning models.
- Natural Language Processing (NLP).
- Computer Vision tools.

31.9 Future of AI Agents

AI agents are becoming more than just software tools.

They are evolving into partners that assist with:

- Strategic decision-making.
- Customer engagement.

Their influence and impact are expected to continue growing across industries.

31.10 Key Takeaways

- AI agents are software programs that interact with their surroundings, process data, and autonomously perform tasks to achieve human-defined goals.
- AI agents can make decisions, solve problems, and adapt to new information without constant human intervention.
- The working process of AI agents consists of perception, understanding, decision-making, action, and learning.
- The four key characteristics of AI agents are social ability, autonomy, reactivity, and proactivity.
- Social ability enables communication and collaboration with other agents and entities.
- Autonomy allows agents to operate independently and make decisions without external assistance.
- Reactivity enables agents to respond immediately to environmental changes.
- Proactivity allows agents to take initiative and make decisions to accomplish objectives.
- Multi-agent systems support distributed problem-solving, cooperative decision-making, and emergent behaviors.
- Multi-agent systems are widely used in online marketplaces, robotic coordination, and traffic management.
- Google uses AI agents in **YouTube**, **Gmail**, and **Google Maps** for recommendations, moderation, navigation, email organization, and spam filtering.
- Amazon uses AI agents in **Alexa**, **Amazon.com**, and **Amazon Web Services (AWS)** for voice recognition, product recommendations, inventory management, Machine Learning, Natural Language Processing, and Computer Vision services.
- AI agents are increasingly becoming strategic partners that enhance decision-making, customer engagement, and operational efficiency.

32 AI Agents and the Shift from Compound AI Systems to Agentic Systems

32.1 Introduction

2024 is expected to be the year of **AI agents**. To understand AI agents, we first need to understand the major shifts taking place in the field of **Generative AI**.

One of the most significant shifts is the movement from **monolithic models** to **compound AI systems**.

32.2 From Monolithic Models to Compound AI Systems

Models on their own are limited by the data they have been trained on. This affects:

- What they know about the world.
- What kinds of tasks they can solve.
- Their ability to adapt to new requirements.

Although we can tune a model, doing so requires a significant investment in data and computational resources.

32.2.1 Example: Checking Remaining Vacation Days

Suppose we want to plan a vacation for the summer and ask:

"How many vacation days do we have left?"

If we simply send this query to a language model, the answer will likely be incorrect because:

- The model does not know who we are.
- It does not have access to our personal vacation records.
- It cannot access sensitive company data.

Language models are excellent for tasks such as:

- Summarizing documents.
- Creating first drafts of emails.
- Writing reports.
- Assisting with general content generation.

However, their full potential is unlocked when we build systems around them and integrate them into existing workflows.

32.3 How a Compound AI System Solves the Problem

Instead of relying only on the language model, we can design a system that gives the model access to the vacation database.

The process works as follows:

- We submit the vacation query.
- The language model generates an appropriate search query.
- The search query is sent to the vacation database.
- The database retrieves the required information.
- The retrieved information is returned to the language model.
- The language model generates a natural language response.

Example response:

"Maya, you have ten days left in your vacation database."

Because the system has access to the correct data source, the answer is accurate.

This is an example of a **compound AI system**.

Compound AI systems recognize that many problems are solved more effectively by combining **system design principles** with language models rather than relying on a model alone.

32.4 Characteristics of Compound AI Systems

Compound AI systems are inherently **modular**.

A system may include:

- A language model.
- Tuned models.
- Large Language Models (LLMs).
- Image generation models.
- Programmatic components.
- Output verifiers.
- Programs that break complex queries into smaller tasks.
- Database search components.
- External tools.

Instead of trying to make one model solve every problem, we can break the problem into parts and choose the appropriate component for each task.

This approach is generally easier and faster than continuously tuning a model.

32.5 Retrieval-Augmented Generation (RAG)

One of the most common examples of a compound AI system is **Retrieval-Augmented Generation (RAG)**.

A RAG system follows a predefined path.

For example, if the system is designed to search a vacation policy database and we ask:

"What is the weather today?"

the system will fail because it is programmed to always search the vacation database, which contains no weather information.

This predefined path that determines how a query is processed is known as the **control logic** of the program (the set of rules that determines how the system handles different requests).

In traditional compound AI systems, this control logic is written by humans.

32.6 From Compound AI Systems to AI Agents

Another way to control a compound AI system is to place a **Large Language Model (LLM)** in charge of the control logic.

This has become possible because modern LLMs have significantly improved reasoning capabilities.

Instead of simply producing an immediate answer, LLMs can:

- Analyze complex problems.
- Break problems into smaller steps.
- Create plans.
- Decide how to solve each step.

This approach is known as the **agentic approach**.

32.7 Fast Thinking vs. Slow Thinking

We can think of AI systems as existing on a spectrum.

32.7.1 Fast Thinking

The system:

- Acts exactly as programmed.
- Follows predefined instructions.
- Does not deviate from the specified workflow.

32.7.2 Slow Thinking

The system:

- Creates a plan.
- Solves one part of the problem at a time.
- Evaluates its progress.
- Adjusts the plan when necessary.
- Continues iterating until it reaches a solution.

Instead of returning the first answer that comes to mind, the model reasons through the problem.

This significantly increases the likelihood of producing a correct answer.

32.8 Components of LLM Agents

32.8.1 Reasoning

Reasoning places the language model at the center of problem solving.

The model is prompted to:

- Create a plan.
- Think through each step.
- Continuously reason about the problem throughout execution.

32.8.2 Acting Through Tools

Agents can perform actions using **tools**.

Tools are external program components that the language model decides when and how to use.

Examples of tools include:

- Web search.
- Database search.
- Calculators.
- Programs that manipulate databases.
- Another language model for specialized tasks such as translation.
- APIs (Application Programming Interfaces that allow software systems to communicate with each other).
- Any external program that extends the capabilities of the language model.

The model determines:

- Which tool to use.
- When to use it.
- How to use it.

32.8.3 Memory

Agents can also access **memory**.

Memory may include:

- Internal reasoning logs (records of the model's reasoning process that can be revisited later).
- Previous conversations with users.
- Historical interactions.
- Previously retrieved information.

Memory enables the system to provide more personalized and context-aware responses.

32.9 ReAct Agents

One of the most popular approaches for building AI agents is **ReAct**.

The name **ReAct** combines:

- **Reasoning**
- **Acting**

32.9.1 ReAct Workflow

The workflow operates as follows:

- A user submits a query.
- The query is sent to the LLM.
- The LLM is instructed not to provide the first answer immediately.
- The LLM creates a plan.
- The LLM decides whether external tools are needed.
- The selected tools are executed.
- The tool outputs are returned to the LLM.
- The LLM observes the results.
- If necessary, it revises its plan.
- The process repeats until a satisfactory final answer is produced.

This creates a continuous cycle of:

- Reason
- Act
- Observe
- Iterate

32.10 Example: Planning a Vacation with an AI Agent

Suppose we ask:

"I'm planning to go to Florida next month. I'll spend a lot of time outdoors, I'm prone to sunburn, and I want to know how many two-ounce sunscreen bottles I should bring."

This is a much more complex problem.

The agent may solve it by completing several tasks.

32.10.1 Determine Available Vacation Days

The agent may retrieve this information from memory because we previously asked about our remaining vacation days.

32.10.2 Estimate Time in the Sun

The agent may:

- Check the weather forecast for Florida next month.
- Estimate the expected daily sunlight hours.

32.10.3 Retrieve Health Recommendations

The agent may visit a public health website to determine:

- The recommended sunscreen usage per hour of sun exposure.

32.10.4 Perform Calculations

Finally, the agent may:

- Calculate the total sunscreen required.
- Determine how many two-ounce bottles are needed.

This demonstrates that there are multiple possible paths to solving a single complex problem.

Because of this flexibility, agentic systems are highly modular and capable of handling much more sophisticated tasks.

32.11 Compound AI Systems vs. Agentic Systems

Compound AI systems are expected to remain an important part of AI development.

However, they are becoming increasingly **agentic**.

We can think of this as a **sliding scale of LLM autonomy** (the degree of independence given to the language model when making decisions).

The system designer chooses how much autonomy to provide depending on the problem being solved.

32.11.1 When a Programmatic Approach Is Better

A traditional programmatic approach is often more efficient when:

- The problem is narrow.
- The workflow is well defined.
- Every query should follow the same execution path.

Example:

A vacation database system where every request should always search the vacation records.

Using an agentic approach here may introduce unnecessary reasoning, looping, and iteration.

32.11.2When an Agentic Approach Is Better

An agentic approach becomes valuable when:

- The system must solve highly complex tasks.
- Users may ask a wide variety of questions.
- Many possible execution paths exist.

Example:

A system that independently solves GitHub issues while handling many different types of requests.

In these situations, manually programming every possible path would require significant effort, making an agentic system more practical.

32.12 Current State of Agentic Systems

We are still in the early stages of agentic systems.

Rapid progress is being made by combining:

- Strong system design.
- Agentic behavior.
- Advanced reasoning capabilities.

In most real-world applications, a **human remains in the loop** (a person supervises or reviews the system's decisions before they are finalized) while the accuracy and reliability of AI agents continue to improve.

32.13 Comparison of Monolithic, Compound, and Agentic AI Systems

32.13.1Architecture Comparison

Feature Aspect /	Monolithic AI Systems	Compound AI Systems	Agentic AI Systems
Core Definition	A single, standalone AI model (such as a base Large Language Model) that relies entirely on its pre-trained internal parameters.	A system that builds modular software components (such as databases, code verifiers, and search APIs) around one or more AI models.	A specialized compound AI system where a reasoning-capable Large Language Model acts as the autonomous central decision-maker.
Control Logic (Workflow Control)	None. The model simply processes incoming text and generates an output	Human-Written Code (Programmatic Logic). Human developers hardcode	Model-Driven (Agentic Logic). The LLM dynamically decides how to solve the problem,

	based on its training data.	the exact step-by-step path the data must follow.	creates a plan, and selects steps in real time.
External Tool & Data Access	Cannot access external tools, live databases, APIs (Application Programming Interfaces), or real-time internet data.	Can access external databases and tools, but only through fixed, pre-programmed execution pipelines.	Can dynamically call, evaluate, and combine external tools (such as calculators, web search, database queries, and custom scripts) as needed.
Problem-Solving Approach	"Fast Thinking." Generates an immediate single response based on statistical patterns in its training dataset.	"Fixed Pipeline." Executes a strict sequence of programmed instructions to fetch data and generate a response.	"Slow Thinking." Prompts the model to reason, plan, execute tools, observe intermediate results, and iterate until reaching a solution.
Flexibility & Adaptability	Low. Adapting the model requires expensive fine-tuning or retraining on new datasets.	Moderate. Components are modular, but adding new execution paths requires manual coding by software developers.	High. The system dynamically adapts its execution path to unexpected queries without requiring code changes.
Handling Complex Queries	Fails or hallucinates (generates false information) if the query requires personal, real-time, or multi-step reasoning.	Solves narrow, specific queries reliably, but fails if a user asks something outside its pre-coded path.	Excels at complex, open-ended multi-step problems by breaking them into manageable sub-tasks.
Primary Use Case	Simple text summarization, drafting emails, standard text generation, and baseline language tasks.	Narrow, well-defined workflows such as standard RAG (Retrieval-Augmented Generation, a method that retrieves specific database records to answer questions) or vacation day lookups.	Open-ended, complex workflows such as independently resolving GitHub software bugs, multi-source research, and detailed trip planning.
Practical Example	Asking a base model, "How many vacation days do we have left?" (Yields an incorrect guess because it has no access to private records).	Asking a vacation bot, "How many vacation days do we have left?" (The system automatically runs a hardcoded database search for our user ID and formats the answer).	Asking an agent, "Calculate how many 2-ounce sunscreen bottles we need for our Florida vacation next month." (The agent checks our remaining vacation days, looks up Florida weather forecasts, queries health guidelines for sunscreen

			application, and performs the math).
--	--	--	--------------------------------------

32.13.2 Summary of Core Differences

To keep the distinction clear for our course notes:

- **Monolithic Systems** rely on a single model working in isolation with no outside tools or system design.
- **Compound Systems** wrap models inside human-coded pipelines to fetch data and verify outputs, following rigid, predetermined steps.
- **Agentic Systems** hand over control logic to a Large Language Model, allowing the model to act as an autonomous brain that plans, calls external tools, observes outcomes, and iterates to solve complex problems.

33 Robotics and Automation

33.1 Introduction

In this topic, we learn about robotics and automation, how robots work, how Artificial Intelligence (AI) technologies are integrated into robotics, how robotics enables automation across various industries, and the concept of Robotic Process Automation (RPA) and its applications.

After studying this topic, we will be able to:

- Define the term **robotics** and understand how robots work.
- Discuss how **Artificial Intelligence (AI)** technologies are integrated into robotics.
- Examine how robotics enables automation across various industries.
- Explore the concept of **Robotic Process Automation (RPA)** and its applications.

33.2 Robotics and Automation in Industry

Across the industrial world, companies are investing heavily in robotics and automation. With the advancement of technology, robots have become more capable of performing complex tasks that were previously done by humans.

This shift is driven by:

- Higher productivity.
- Improved quality.
- The ability to operate continuously.

According to a recent [global industrial robotic survey conducted by McKenzie](#), many industries plan to spend **25% of their budgets on automated systems over the next four years**.

33.3 Robotics

Robotics involves designing, constructing, and operating robots, which are machines that can perform tasks by themselves or with some help.

Robots are complex systems made up of three key components:

- Sensors
- Actuators
- Controllers

33.3.1 Sensors

Sensors act as the robot's **eyes and ears**, gathering information about the surrounding environment.

Examples include:

- A **camera sensor** helps a robot navigate by capturing images and detecting obstacles.
- **Temperature sensors** ensure the robot operates within safe limits.

33.3.2 Actuators

Actuators are the robot's **muscles**, enabling movement and interaction with the environment.

Examples include:

- Electric motors that move robotic arms and wheels.
- Hydraulic actuators.
- Pneumatic actuators, which control various parts of the robot.

33.3.3 Controllers

Controllers act as the robot's **brain** and run software that determines the robot's actions.

For example:

- In a robotic arm, the controller interprets sensor data and directs the motors to move the arm accurately.

33.4 AI Technologies Used in Robotics

Robots use their sensors, actuators, and controllers together with various AI technologies to enable automation.

These AI technologies include:

- Artificial Intelligence (AI)
- Machine Learning (ML)

- Natural Language Processing (NLP)

33.5 Examples of AI-Powered Robots

33.5.1 Robot Vacuum Cleaner

A robot vacuum cleaner uses AI to clean a home efficiently.

It includes:

- Sensors to detect dirt, walls, and furniture.
- An AI controller that maps the home and chooses the best cleaning path.
- Actuators that adjust the vacuum's wheels and brushes to clean surfaces such as carpets and hardwood floors.

33.5.2 Smart Lawn Mower

A smart lawn mower uses Machine Learning to mow lawns efficiently.

It includes:

- Sensors to detect grass height, obstacles, and lawn boundaries.
- A controller that learns the lawn's layout and grass growth patterns to optimize mowing routes and schedules.
- Actuators that adjust the mower's wheels and blades to ensure an even cut.

33.5.3 Voice-Activated Smart Speaker Robot

A voice-activated smart speaker robot uses **Natural Language Processing (NLP)** to understand and respond to spoken commands.

It includes:

- Microphones as sensors to capture voice.
- An NLP controller to process and understand language.
- Actuators that perform actions such as:
 - Playing music.
 - Providing updates.
 - Controlling smart home devices.

These devices act as **self-operating robots**, meaning they do not require human intervention because they can figure things out on their own.

33.6 Self-Operating Robots

Self-operating robots are used in environments that are difficult or dangerous for humans.

Examples include:

33.6.1 Space Exploration

Space rovers, such as those on **Mars**, explore other planets, collect valuable information, and send it back to Earth independently.

33.6.2 Deep Sea Exploration

Underwater drones explore deep parts of the ocean, collect samples, and perform research in places that are difficult for people to reach.

33.7 Collaborative Robots (Cobots)

Robots that work directly alongside humans are called **Cobots (Collaborative Robots)**.

Unlike traditional robots, cobots are designed to interact directly with people, making automation more accessible and versatile.

Cobots are equipped with:

- Advanced sensors.
- AI technologies.

These capabilities allow cobots to communicate and coordinate their actions to complete tasks requiring teamwork.

Examples include:

- **Automobile manufacturing:** One cobot can hold a part in place while another welds it, ensuring precision and speed.
- **Logistics:** Cobots work together to sort and move packages efficiently.

33.8 Applications of Robotics in Different Industries

33.8.1 Healthcare

Robotics is used in healthcare to:

- Perform precise surgical operations using surgical robots.
- Reduce recovery times.
- Improve patient outcomes.
- Help patients recover mobility and strength using rehabilitation robots.
- Offer personalized therapy.

33.8.2 Agriculture

Agricultural robots are used to:

- Plant crops.
- Harvest crops.

- Monitor crops with precision.
- Boost yields.
- Reduce waste.
- Optimize irrigation.
- Apply fertilizers accurately.
- Monitor crop health in real time.

33.8.3 Retail

In retail, robots are used for:

- Self-checkout kiosks.
- Inventory management.
- Customer assistance.

These applications:

- Streamline operations.
- Improve customer service.
- Reduce wait times.
- Ensure stock availability.

33.9 Robotic Process Automation (RPA)

33.9.1 Definition

Robotic Process Automation (RPA) is a type of computer software that helps create, use, and control **virtual robots**.

These virtual robots behave like humans when interacting with digital systems and software.

According to a **PEGA system survey** conducted among **500 decision-makers** employing various RPA solutions:

- **63%** viewed RPA as a major component of overall digital transformation.
- **23%** said it is the most important component.
- **14%** said it is a minor component.

33.9.2 Applications of RPA

RPA automates repetitive and rule-based tasks across various industries.

33.9.2.1 Finance and Accounting

RPA can automate:

- Invoice processing by capturing, validating, and posting invoices.
- Payroll calculations, deductions, and payments.
- Financial report generation by extracting data from multiple sources.

33.9.2.2 Human Resources

RPA streamlines:

- Employee onboarding by collecting documents, entering data, and setting up profiles.
- Leave management by handling leave requests and updating records.
- Recruitment by screening resumes, scheduling interviews, and sending communications.

33.10 Key Takeaways

- Robotics involves designing, constructing, and operating robots to perform tasks with or without human assistance.
- Robots use **sensors** to observe their environment, **controllers (processors for decision-making)**, and **actuators** for movement.
- Robotics integrates **Artificial Intelligence (AI)**, **Machine Learning (ML)**, and **Natural Language Processing (NLP)** to enable automation.
- **Cobots (Collaborative Robots)** work alongside humans using advanced sensors and AI technologies, communicating and coordinating tasks that require teamwork.
- **Robotic Process Automation (RPA)** automates repetitive, rule-based digital tasks by creating, using, and controlling virtual robots.

34 Transforming Businesses Through AI

34.1 AI as a Business Transformation Tool

Artificial Intelligence (AI) is transforming business operations through:

- Smart workflow automation
- Powerful data analysis
- Efficient decision-making

This transformation helps companies achieve new levels of success and remain competitive in fast-paced markets.

According to **Accenture**, AI in business operations is predicted to:

- Double workforce efficiency
- Boost profitability by **38%** over the next **10 years**

34.2 AI for Workflow Automation

Organizations handle a wide range of tasks daily. AI automates repetitive tasks such as:

- Data entry
- Scheduling
- Report generation

This allows employees to focus more on creative and strategic work.

Depending on which tasks we are looking to automate, a wide range of AI tools is available.

34.3 AI in Customer Service

AI chatbots understand:

- Natural language
- Customer sentiment

This enables them to offer personalized support, helping businesses:

- Resolve issues faster
- Improve customer satisfaction
- Reduce customer service costs compared to human agents

34.3.1 Real-World Example: AirHelp

AirHelp, a global service, uses AI to assist airline passengers with:

- Flight status updates
- Speedy customer support
- Cost-efficient customer service

This boosts customer satisfaction and strengthens the company's brand reputation.

34.4 AI in Recruitment

AI recruitment tools analyze:

- Resumes
- Video interviews

This helps businesses:

- Find the best candidates
- Fill positions faster
- Reduce productivity gaps
- Ensure quality hires

34.4.1 Real-World Example: Hilton

Hilton's AI chatbot accelerates hiring by:

- Answering candidates' questions
- Scheduling interviews

This saves Human Resources (HR) teams time and can potentially lower hiring costs.

34.5 AI in Accounting

AI platforms automate:

- Payroll
- Auditing

Benefits include:

- Reduced risk of human error in manual processes
- Enhanced accuracy

34.5.1 Real-World Example: EY

EY uses AI to:

- Evaluate fraud risks
- Detect hidden patterns in financial data
- Identify financial misstatements early

This helps protect clients and investors.

34.6 AI for Data Analysis and Decision Making

AI enables businesses to make informed decisions by:

- Swiftly analyzing large volumes of data
- Identifying patterns and trends
- Predicting outcomes based on historical data

Benefits include:

- Fewer errors
- Better predictions
- Improved decision-making

AI-driven tools also provide:

- Real-time analytics
- Interactive dashboards

These help organizations:

- Monitor performance
- Identify opportunities for improvement

34.7 AI in Marketing and Sales

AI tools analyze:

- Purchasing patterns
- Customer purchase history
- Browsing history

Based on this analysis, AI suggests relevant products.

Benefits include:

- Increased sales
- Higher conversion rates
- Improved customer satisfaction

34.7.1 Real-World Example: Amazon

Amazon uses AI for:

- Personalized product recommendations
- Inventory management

This helps:

- Reduce the risk of stock-outs
- Ensure the right products are available
- Increase sales
- Improve customer satisfaction

34.8 AI in Manufacturing

AI supports predictive maintenance by using:

- Sensors
- Data analysis

These technologies monitor equipment to prevent costly breakdowns.

Benefits include:

- Proactive maintenance scheduling
- Prevention of catastrophic equipment failures

34.8.1 Real-World Example: BMW

BMW uses AI to inspect parts on the production line.

This enables the company to:

- Quickly identify quality issues
- Produce higher-quality vehicles
- Improve manufacturing efficiency

34.9 AI in Healthcare and Research

AI accelerates:

- Drug discovery
- Drug development

It does this by analyzing medical data, enabling businesses to develop successful drugs more quickly.

34.9.1 Real-World Example: IBM

IBM's healthcare technology uses AI to:

- Analyze large volumes of medical data
- Support genomic analysis (the study and interpretation of an organism's complete genetic material)
- Assist in disease diagnosis

This attracts new clients and strengthens IBM's reputation within the healthcare industry.

34.10 AI for Creativity and Innovation

AI enhances creativity by:

- Generating innovative ideas
- Overcoming creative blocks

This helps businesses better connect with their target audience, resulting in:

- Higher engagement
- Increased sales

34.10.1 Real-World Example: Canva

Canva's Magic Design feature suggests:

- Layouts
- Color schemes
- Fonts

These suggestions are based on:

- Content type
- User preferences

This improves the user experience and keeps users engaged with Canva.

34.11 Benefits of AI in Business

AI provides numerous benefits across industries, including:

- Improved operational efficiency
- Enhanced customer experience
- Significant cost savings
- Better decision-making
- Higher workforce productivity
- Greater business competitiveness
- Increased customer engagement
- Higher sales and conversions
- Reduced human errors
- Improved product and service quality

34.12 Key Takeaways

- AI automates repetitive business tasks such as data entry, scheduling, and report generation, allowing employees to focus more on creative and strategic work.
- AI enables businesses to analyze large volumes of data quickly by identifying patterns and trends, resulting in fewer errors and better predictions.
- AI enhances decision-making through real-time analytics, dashboards, and predictive insights.
- AI is widely used across customer service, recruitment, accounting, marketing and sales, manufacturing, healthcare, and creative design.
- Real-world organizations including **Accenture**, **AirHelp**, **Hilton**, **EY**, **Amazon**, **BMW**, **IBM**, and **Canva** demonstrate how AI improves efficiency, customer satisfaction, quality, and profitability.
- AI generates innovative ideas and helps overcome creative blocks, enabling businesses to better engage their target audience and increase sales.

35 Rise of Generative AI for Business

35.1 Learning Objectives

After studying this topic, we will be able to:

- Recognize the transformative impact of Generative AI on businesses.
- Identify the practical applications of Generative AI in various business contexts.

35.2 The Rise of Generative AI in Business

Businesses are achieving remarkable milestones and transforming the way technology is applied by harnessing Generative AI.

Arthur C. Clarke famously said:

"Any sufficiently advanced technology is indistinguishable from magic."

For many people, the first experience with Generative AI evokes a sense of magic. For the first time in history, we have technology that can:

- Speak our languages.
- Understand our requests.
- Produce entirely novel (completely new and original) outputs.

Generative AI is capable of:

- Writing poetry.
- Creating otherworldly images (highly imaginative or surreal images).
- Writing code.
- Producing original jokes.
- Composing original music.

The ability of AI to create often inspires a sense of wonder.

35.3 Generative AI Beyond the Digital World

Artificial Intelligence is not limited to the digital world. It also has significant potential in the physical world.

When applied properly, AI can accelerate the pace of discovery and innovation in areas such as:

- [Discovering new materials.](#)
- [Medicine.](#)
- [Energy.](#)
- [Climate.](#)
- [Addressing many of the pressing challenges faced as a species.](#)

Ultimately, our success depends on how we approach and use AI.

35.4 Generative AI Adoption Across Industries

A survey by **JP Morgan** indicates that investors predict the highest Generative AI technology adoption in the following industries:

Industry	Predicted Adoption
Marketing	28%
Legal Services	21%
Insurance	21%
Media	20%
Data Analytics	18%
Consumer Technology	13%

These predictions indicate significant transformations ahead for these industries.

35.5 How Generative AI Is Transforming Business

In the modern business landscape, companies constantly strive to become:

- More efficient.
- More innovative.
- More customer-centric.

Generative AI is making significant contributions in several key business areas.

35.6 Content Generation

Businesses constantly require fresh content, including:

- Marketing materials.
- Product descriptions.
- Social media posts.

Producing high-quality content is often time-consuming and resource-intensive.

Generative AI platforms can:

- Analyze target audience data.
- Create unique and engaging content.
- Help organizations overcome creative blocks.
- Maximize audience reach and impact.

35.6.1 Real-World Example: Persado and Vanguard

Persado, a Generative AI platform, helped **Vanguard** increase its sales conversion rate by **15%**.

Vanguard used Persado to:

- Find the most effective phrases that captivate customers.
- Craft messages with the right emotion and action.
- Deliver impactful content.
- Seamlessly integrate data-driven insights into its marketing strategy.

35.7 Data Analysis and Business Insights

Businesses collect enormous amounts of data, but transforming this information into actionable insights can be overwhelming.

Generative AI helps by:

- Analyzing complex datasets.
- Identifying hidden patterns.
- Generating reports.

- Translating data into actionable insights.

According to a survey by **Deloitte**, employees using Generative AI save **1.75 hours per day**, which is equivalent to **one full workday every week**.

35.7.1 Real-World Example: Salesforce

Salesforce integrated Generative AI capabilities into its CRM (Customer Relationship Management) platform to:

- Analyze customer data.
- Enhance sales strategies.
- Improve marketing strategies.
- Identify potential leads.
- Predict customer behavior.
- Personalize customer interactions.
- Improve conversion rates.

35.8 Customer Service

Traditional customer service can be frustrating for both customers and support agents.

Repetitive activities such as:

- Answering frequently asked questions.
- Performing basic troubleshooting.

can reduce the time available for handling complex customer issues.

Generative AI-powered chatbots help by:

- Handling repetitive tasks.
- Providing fast and personalized answers.
- Responding to frequently asked questions.
- Analyzing customer data.
- Suggesting real-time troubleshooting steps.
- Helping agents resolve issues more efficiently.

35.8.1 Real-World Example: Sephora

Sephora leverages Generative AI to deliver personalized beauty recommendations.

When customers shop at Sephora, information such as:

- Personal preferences.
- Purchase history.
- Skin tone.
- Skin concerns.

is stored in a database.

Generative AI processes this data to create personalized emails that recommend products specifically suited to each customer.

35.9 Product Development

Product development teams often face challenges such as:

- Limited design options.
- Slow iterative prototyping process (the repeated process of designing, testing, and improving prototypes).

Generative AI models can:

- Produce multiple design variations.
- Help companies shortlist the most appealing designs.
- Streamline the product development process.

35.9.1 Real-World Example: Nike

Nike uses Generative AI to innovate its product design process.

By using data related to:

- [Materials](#).
- [Biomechanics](#) (the study of movement and mechanical principles in living organisms).
- [Fashion trends](#).

Generative AI generates numerous product design possibilities.

35.10 Generative AI for Startups

Generative AI is not only benefiting large organizations but is also creating opportunities for startups.

Young startups often face challenges such as:

- Limited resources.
- Tight budgets.
- High costs of creative content development.
- High costs of product design.

With Generative AI:

- Team members spend less time on mundane and repetitive tasks.
- More time can be dedicated to strategic planning.
- More focus can be placed on refinement and innovation.

Benefits for startups include:

- Cost-effectiveness.
- Increased efficiency.
- Improved innovation.
- Competitive advantage through early adoption.

35.10.1 Real-World Example: Stitch Fix

Stitch Fix, a full-stack startup, uses Generative AI to:

- Help stylists interpret customer feedback.
- Provide better product recommendations.

35.11 Future of Generative AI in Business

The possibilities with Generative AI continue to evolve.

As the technology matures, its benefits are expected to become even more significant.

The current focus is on:

- Practical business applications.
- Addressing existing challenges.
- Sustainable AI adoption.
- Responsible AI usage.
- Creative applications across industries.

35.12 Key Takeaways

Businesses are leveraging Generative AI in several important areas, including:

- **Content Generation**
 - Creating high-quality marketing materials.
 - Generating product descriptions.
 - Producing social media content.
- **Data Analysis**
 - Identifying hidden patterns.
 - Generating actionable insights.
- **Customer Service**
 - Automating repetitive tasks.
 - Providing personalized responses.
 - Improving customer satisfaction.
- **Product Development**
 - Generating multiple design variations.
 - Streamlining the prototyping process.
- **Startups**
 - Improving cost-effectiveness.
 - Increasing operational efficiency.

- Driving innovation.
- Gaining a competitive advantage through early adoption of Generative AI.

36 Generative AI for Business Transformation

36.1 Learning Objectives

After completing these exercises, we will be able to:

- Utilize Generative AI to accelerate innovation by analyzing market and technology trends.
- Leverage Generative AI to enhance various aspects of e-commerce operations, such as dynamic pricing strategies.
- Support strategic decision-making by providing data-driven insights.

36.2 Introduction

In today's competitive e-commerce landscape, Artificial Intelligence (AI) has become an essential tool for driving innovation and business growth. AI enables organizations to analyze market and technology trends, identify new opportunities, enhance operational efficiency, and improve decision-making.

In this topic, we explore how Generative AI can assist businesses in market analysis, technology adoption, and AI-driven dynamic pricing strategies.

36.3 Marketing and Technology Trends

This section explores how Generative AI can analyze market trends and emerging technologies to help an e-commerce company identify opportunities, support innovation, and make informed business decisions.

36.3.1 Analyze Recent E-Commerce Market Trends

Prompt

I'm interested in understanding the latest trends in the e-commerce industry. Please provide a summary of recent trends based on recent articles, blogs, and industry reports.

36.3.2 Analyze Recent Technological Innovations

Prompt

What are the latest technological innovations driving the e-commerce industry forward? Please provide examples of how these technologies are transforming e-commerce.

36.3.3 Implementing AR/VR in an E-Commerce Platform

Prompt

We're considering implementing AR/VR (Augmented Reality/Virtual Reality) features on our e-commerce platform. Can you provide detailed steps and best practices for implementing these features?

36.3.4 Best Practices for Implementing AR/VR

- Clearly define the business objectives for implementing AR/VR.
- Select the appropriate AR or VR technology based on business requirements.
- Ensure compatibility with supported hardware and software platforms.
- Create high-quality 3D product models.
- Optimize the user experience with intuitive interactions and realistic rendering.
- Prioritize mobile optimization where appropriate.
- Conduct thorough testing across multiple devices and environments.
- Continuously gather user feedback and refine the solution.

Consider privacy, development cost, device compatibility, and content creation challenges during implementation.

36.4 AI-Driven Dynamic Pricing

Dynamic pricing is a pricing strategy in which product prices are adjusted in real time based on factors such as market demand, competitor pricing, and customer behavior.

In this section, we explore how Generative AI can assist organizations in planning and implementing AI-driven dynamic pricing for e-commerce businesses.

36.4.1 Generate a Dynamic Pricing Implementation Plan

Prompt

We are planning to implement dynamic pricing using AI to optimize sales of our best-selling electronics and home appliances during the holiday season.

36.4.2 Key Steps for Implementing Dynamic Pricing

- Collect relevant business and market data.
- Select suitable AI tools and platforms.
- Design and train the pricing algorithm.
- Pilot and evaluate the pricing strategy.
- Deploy the solution and continuously monitor and improve its performance.

36.4.3 Identify Data Required for Dynamic Pricing

Prompt

What data should we consider collecting for dynamic pricing?

Important Data Sources

- Historical sales data
- Competitor pricing
- Market trends
- Customer demographics and purchasing behavior
- Promotions and discount calendars

36.5 Summary

In this topic, we explored how Generative AI supports business transformation by:

- Analyzing market and technology trends.
- Identifying emerging innovations.
- Assisting organizations in implementing AI-powered solutions such as AR/VR experiences and dynamic pricing.
- Providing data-driven insights that help businesses improve customer experiences.
- Optimizing pricing strategies.
- Supporting strategic decision-making.

37 Fire, Generative AI, and the Foundation of AI Value Creation

37.1 Fire as the Foundation of Human Progress

When oxygen, heat, and fuel combine, we get fire. It is basic, it is primal, and it is the key that unlocked human progress.

Fire provided:

- Light
- Heat
- Protection

With fire, our ancestors could move to new climates and eat new foods.

Pottery, metallurgy, chemistry, rapid transportation, and many other technologies all start with fire.

Imagine if fire had been proprietary. What if that knowledge had not been shared? What if there had been just a few keepers of the fire? What would we be?

37.2 A Landmark Moment for Generative AI

We are in a landmark moment with generative AI that will shape our society for generations to come.

Chances are we are already AI users. Even if we do not have a formal AI initiative or an AI team in place, AI is already in use across our business.

It might be baked into our off-the-shelf applications. It may appear in:

- Chatbots
- HR self-service portals
- Transcription services

These are what we might call **traditional AI**.

37.3 Traditional AI

Traditional AI focuses on executing discrete tasks.

Its characteristics include:

- Each AI instance is usually trained individually.
- Each model is trained using its own data.
- The required data must be compiled and manually labeled before training.

37.4 What Makes Generative AI Different

Generative AI is enormously flexible and is not limited to narrow tasks.

Instead of simply filling specific blanks, it can:

- Write an entire document from scratch.
- Create or generate something entirely new.

What makes this possible are **foundation models**.

37.5 Foundation Models

Foundation models are not trained the same way as traditional AI.

Instead, they are trained using **self-supervised learning** (a training approach where the model learns patterns from unlabeled data without requiring humans to manually label every example).

This means we do not have to manually annotate a massive amount of data.

Instead, we tell the models to read enormous amounts of text, and they do.

The result is:

- A large model
- A versatile model
- More human-like language capabilities

Algorithms use mathematical models to represent the relationships between the words they ingest.

If we provide the model with a few words in a prompt, it can mathematically predict the likelihood of words appearing in the response.

Instead of building one AI model for each specific task, we can train one model and adapt it to many downstream tasks (specialized tasks performed after pre-training).

We moved from:

- One task → One model

to:

- One model → Many tasks

37.6 One Foundation Model, Many Applications

The same foundation model can power multiple business applications.

For example, the same model can support:

- Chatbots
- HR self-service systems
- Applications that write marketing emails
- Applications that summarize legal documents

This is the first critical point.

Ideally, a model is not the final form of AI. It is simply the foundation that we build upon.

How we use it is up to us.

37.7 Three Modes of AI Consumption

When it comes to using AI, there are basically three modes of consumption.

- Embedded AI
- API-based AI
- Platform AI

37.7.1 Embedded AI

Embedded AI is baked into off-the-shelf software.

[In this model:](#)

- [The software vendor creates the AI.](#)

- We use that AI within our business.

Examples include:

- Writing assistants that help us strike the right tone in emails.
- Image-editing software that automatically processes images.

These capabilities can significantly improve productivity.

However, there is an important caveat.

Whatever AI capabilities we purchase are also available to everyone else.

As a result:

- These capabilities do not become competitive differentiators.
- Instead, they establish a higher productivity baseline for everyone.

37.7.2 API-Based AI

The second model of AI consumption is through API (Application Programming Interface) calls.

As we develop custom applications for our business, those applications can call another company's AI service using that company's models and processes and then return the results.

This is also a viable way of consuming AI.

Depending on:

- How cleverly we use the APIs
- The diversity of AI service providers we choose

we can begin differentiating how AI is applied relative to our competitors.

However, this approach also comes with important caveats.

First, just like off-the-shelf software, the models and services we access through APIs are available to everyone.

Second, when an API call is made, it connects to what is essentially a black box.

We do not necessarily know:

- What is happening on the other end.
- The provenance (origin and history) of the data being used.
- The governance (management, control, and oversight) of that data.

This uncertainty can make organizations nervous because our business remains accountable for the final outcomes.

Another important caution concerns long-term value creation.

Historically, many technology businesses have followed value-extractive business models.

Another company offers us a service, such as an AI API.

While we gain value from using that service, the provider also extracts value from:

- Our usage
- Our data

Over time, that provider continues accumulating more value.

The imbalance becomes visible over time, even in company stock prices.

This unequal relationship can have long-term consequences for:

- Our individual business
- The broader economy
- The overall progress of technology

37.8 Platform AI: Becoming Our Own AI Value Creator

The third model of AI consumption is the **platform model**.

This is the most comprehensive approach. It is how we become our own AI fire starter.

This does **not** mean doing everything alone, reinventing AI from scratch, or spending years and millions of dollars building our own models.

With a platform, we have all the essential elements and ingredients needed to build our own AI solutions.

These include:

- Foundation models
- Tools to improve and customize models
- Processes to build tailored AI solutions

Most importantly, we create and accrue value that is unique to our business.

Ultimately, most businesses should end up with a mix of all three AI consumption models.

We will:

- Use third-party software with embedded AI.
- Sometimes use someone else's AI through APIs where appropriate.
- Adopt a platform approach to create and own our AI in order to fully realize AI's value and differentiate ourselves from competitors.

37.9 Foundation Models as the Basis for AI Value Creation

Let us go a bit deeper into AI value creation, starting with foundation models.

Foundation models are **large-scale deep neural networks** trained with large amounts of **unlabeled data** and subsequently adapted to many downstream tasks.

A foundation model may be:

- A broad general model
- A narrow, deeper model

The key characteristic is that it is pre-trained with the expectation that we can further enhance it using our own proprietary data.

A useful comparison is a new employee joining our business.

A new employee arrives with:

- General skills as a foundation
- The ability to learn

As that employee learns more about our business, they gain institutional knowledge and expertise, allowing them to deliver greater value.

The same principle applies to a foundation model.

Using our AI platform, we tune the model with:

- Our specific business data
- Our proprietary knowledge
- Our organizational expertise

As a result, the model becomes:

- More knowledgeable
- More specialized
- More valuable for our business over time

Because we control the platform, the processes, and the data, we continue to accrue increasing value over time.

37.10 Risks of Losing Control Over AI

With some consumer AI systems already available in the market, we have seen what can happen when we surrender that control.

Potential risks include:

- Bad data leading to bad outcomes.
- Hallucinations (AI generating very confident but incorrect answers).
- Inadvertently using someone else's rights-managed content.
- Proprietary or sensitive data being unintentionally leaked into public spaces.

These risks highlight why we need to understand:

- How our model is built.
- What data was included during training.

They also demonstrate why tight control over sensitive business data should be a priority.

37.11 The Importance of AI Governance

Strong AI governance is absolutely critical.

Now is the time to adopt AI, but we should look before we leap.

We should ensure that we invest in an approach that is:

- Smart
- Safe
- Sustainable

The primary beneficiaries of our AI strategy should be:

- Our business
- Our clients
- Our customers

37.12 The Future of AI Is Multi-Model

There is a common myth about AI today—a basic misunderstanding.

For many people, generative AI appears to have emerged suddenly.

Some believe that only a handful of consumer-oriented AI experiences will dominate the future and that one model—or only a few leaders—will ultimately win.

However, that is unlikely to be how AI evolves.

The future of AI is **not about one model**.

It is **multi-model**.

Our business will use multiple fine-tuned models to achieve the best possible results for specific use cases.

This is another reason why the platform approach is so important.

37.13 Open Science and Open Source Will Shape AI

Realistically, the future of AI is not only proprietary.

It will also be powered by:

- Open science
- Open source

Proprietary models will certainly play a role.

However, much of AI's future development will occur openly rather than behind closed doors.

It will happen with:

- Full transparency
- Greater accountability

These are important strengths provided by open-source development.

The energy within the open-source community is phenomenal.

Innovation is being driven by:

- Distributed projects
- University projects
- Corporate efforts

Together, these communities are producing foundation models that we can tune and deploy for our own use cases.

One example is **Hugging Face**, which is described as being like **GitHub** for foundation models.

At the time of this discussion, there are **over 325,000 open-source models available**, with **thousands more being added**.

This is exactly how AI should evolve.

For the long-term benefit of society, we do not want only one or a few companies defining what AI is and dictating how it should be used.

Instead, what we are likely to see is the **democratization of AI**.

37.14 Becoming an AI Value Creator

We have covered many different topics, and it may feel like there is a lot of work required to become great at AI.

However, that is not the intention.

The goal is to take down barriers to participation, not put them up.

There is a sense of urgency and fear about waiting too long and missing the moment. That is okay.

Everyone is in the same situation and experiencing similar emotions.

Instead of rushing into fast and easy options simply to check the box and say, "**Yes, we put AI in the business,**" we should take a thoughtful, deliberate, and strategic approach.

Our focus should be on:

- An AI platform
- Data management
- AI governance

If we become AI value creators, we will be in a position to succeed over the long term.

We will not have to start over every time the wind changes direction.

37.15 Building AI Capabilities for the Future

We may not have all the products and processes we need today.

We may not have all the necessary skills in place.

That is fine.

If we are not ready to start alone, we should find a partner we can trust and begin the journey together.

Doing so allows us to take advantage of the extraordinary opportunities that generative AI will unlock for our business.

37.16 The Long-Term Impact of Generative AI

This chapter in technology is incredibly exciting.

Together, we are going to use AI to reshape not only our digital world but also our physical world.

We are going to use AI to help tackle some of our toughest challenges, including:

- Social problems
- Medical problems
- Environmental problems
- Many other important challenges

We will achieve this through science while also empowering businesses like ours to do more, faster, and more responsibly.

Whatever work we do, AI will become a powerful new tool that helps us do it better.

37.17 AI Value Creators Will Shape the Future

The organizations and individuals who will make the biggest impact are the **AI value creators**.

These are the people who take the amazing foundational technology of generative AI and use it to build entirely new solutions.

That is why the goal is to make AI accessible to everyone and put it into our hands.

The future of AI depends not only on the technology itself but also on how thoughtfully, responsibly, and creatively we choose to use it.

We cannot wait to see what we will do with it.

38 Introduction to Retrieval-Augmented Generation (RAG)

38.1 Introduction

Retrieval-Augmented Generation (RAG) is a cutting-edge approach in the field of Natural Language Processing (NLP) that combines the strengths of retrieval-based and generation-based models. This hybrid method is particularly effective for tasks that require generating informative and contextually relevant text, such as question answering, dialogue systems, and content creation.

38.2 Objectives

After completing this topic, we will be able to:

- Explain the working of the RAG model.
- Describe how RAG addresses the limitations of generative AI models.
- Explore the applications and use cases of RAG.

38.3 Retrieval-Augmented Generation (RAG)

RAG is a hybrid approach that combines generative models with retrieval-based methods.

38.3.1 How RAG Works

RAG works through a three-step process:

1. **Retrieve Information**

- Given a query, the model retrieves relevant documents or pieces of information from a predefined corpus or database.
- 2. **Augment the Input**
 - The retrieved documents are used to augment the input to the generative model, providing it with additional context.
- 3. **Generate the Response**
 - The generative model uses both the original query and the retrieved information to generate a response, ensuring that the output is both contextually rich and factually grounded.

This process enhances the accuracy and relevance of the generated content, addressing many limitations of standalone generative models.

38.4 Limitations of Generative AI Models

Generative AI models, such as **GPT-3** and **GPT-4**, have several limitations:

- They can produce information that sounds plausible but is factually incorrect or fabricated, a phenomenon known as **hallucination** (generation of false or invented information presented as factual).
- They have a fixed knowledge cutoff date, meaning they lack access to information beyond their last training update, making them outdated for generating current content.
- They have limited context windows, making them struggle with tasks requiring long-term context or detailed memory over extended conversations or documents.
- While they can generate fluent and coherent text, they might lack the depth and specificity required for certain queries, especially those requiring detailed or specialized knowledge.
- Generating high-quality text can be computationally expensive and slow, particularly for complex or long-form content.

38.5 How RAG Addresses the Limitations of Generative AI Models

RAG effectively overcomes many limitations of traditional generative AI models.

- By retrieving relevant documents or data from a corpus before generating a response, RAG grounds the generation process in factual and contextually accurate information, thereby reducing hallucination.
- By leveraging real-time retrieval from up-to-date databases, RAG provides more current and relevant information, mitigating the issue of the knowledge cutoff.
- Retrieval mechanisms enhance contextual understanding by pulling relevant documents or chunks of text, extending the effective context window and maintaining coherence over longer interactions.
- The retrieval step ensures that the generative model has access to specific and detailed information relevant to the query, improving the specificity and depth of responses.
- RAG models can also be more efficient because the retrieval step narrows the information space, allowing the generative component to focus on synthesizing relevant information instead of generating everything from scratch.

38.6 Key Components of RAG

38.6.1 Retrieval Component

Function

The retrieval component is responsible for searching and extracting relevant information from a large corpus of documents or a knowledge base. This ensures that the model has access to factual and contextually appropriate information.

Mechanism

Typically, this involves using a retriever model such as:

- BM25
- Dense retrievers based on neural networks

These retrievers find the most relevant passages or documents that match a given query.

38.6.2 Generation Component

Function

The generation component takes the retrieved information and uses it to generate coherent and contextually appropriate responses or text. This is achieved using a generative model such as **GPT-3** or **BERT**.

Mechanism

The generative model leverages the context provided by the retrieved documents to produce more accurate and relevant outputs by blending retrieval results with its generative capabilities.

38.7 Benefits of RAG

RAG provides several important benefits:

- Improves the accuracy of responses by grounding generated text in actual data, enhancing factual correctness.
- Boosts contextual relevance by incorporating real-time information retrieval, ensuring responses are pertinent to the query.
- Offers flexibility, allowing it to be adapted for various NLP tasks.
- Retrieves the latest information, ensuring responses remain dynamically updated and current.

38.8 Applications of RAG

RAG has numerous real-world applications, including:

- **Question-answering systems:** Retrieves relevant documents to answer complex questions accurately.
- **Content creation:** Assists in generating detailed and informative content for articles, reports, and creative writing.
- **Customer support:** Provides accurate and contextually relevant answers by retrieving the latest information from knowledge bases.
- **Search engines:** Improves search results by providing detailed and accurate answers based on retrieved documents.

38.9 Implementation of RAG on Google Cloud

Google Cloud provides robust tools and frameworks for implementing RAG models, including **Vertex AI** and **BigQuery**. These platforms facilitate the integration of retrieval mechanisms with Large Language Models (LLMs), enabling the development of scalable and efficient RAG applications.

38.9.1 Vertex AI

- Offers a comprehensive suite for building and deploying machine learning models.
- Includes support for RAG frameworks.

38.9.2 BigQuery

- Enables efficient querying and retrieval of large datasets.
- Provides a robust backend for the retrieval component of RAG models.

38.10 Key Features of RAG on Google Cloud

- **Scalability:** Google Cloud infrastructure enables RAG models to handle large-scale data retrieval and processing.
- **Integration:** Seamless integration with various data sources and APIs facilitates comprehensive retrieval capabilities.
- **Customization:** Provides tools and frameworks for customizing RAG models according to specific business needs and applications.

38.11 Example

Consider a RAG-based system designed to answer questions about historical events.

When asked:

"What were the key causes of World War II?"

The system:

1. Retrieves documents or passages from a history database discussing the causes of World War II.
2. Generates a comprehensive answer based on the retrieved information.

3. Ensures that the response is both accurate and informative.

38.12 Example Use Case: Enhancing LLMs with RAG

By integrating RAG with **Google Cloud BigQuery**, organizations can enhance the capabilities of their Large Language Models (LLMs).

For example:

A customer support application can use RAG to retrieve the latest policy documents from a database, ensuring that the LLM's responses remain accurate and current.

38.13 Summary

In this topic, we learned the following:

- RAG is an advanced method that combines the strengths of generative AI models with external information retrieval mechanisms to address several limitations inherent in generative models.
- RAG addresses the limitations of generative AI models by retrieving relevant documents or data from a corpus before generating a response.
- RAG incorporates a retrieval step to ensure that the generative model is grounded in actual data, reducing the risk of hallucinations.
- Major applications of RAG include:
 - Question-answering systems
 - Content creation
 - Customer support
 - Search engines

39 Adopting AI in Your Business

39.1 Learning Objectives

After studying this topic, we will be able to:

- Explain the key benefits of adopting AI in business operations.
- Explore how businesses can effectively integrate Artificial Intelligence (AI) into their operations.
- Describe the general steps required to adopt AI in a business.

39.2 Introduction

AI is transforming business operations across industries by helping businesses:

- Surpass competitors.
- Improve customer interactions.
- Develop new ideas.
- Create new opportunities for efficiency and growth.

Examples include:

- E-commerce companies leveraging AI for personalized recommendations.
- Manufacturing firms using AI-powered robots to improve operations.

The possibilities of AI adoption in business are extensive and continue to expand.

39.3 Benefits of Adopting AI in Business

Businesses gain several advantages by adopting AI.

- **Increases efficiency** by automating repetitive tasks, freeing employees to focus on higher-value work.
- **Enhances decision-making** by analyzing vast amounts of data to predict market trends and consumer behaviors.
- **Improves customer service** through data-driven personalized recommendations, increasing customer satisfaction.
- **Accelerates innovation** by enabling businesses to create unique offerings and differentiate themselves from competitors.

39.3.1 Industry Insight

A recent global study conducted by **Tata Consultancy Services (TCS)** found that:

- More than **8 out of 10**, or **86%**, of senior business leaders worldwide have adopted AI to:
 - Boost existing revenue streams.
 - Generate new revenue opportunities.

39.4 Real-World Examples of AI Adoption

39.4.1 Amazon

[Amazon utilizes AI across multiple areas of its business operations.](#)

39.4.1.1 Personalized Product Recommendations

Amazon's recommendation engine uses AI algorithms to analyze:

- User browsing history.
- Purchase patterns.
- Demographic data.

This enables Amazon to recommend products that match customer preferences.

Benefits include:

- Increased sales.
- Enhanced customer satisfaction.

- Reduced search time for relevant products.

39.4.1.2 Alexa Voice Assistant

Amazon's AI-powered voice assistant, **Alexa**, uses:

- Natural Language Processing (NLP) (*technology that enables computers to understand and respond to human language*).
- Machine Learning.

Alexa enables users to:

- Place orders.
- Control smart home devices.
- Access information using voice commands.

39.4.1.3 Supply Chain and Inventory Management

Amazon uses AI to optimize:

- Inventory management.
- Warehouse operations.
- Supply chain processes.

AI-powered forecasting models analyze:

- Historical data.
- External factors.

These models help to:

- Predict product demand.
- Optimize inventory levels.
- Minimize stock-outs.
- Streamline logistics.

As a result, Amazon reduces operational costs and improves overall delivery efficiency.

39.4.2 Tesla

Tesla leverages AI extensively in the automotive industry, particularly in autonomous driving technology.

39.4.2.1 Autopilot System

Tesla's Autopilot uses AI algorithms, including:

- Computer Vision (*AI that enables computers to interpret images and videos*).

- Deep Learning (*a subset of machine learning that uses neural networks to learn complex patterns*).

The system analyzes data collected from:

- Cameras.
- Radar.
- Sensors.

It performs tasks such as:

- Navigation.
- Obstacle detection.
- Real-time driving decisions.

These capabilities improve both driving safety and efficiency.

39.4.2.2 Summon Feature

Tesla's **Summon** feature uses AI to detect and react to its surroundings.

It allows drivers to remotely park their vehicles using a smartphone application within a specified range.

39.4.2.3 Battery Management System (BMS)

Tesla's **Battery Management System (BMS)** (*a system that monitors and manages battery performance and efficiency*) uses AI to:

- Predict energy requirements based on:
 - Driving conditions.
 - Driver behavior.
- Adjust acceleration.
- Control climate settings to conserve energy.

It also assists the navigation system by planning efficient routes while considering:

- Traffic conditions.
- Elevation.

39.5 General Roadmap for Adopting AI in Business

Adopting AI is a multifaceted process that requires careful planning and execution.

The major steps include:

- Defining business goals.
- Identifying suitable use cases.
- Preparing data.

- Building AI capabilities.
- Deploying AI solutions.
- Monitoring and optimizing AI systems.

Each step plays an important role in achieving successful AI implementation.

39.6 Define Business Goals

The journey begins by clearly defining business goals.

Businesses should identify the specific problems they want AI to solve by engaging stakeholders to understand their needs and establish clear objectives.

Important questions to consider include:

- What are the pain points?
- How can AI help achieve business goals?

39.7 Identify Suitable Use Cases

After defining business goals, the next step is identifying suitable use cases.

Businesses should perform a thorough analysis of their processes to determine areas where AI can be applied effectively.

Priority should be given to tasks that are:

- Repetitive.
- Time-consuming.
- Data-driven.
- Dependent on complex decision-making.

39.8 Prepare Data (Data Readiness)

Data readiness is one of the most critical steps in AI adoption.

Businesses should:

- Collect data.
- Clean data.
- Organize data.

The data must be ready for AI algorithms.

Data readiness involves:

- Removing inconsistencies.
- Filling missing gaps.
- Ensuring the data is relevant to the business problem.

39.9 Build AI Capabilities

Organizations should develop AI capabilities internally.

This includes:

- Training employees on AI concepts.
- Learning AI tools.
- Understanding AI techniques.

Investing in workforce upskilling ensures employees can effectively use and manage AI technologies.

Training programs help employees understand how AI can be applied to specific business tasks.

39.10 Deploy AI Solutions

The next step is deploying AI solutions.

Businesses should:

- Select AI tools and technologies that match business requirements.
- Deploy solutions gradually, beginning with pilot projects before scaling.

It is important to ensure AI solutions integrate seamlessly with existing systems and business processes.

39.11 Monitor and Optimize AI Systems

After deployment, AI systems require continuous monitoring and optimization.

Businesses should regularly:

- Track AI performance.
- Identify issues.
- Resolve problems.
- Update AI models.
- Refine AI systems.
- Make necessary adjustments over time.

Continuous optimization ensures AI systems continue to perform effectively.

39.12 Key Takeaways

AI adoption enables businesses to:

- Increase operational efficiency.

- Enhance decision-making.
- Improve customer service.
- Drive innovation.

Real-world organizations such as **Amazon** and **Tesla** demonstrate how AI can deliver measurable business value across multiple operations.

A general roadmap for adopting AI includes:

1. Define clear business goals.
2. Identify suitable AI use cases.
3. Prepare data by collecting, cleaning, and organizing it.
4. Build AI capabilities by training employees and developing the necessary organizational setup.
5. Deploy AI solutions and integrate them with existing systems.
6. Monitor and optimize AI systems to ensure long-term performance and continuous improvement.

40 Frameworks for AI Adoption

40.1 Overview

Organizations are increasingly adopting **Artificial Intelligence (AI)** to drive innovation and gain a competitive edge. However, successfully harnessing the power of AI can be challenging. To address this, organizations are turning to **AI adoption frameworks** to guide their AI implementation efforts.

These frameworks provide a structured approach to:

- Identifying the right AI use cases
- Aligning AI initiatives with business goals
- Creating a culture that supports responsible AI development and deployment

One comprehensive framework designed to guide enterprises through this transformative journey is the **IBM AI Ladder**.

40.2 IBM AI Ladder

The **IBM AI Ladder** is a structured framework that guides organizations from basic data collection to full-scale AI integration through four key stages:

- Collect
- Organize
- Analyze
- Infuse

40.2.1 Collect

The **Collect** stage involves gathering and storing high-quality data from diverse sources.

IBM tools:

- IBM Cloud
- watsonx.data

Key actions:

- Identify data sources such as databases and IoT (Internet of Things) devices.
- Implement data collection tools.
- Ensure data quality.

Example

A retail company uses **watsonx.data** to gather sales data from various channels.

40.2.2 Organize

The **Organize** stage focuses on cleaning, categorizing, and making data accessible for analysis.

IBM tools:

- IBM Cloud Pak for Data
- IBM InfoSphere
- IBM DataOps
- IBM watsonx.governance

These tools help with:

- Data management
- Data governance
- Data security

Key actions:

- Cleanse data.
- Categorize data for easy retrieval.
- Implement governance policies.
- Centralize data storage.

40.2.3 Analyze

The **Analyze** stage harnesses AI's potential through advanced analytics and machine learning models.

IBM tools:

- IBM SPSS
- IBM Cognos Analytics
- IBM Watson Machine Learning

- IBM watsonx.ai

These tools enable efficient:

- Model building
- Model testing
- Model deployment

Key activities:

- Apply statistical and machine learning algorithms.
- Develop predictive models.
- Visualize insights.
- Refine models based on performance.

Example

Data scientists use **watsonx.ai** to analyze customer behavior.

40.2.4 Infuse

The **Infuse** stage integrates AI into daily business operations to enhance decision-making and automate tasks.

IBM tools:

- IBM watsonx APIs
- Robotic Process Automation (RPA)
- IBM Cloud Pak for Automation

These tools help with:

- Embedding AI models into business applications
- Automating business processes
- Monitoring AI models in real time
- Training employees in AI utilization

Example

AI insights are integrated into **Customer Relationship Management (CRM)** systems to:

- Automate marketing campaigns
- Optimize inventory management

40.3 Benefits of the IBM AI Ladder

The IBM AI Ladder transforms data into actionable insights and business strategies.

It supports organizations across industries, including:

- Healthcare
- Finance
- Manufacturing
- Retail

40.4 From +AI to AI+

In the current business landscape, where AI is becoming increasingly prevalent, company leaders need to shift their mindset from the **+AI** model to the **AI+** approach.

40.4.1 +AI Approach

In the **+AI** model:

- AI is used as a supplementary tool.
- AI is incorporated into existing workflows.
- Organizations gain only limited, incremental advantages.

40.4.2 AI+ Approach

The **AI+** approach integrates AI holistically throughout the organization.

Instead of treating AI as an afterthought or add-on, business processes are designed with AI as a fundamental component from the outset.

AI becomes a core part of the business rather than just a supporting tool.

40.5 Key Elements of the AI+ Approach

40.5.1 Holistic Integration

AI should be seamlessly integrated across all business functions.

This includes:

- Embedding AI into product innovation
- Enhancing operational efficiency
- Fostering a culture that embraces AI-driven decision-making

The goal is to make AI a core part of the organization's DNA, influencing every aspect of the business, from strategy to execution.

40.5.2 Identify High-Value AI Use Cases

Organizations should identify AI use cases that deliver significant returns on investment (ROI).

The recommended approach is to:

- Begin with projects that have clear, measurable outcomes.

- Focus on projects with substantial business impact.
- Scale successful projects across the organization.

This ensures AI investments are strategically aligned with business goals and provide tangible benefits.

40.5.3 Choose the Right AI Technology

Successful AI integration requires selecting appropriate AI technologies and architectures.

Organizations should:

- Evaluate different AI tools and platforms.
- Ensure technologies align with specific business needs and goals.
- Select technologies that are scalable and flexible.

40.5.4 Build a Robust Data Foundation

A robust data foundation is essential for effective AI implementation.

Organizations should ensure that data is:

- Trustworthy
- Transparent
- Well-governed

High-quality data is critical for:

- Training AI models accurately
- Making reliable predictions

Organizations should also comply with:

- Data privacy regulations
- Data security regulations

This helps avoid legal and ethical issues.

40.5.5 Continuous Innovation and Modernization

Organizations should continuously:

- Develop new AI applications.
- Modernize existing AI applications.

Leveraging a **Hybrid Cloud** platform provides the flexibility and scalability needed to support ongoing AI initiatives.

Continuous improvement and adaptation ensure AI solutions remain relevant and effective in addressing evolving business challenges.

40.6 Key Takeaways

40.6.1 IBM AI Ladder

The IBM AI Ladder provides a structured framework for AI adoption through four stages:

- **Collect** – Gather and store high-quality data.
- **Organize** – Make data simple, accessible, and business-ready.
- **Analyze** – Build and scale AI with trust and transparency.
- **Infuse** – Operationalize AI throughout the business.

40.6.2 AI+ Approach

The AI+ approach focuses on:

- Holistic AI integration across the organization
- Identifying high-value AI use cases
- Choosing the right AI technologies
- Building a robust and trustworthy data foundation
- Continuous innovation and modernization

41 AI Adoption Frameworks: Amazon, OpenAI, and Facebook

41.1 Introduction

AI adoption frameworks provide a comprehensive guide for organizations to integrate AI technologies effectively and ethically into their operations. These frameworks help organizations navigate the complex landscape of AI implementation, ensuring that AI projects are aligned with business goals while adhering to ethical standards.

41.2 Objective

After completing this topic, we will be able to explain the Amazon AI Services, OpenAI, and Facebook AI integration frameworks for adopting AI.

41.3 Amazon AI Services Framework

The Amazon AI Services Framework provides a structured approach for organizations adopting AI technologies. It is divided into four phases involving data preparation, model development, deployment, and optimization.

41.3.1 Data Preparation

Data preparation involves gathering, storing, and preparing data for AI applications.

Actions:

- Collect data from various sources.
- Clean and preprocess data for quality assurance.
- Securely store the prepared data.

Tools:

- Amazon S3
- AWS Glue
- Amazon Redshift

Example:

An e-commerce company uses **Amazon S3** to collect customer data, cleans it with **AWS Glue**, and stores it in **Amazon Redshift** for analysis.

41.3.2 Model Development

Model development focuses on building and training machine learning models.

Actions:

- Select appropriate machine learning algorithms.
- Train models using prepared datasets.
- Refine models to improve accuracy.

Tools:

- Amazon SageMaker
- AWS Deep Learning AMIs
- AWS Lambda

Example:

An e-commerce company uses **Amazon SageMaker** to create a recommendation engine, refining its accuracy through iterative testing on historical data.

41.3.3 Deployment

Deployment integrates AI models into production environments.

Actions:

- Integrate AI models with existing applications.
- Automate deployment processes.
- Monitor model performance in real time.

Tools:

- Amazon SageMaker
- AWS Lambda
- Amazon CloudWatch

Example:

A recommendation engine is deployed using **Amazon SageMaker**, updates are automated with **AWS Lambda**, and performance metrics are monitored using **Amazon CloudWatch**.

41.3.4 Optimization

Optimization focuses on the continuous improvement of AI models and business processes.

Actions:

- Analyze model performance data.
- Optimize models for efficiency and accuracy.
- Scale AI applications across the organization.

Tools:

- Amazon SageMaker Debugger
- Amazon Personalize
- AWS Step Functions

Example:

Amazon SageMaker Debugger is used to enhance model performance, **Amazon Personalize** scales recommendation capabilities, and **AWS Step Functions** integrates AI workflows across business units.

41.3.5 Key Benefit

The Amazon AI Services Framework helps organizations leverage AI to improve operational efficiency and enhance customer experience across various industries.

41.4 OpenAI Framework

The OpenAI Framework provides a structured approach to AI adoption through four phases:

- Data preparation
- Model development
- Model deployment
- Continuous improvement

41.4.1 Data Preparation

Data preparation focuses on gathering, cleaning, and preparing data for AI applications.

Actions:

- Collect data from diverse sources.
- Preprocess data using appropriate tools.
- Ensure data quality and governance.

Tools:

- OpenAI API
- Pandas
- NumPy

Example:

A content company collects articles through web scraping, cleans the data using **Pandas** and **NumPy**, and securely stores it in cloud storage.

41.4.2 Model Development

Model development focuses on building and training AI models.

Actions:

- Select AI models such as OpenAI GPT-3 and Codex.
- Train models using cleaned datasets.
- Refine models using iterative experimentation.

Tools:

- OpenAI GPT-3
- Codex
- Jupyter Notebooks

Example:

A content company uses **GPT-3** to generate marketing copy and refines the model using **Jupyter Notebooks**.

41.4.3 Model Deployment

Model deployment integrates AI into operational environments.

Actions:

- Integrate AI models into applications.
- Scale deployments for production workloads.
- Ensure reliable performance.

Tools:

- OpenAI API
- Docker
- Kubernetes

Example:

GPT-3 is deployed through the **OpenAI API**, packaged with **Docker**, and scaled using **Kubernetes** to maintain reliable performance under heavy workloads.

41.4.4 Continuous Improvement

Continuous improvement focuses on refining AI models and business processes.

Actions:

- Collect user feedback.
- Monitor model performance.
- Optimize models using real-time data.

Tools:

- OpenAI API
- TensorBoard
- Google Analytics

Example:

TensorBoard and **Google Analytics** are used to refine GPT-3 using real-time data, resulting in improved user engagement and higher conversion rates.

41.5 Facebook AI Integration Framework

Facebook integrates AI technologies throughout its platform to improve user experience and business functionality. Its approach focuses on AI-driven data processing, content recommendation, and user interaction.

41.5.1 Data Integration

Data integration prepares data for AI applications.

Actions:

- Collect user interaction data.
- Gather content preferences and behavioral insights.
- Clean, preprocess, and format data for AI analysis.
- Ensure compliance with privacy regulations.

Tools:

- Facebook Graph API
- Facebook Analytics

Important Note:

Facebook maintains data privacy and complies with regulations such as **GDPR (General Data Protection Regulation)** while collecting insights into user preferences and behaviors.

41.5.2 AI Model Development

AI model development focuses on creating models for personalized recommendations and content moderation.

Actions:

- Develop machine learning models.
- Train models using extensive datasets.
- Validate models for performance and scalability.

Tools:

- FAIR's AI Research Tools (Facebook AI Research Tools)
- PyTorch

Example:

Facebook trains AI models on large datasets containing user interactions and content features to deliver accurate personalized recommendations while improving content moderation.

41.5.3 AI Model Deployment

AI model deployment integrates AI models into Facebook's platform for real-time applications.

Actions:

- Integrate models into News Feed algorithms.
- Integrate models into ad targeting systems.
- Automate deployment and model updates.
- Continuously adapt models using user interactions and feedback.

Tools:

- Facebook Developer API
- AI Infrastructure

Outcome:

These deployments improve content personalization and increase the effectiveness of advertisement targeting.

41.5.4 Continuous Improvement

Continuous improvement optimizes AI models and algorithms over time.

Actions:

- Collect user feedback.
- Analyze performance metrics.
- Retrain models.
- Update algorithms based on insights.

Tools:

- Facebook Analytics
- AI Research Tools

Outcome:

Continuous retraining improves content recommendations and enhances user engagement metrics.

41.6 Summary

The major components of AI adoption frameworks are:

41.6.1 Data Preparation

- Collecting, cleaning, and organizing data.
- Ensuring data quality and relevance for model training.

41.6.2 Model Development

- Designing and training AI models.
- Using algorithms and frameworks to build predictive models.

41.6.3 Deployment

- Integrating AI models into production environments.
- Ensuring models are scalable and reliable.

41.6.4 Optimization

- Continuously monitoring and improving model performance.
- Updating models using new data and user feedback.

42 Transforming Work Through AI Tools

42.1 Learning Objectives

After completing this topic, we will be able to:

- List the key AI tools utilized by various professionals.
- Demonstrate the practical applications of AI tools across different industries.

42.2 AI Tools in Today's Digital Age

In today's digital age, AI tools and platforms play vital roles across multiple industries.

A report published by **Statista** shows a huge rise in global AI tool users. It is expected to reach **over 700 million by the end of the decade**.

Professionals use AI tools to:

- [Work smarter.](#)
- [Be more creative.](#)
- [Stay ahead of the curve.](#)
- [Unlock new growth opportunities.](#)

Popular AI tools such as **Microsoft Copilot** and **Google Gemini** help professionals with tasks including:

- Answering questions.
- Generating text.
- Improving productivity.

42.3 AI Tools for Research and Healthcare Professionals

42.3.1 Research Analyst: Microsoft Copilot

[Mark](#), a research analyst at a finance startup, is responsible for investigating ongoing market trends. He uses [Microsoft Edge](#) for research.

To grasp concepts quickly and save time, Mark uses **Microsoft Copilot**, a powerful AI tool that integrates seamlessly with:

- Microsoft Edge
- Microsoft Word
- Microsoft Excel
- Microsoft PowerPoint
- Microsoft Outlook

Copilot helps Mark by:

- Summarizing an article's key points.
- Defining unfamiliar terms.

- Streamlining his daily tasks.

42.3.2 Healthcare Professional: Google Gemini

Vijay, a healthcare professional, manages patient records and contact details.

He relies on **Google Gemini** to create a standardized patient intake form template.

Using a consistent template enables Vijay to:

- Capture all necessary patient information uniformly.
- Maintain accuracy.
- Improve efficiency in his work.

Microsoft Copilot and Google Gemini are valuable not only for research analysts and healthcare professionals but also for:

- Content creators
- Business operations professionals
- Customer service agents
- Many other professionals

42.4 AI Tools for Marketers, Content Creators, and Editors

Marketers, content creators, and editors use AI tools to streamline processes, generate innovative content, and achieve marketing goals.

42.4.1 Sarah's AI-Powered Content Creation Workflow

Sarah manages the social media team at a tech startup.

Her team uses the following AI tools:

AI Tool	Purpose	Practical Application
Synthesia	AI video creation	Generates digital avatars, simplifying video production without requiring extensive technical skills.
QuillBot	AI-powered writing assistant	Refines and condenses text while ensuring originality and coherence in social media posts.
Grammarly	Grammar and style checking	Ensures content is clear, correct, and stylistically consistent before publishing.

These AI tools simplify the team's workload by:

- Streamlining content creation processes.
- Producing impactful content.
- Effectively engaging audiences online.

42.5 AI Tools for Language Learning and Translation

Language learners and translators employ AI tools to personalize their learning experiences and overcome language barriers.

42.5.1 James' Language Learning and Translation Workflow

James, an international business consultant, frequently travels and communicates with clients from diverse linguistic backgrounds.

To excel in his role, James relies on several AI tools.

42.5.1.1 Duolingo

James refines his language skills using **Duolingo's AI-driven platform**.

The personalized lessons:

- Cater to his learning needs.
- Focus on work-specific vocabulary.
- Teach phrases essential for presentations.

42.5.1.2 Google Translate

During client calls, James relies on **Google Translate** to accurately translate messages in real time.

Google Translate helps him:

- Communicate seamlessly across multiple languages.
- Understand client requirements accurately.

By integrating these AI tools into his professional routine, James enhances his ability to communicate effectively with international clients.

42.6 AI Tools for Customer Support

Customer support representatives use AI tools to enhance efficiency and transform their workflow.

42.6.1 Bella's Customer Support Workflow

Bella is a customer support representative who handles a high volume of customer inquiries during her shift.

42.6.1.1 AI-Powered Chatbots

AI-powered chatbots instantly generate accurate responses for common issues, including:

- Account inquiries.
- Product troubleshooting.

This:

- Saves time.
- Ensures customer satisfaction.

42.6.1.2 Zendesk AI

Bella also uses **Zendesk's AI capabilities**, which:

- Suggest pertinent knowledge base articles.
- Recommend relevant solutions based on customer inquiries.
- Streamline issue resolution.
- Minimize escalation needs.
- Enhance response times.

42.6.1.3 LivePerson

For more personalized interactions, Bella uses **LivePerson's AI-powered messaging** with **sentiment analysis** (the process of identifying and evaluating emotions expressed in customer messages).

LivePerson:

- Analyzes customer messages.
- Evaluates customer emotions.
- Tailors responses accordingly.

42.7 AI Tools for Data Analysis and Visualization

Data professionals also use AI tools for data analysis and visualization.

42.7.1 Emily - Data Analyst

Emily, a data analyst, uses **Tableau** for intuitive data visualization.

Using **natural language queries** (asking questions in everyday language instead of a programming language), she quickly creates interactive dashboards that empower her team with actionable insights.

42.7.2 David - Data Scientist

David, a data scientist, uses **Microsoft Power BI** to:

- Analyze client data trends.
- Create clear reports.
- Build analytical models.

These visualizations help him understand:

- Client demographics.
- Resource needs.

42.8 AI Tools for Human Resources (HR)

HR professionals use AI tools to streamline recruitment, automate tasks, and enhance employee engagement efficiently.

42.8.1 Miranda's HR Workflow

Miranda is an HR manager in a tech company.

Her organization faces challenges such as:

- Low employee engagement.
- Recruitment pressures.

42.8.2 Culture Amp

Miranda implements **Culture Amp**, an AI platform that provides insights into employee engagement.

Culture Amp helps to:

- Address employee engagement issues.
- Improve workplace culture.

Such AI platforms often:

- Automate the initial support process.
- Alert management when intervention is needed.
- Enable more meaningful personal interactions.

42.8.3 IBM watsonx Orchestrate

To address recruitment pressures, Miranda leverages **IBM watsonx Orchestrate**.

Orchestrate:

- Handles interview scheduling.
- Streamlines recruitment communications.
- Works continuously to ensure the recruitment process runs smoothly.

These tools help Miranda and her team overcome challenges and reduce recruitment pressure.

42.9 Summary

Professionals from different domains leverage AI tools and platforms to improve productivity and efficiency.

The following table summarizes the AI tools and their primary applications.

Profession	AI Tool	Primary Application
Research Analyst	Microsoft Copilot	Summarizing articles, defining unfamiliar terms, streamlining research tasks
Healthcare Professional	Google Gemini	Creating standardized patient intake form templates
Marketers & Content Creators	Synthesia	Generating digital avatars for video production
Marketers & Content Creators	QuillBot	Refining and condensing written content
Marketers & Content Creators	Grammarly	Grammar and style checking
Language Learners	Duolingo	Personalized AI-driven language learning
Translators & Business Consultants	Google Translate	Real-time language translation
Customer Support Representatives	AI Chatbots	Automated responses to common customer inquiries
Customer Support Representatives	Zendesk AI	Knowledge base suggestions and issue resolution
Customer Support Representatives	LivePerson	AI-powered messaging with sentiment analysis
Data Analysts	Tableau	Interactive data visualization and dashboards
Data Scientists	Microsoft Power BI	Data analysis, reporting, and visualization
HR Professionals	Culture Amp	Employee engagement insights and workplace improvement
HR Professionals	IBM watsonx Orchestrate	Interview scheduling and recruitment automation

42.10 Key Takeaways

- AI tools are becoming increasingly important across industries.
- According to **Statista**, global AI tool users are expected to exceed **700 million by the end of the decade**.
- **Microsoft Copilot** and **Google Gemini** support a wide variety of professional tasks.
- **Synthesia**, **QuillBot**, and **Grammarly** assist marketers, editors, and content creators.
- **Duolingo** and **Google Translate** support language learning and multilingual communication.
- **AI chatbots**, **Zendesk**, and **LivePerson** improve customer service efficiency.
- **Tableau** and **Microsoft Power BI** enable effective data analysis and visualization.
- **Culture Amp** and **IBM watsonx Orchestrate** help HR professionals improve employee engagement and streamline recruitment.

43 Career Opportunities with AI

43.1 Learning Objectives

After studying this topic, we will be able to:

- Identify key AI-related roles across various industries.
- Recognize the foundational skills necessary to pursue a career in AI.
- Explore how we can effectively switch our career to AI.

43.2 AI and the Future of Jobs

Many people worry that AI might take away their jobs. While it is true that AI will change the job landscape by automating multiple tasks, it is also important to understand that it will create new opportunities.

Throughout history, technological advancements have transformed industries and the job market.

- The **Industrial Revolution** replaced some manual labor jobs but also created entirely new industries that did not exist before.
- During the **Information Age**, many people feared that computers would replace human jobs and lead to widespread unemployment. Instead, computers created career opportunities such as:
 - IT Support
 - Web Development
 - Many other technology-related professions

Similarly, AI is creating a plethora of career opportunities that were unimaginable just a few years ago. From healthcare to finance, and from entertainment to education, the demand for AI professionals is continuously increasing.

43.3 Technical AI Career Opportunities

Some of the most sought-after AI careers include:

- AI Engineer
- Data Scientist
- Robotics Engineer
- NLP (Natural Language Processing) Engineer
- AI Application Developer
- AI Research Scientist

43.3.1 AI Engineer

AI engineers specialize in the design and maintenance of AI systems.

They leverage expertise in:

- Machine Learning
- Neural Networks
- Programming languages such as:
 - Python
 - Java

43.3.2 Data Scientist

Data scientists are responsible for analyzing huge datasets using:

- Advanced Mathematics
- Statistics
- Data Visualization tools
- Programming languages such as:
 - Python
 - SQL

They use these skills to develop AI models.

43.3.3 Robotics Engineer

Robotics engineers design and build robots.

Required skills include:

- Robotics
- Machine Learning
- Mechanical Engineering
- Programming

43.3.4 NLP Engineer

NLP engineers create AI systems that process human language.

They possess expertise in:

- Machine Learning
- Neural Networks
- Programming languages such as:
 - Python
 - Java

43.3.5 AI Application Developer

AI application developers develop software applications that incorporate AI.

Required expertise includes:

- Programming languages

- AI Frameworks
- Data Analysis
- APIs (Application Programming Interfaces)
- Software Development

43.3.6 AI Research Scientist

AI research scientists advance AI through research by developing new algorithms and models.

If we are fascinated by:

- Machine Learning
- Neural Networks
- Natural Language Processing (NLP)

we can consider becoming an AI researcher.

This role demands a strong research background.

AI research scientists work in industries such as:

- Technology
- Finance
- Healthcare
- Automobile Research
- Many other industries

43.4 AI Careers That Do Not Require Strong Programming Skills

Even if we do not have strong programming skills, there are plenty of emerging AI roles in various domains.

Examples include roles related to:

- Ethics
- Management
- Communication
- Application of AI technologies

These dynamic and in-demand job roles exist across many industries and do not require highly technical or programming experience, making them accessible for people with varied backgrounds.

43.4.1 AI Ethicist

AI ethicists ensure that AI systems are developed and used ethically while considering social and moral implications.

Suitable backgrounds include:

- Ethics
- Philosophy
- Law

Important skills:

- Strong communication skills

43.4.2 AI Product Manager

AI product managers oversee the development and implementation of AI products while ensuring they meet market needs.

Required skills include:

- Project Management
- Business Analysis
- Understanding of AI technologies

43.4.3 AI Strategist

AI strategists develop long-term strategies for implementing AI to achieve business goals.

Required skills include:

- Strategic Planning
- Business Acumen
- Good understanding of AI applications

43.4.4 AI Marketing Specialist

AI marketing specialists utilize AI to:

- Analyze consumer data
- Optimize marketing strategies

Useful skills include:

- Data Analysis
- Marketing knowledge
- Familiarity with AI tools

These professionals commonly work in industries such as:

- Technology
- Marketing
- Media
- Consumer Goods
- Healthcare

- Consulting

43.5 How to Switch to an AI Career

If we are thinking about switching our career to AI, we should follow these steps.

43.5.1 Identify Existing Skills

First, identify the skills we already possess that can be leveraged.

Examples include:

- Problem-solving
- Analytical thinking
- Programming experience (if any)

43.5.2 Learn Core AI Concepts

Develop a detailed understanding of fundamental AI concepts such as:

- Machine Learning
- Deep Learning
- Natural Language Processing (NLP)

Note: Python is the most widely used programming language in AI.

43.5.3 Apply Skills Practically

Strengthen our understanding by:

- Contributing to open-source AI projects on **GitHub** (<https://github.com>)
- Developing small AI projects to solidify our learning

43.5.4 Stay Updated

Since AI is constantly evolving, it is important to stay updated.

Ways to remain current include:

- [Networking with AI professionals](#)
- [Attending conferences](#)
- [Reading publications about the latest AI trends](#)

43.5.5 Specialize in an AI Domain

As we gain experience, we can explore specialized AI fields that align with our interests, such as:

- Robotics

- Natural Language Processing (NLP)

43.6 Real-World Career Transition Examples

43.6.1 Emma

Emma leveraged her research and analytical skills from law to transition into AI for the legal field.

She now helps develop AI systems that analyze legal documents to support lawyers in their work.

43.6.2 Sarah

Sarah leveraged her strong communication and customer service skills from marketing to transition into AI.

She acquired expertise in:

- Natural Language Processing (NLP)
- Machine Learning
- Python

She now builds AI-powered chatbots that answer customer queries and provide support on company websites.

43.7 Key Takeaways

Some of the most sought-after AI careers include:

- AI Engineer
- Data Scientist
- Robotics Engineer
- NLP Engineer
- AI Application Developer
- AI Research Scientist

Several AI-related careers do not require programming knowledge, including:

- AI Ethicist
- AI Product Manager
- AI Strategist
- AI Marketing Specialist

To successfully transition into an AI-related career, we should:

- Identify our existing transferable skills.
- Learn core AI concepts.
- Apply our knowledge through practical projects.
- Stay updated with the latest AI trends.

- Develop expertise in specialized AI domains that align with our interests.

44 Using Generative AI for Your Work

44.1 Learning Objectives

After completing these exercises, we will be able to:

- Leverage generative AI to create resume templates.
- Apply generative AI to generate email templates for various professional scenarios.
- Use generative AI to identify and enhance skill sets for specific roles.

44.2 Introduction

Generative AI has become a powerful tool for enhancing productivity and creativity in professional work. These exercises demonstrate practical scenarios where generative AI can assist with resume writing, professional email drafting, and career skill development. Using prompt-based interactions, we can generate professional content, customize it for specific situations, and receive personalized recommendations for career growth.

44.3 Exercise 1: Writing and Refining Resumes

44.3.1 Purpose

In this exercise, we use generative AI to create a professional resume template, customize it according to a job description, and generate a cover letter tailored to the same position.

44.3.2 Prompt 1: Generate a Resume Template

Design a resume template suitable for a **Data Scientist** with good knowledge of **Python** and **SQL programming** and **7 years of experience**.

44.3.3 Prompt 2: Customize the Resume Based on a Job Description

Modify the above resume for the following job description:

- Lead and mentor a team of data scientists and analysts, fostering a collaborative and innovative environment.
- Develop and implement advanced machine learning models and statistical analyses to solve complex business problems.
- Collaborate with cross-functional teams to identify opportunities for leveraging data to drive business solutions.
- Design and execute experiments, **A/B tests** (controlled experiments that compare two versions to determine which performs better), and other quantitative analyses to inform strategic decisions.
- Communicate findings and insights to stakeholders through clear and compelling data visualizations and reports.
- Stay up-to-date with the latest advancements in data science and machine learning, and apply them to improve existing processes and solutions.
- Ensure the quality, integrity, and security of data used in analysis and modeling.

- Drive the development and maintenance of data infrastructure, tools, and processes to support scalable data science efforts.

44.3.4 Prompt 3: Generate a Cover Letter

Generate a cover letter for the above resume.

44.4 Exercise 2: Drafting Emails

44.4.1 Purpose

Professional email communication is an essential workplace skill. Generative AI can help draft clear, professional, and context-aware emails for different business scenarios, reducing the time required to compose them while maintaining an appropriate tone.

44.4.2 Prompt 1: Job Application Email

Write a professional email for **Jill**, who is applying for the position of **Lead Data Scientist**.

The email should:

- Include a brief introduction of the candidate.
- Mention the job role she is applying for.
- Summarize her relevant experience and skills.
- Express enthusiasm for the opportunity.
- Mention the attached resume.
- Offer to provide additional information if needed.

44.4.3 Prompt 2: Team Introduction Email

Write an email for **Jill**, who has recently joined an organization as the **Lead Data Scientist**.

The email should:

- Include a brief introduction of her background and experience.
- Express her excitement about joining the team.
- Describe her initial goals or areas of focus.
- Invite the team to connect with her.

Tone: Friendly and professional.

44.4.4 Prompt 3: Meeting Request Email

Generate an email to schedule a meeting with a potential client.

Assume:

- **Sender:** A business executive from **Xenkee**.
- **Recipient:** A key decision-maker at **Vinnel**.

Tone: Formal and professional.

44.5 Exercise 3: Enhancing Your Skills

44.5.1 Purpose

Generative AI can help identify skill gaps for a desired role and recommend practical methods, learning resources, certifications, and personalized learning paths to support professional growth and career advancement.

44.5.2 Prompt 1: Identify Skill Gaps and Development Methods

A candidate is experienced in **digital marketing** but needs to enhance their skills in:

- Market research
- Analytical thinking
- Leadership

Provide:

- A brief description of each skill.
- Suggested methods for the candidate to develop these skills.

44.5.3 Prompt 2: Skills Required for Career Advancement

A candidate is a **Human Resource Executive** with **3 years of experience**. She aims to be promoted to a **Human Resources Manager** role. Identify the skills she needs to enhance to achieve this goal.

44.5.4 Prompt 3: Personalized Learning Path

Suggest methods and personalized learning paths for the candidate to develop these skills.

44.6 Summary

In this topic, we explored how generative AI can support professional work by:

- Generating resume templates.
- Customizing resumes according to job descriptions.
- Creating cover letters.
- Drafting professional emails for different workplace scenarios.
- Identifying skill gaps.
- Recommending personalized learning paths for career development.

45 Human, AI, or Augmented Intelligence: Who Should Make the Decision?

45.1 Decision Making: Human vs. AI

A decision needs to be made. The question is: **Who should make it?** Should it be **a human, an Artificial Intelligence (AI) system, or a combination of both?**

We have discussed that humans can outperform AI at some tasks, while statistically, AI performs better for other tasks. Therefore, for a single decision, we need to determine who is most likely to make the best decision.

The answer is a fascinating combination of **performance curves** and **human cognitive bias**.

45.2 Fraud Detection Example

Consider a **fraud detection system**.

- The system generates alerts for potentially fraudulent transactions.
- Financial analysts review each alert.
- Thousands of events are generated every day.
- Analysts are overwhelmed because **90% of these alerts are false positives**.

An AI system can help reduce the analysts' workload. However, an important question remains:

- Which alerts should the AI handle?
- Which alerts should be processed by a skilled financial analyst?

To answer this, we examine the relationship between **confidence scores** and **decision accuracy**.

45.3 AI Confidence Score and Success Rate

Imagine a graph with two axes:

- **Y-axis:** Success rate (whether the prediction turns out to be correct).
- **X-axis:** Confidence score.

The confidence score represents how certain the prediction is.

- **0% confidence** means the prediction believes the alert is definitely **not** real (a false positive).
- **100% confidence** means the prediction is certain that the alert **is** real.

45.3.1 AI Performance Curve

A typical AI performance curve behaves as follows:

- Very **low confidence scores** (strong confidence that the alert is not real) have a **high success rate**.
- Very **high confidence scores** (strong confidence that the alert is real) also have a **high success rate**.
- Around the middle confidence levels, the success rate drops significantly.

When the AI has approximately **50% confidence**, it is effectively saying:

"I don't know."

In these uncertain situations, the AI becomes less accurate.

45.4 Human Performance Curve

Human performance curves are generally flatter than AI performance curves.

Humans are:

- Often **not quite as accurate** as a highly confident AI.
- Better at making decisions when the AI is uncertain.

At approximately **50% confidence**, humans are likely to make better decisions than AI.

45.4.1 Why Humans Perform Better in Uncertain Cases

When AI is highly confident:

- It is highly accurate.
- It often outperforms humans.
- It remains consistent.
- It does not lose focus or become distracted.

Humans, however, may lose consistency or attention.

When AI is uncertain:

- The cases are often statistically rare or highly complex.
- Humans can introduce additional information and context.
- We can:
 - Look up additional information.
 - Ask a colleague.
 - Apply reasoning beyond the available data.

Meanwhile, the AI continues using the same decision logic and information it was trained on.

Therefore:

- If the AI assigns a **very high** or **very low confidence score**, statistically it is more likely than a financial analyst to correctly determine whether the alert is real or a false positive.

- If the AI has **low confidence in its own prediction**, a human is often the better decision maker.

45.5 Augmented Intelligence

The choice does not have to be **AI or human**.

There is a third option:

Augmented Intelligence

Augmented Intelligence combines:

- Human decision making.
- AI recommendations.

Its performance curve lies between the human and AI performance curves.

For many predictions with **somewhat low** or **somewhat high confidence scores** (which represent a significant proportion of predictions), **Augmented Intelligence achieves the highest success rate**.

By combining the strengths of humans and AI, we can improve overall decision-making performance.

45.6 Human Cognitive Bias

For Augmented Intelligence to be most effective, we must account for **human cognitive bias**.

Humans are not always good at following recommendations objectively.

The way AI information is presented to human decision makers significantly influences how effectively that information is used.

45.7 Forced Display vs. Optional Display

Two common methods are used to present AI recommendations.

45.7.1 Forced Display

In a **Forced Display**, the AI recommendation is shown at the same time as the decision case.

For every fraud alert:

- The analyst sees the case.
- The analyst also immediately sees the AI recommendation.

This can lead to **Automation Bias**.

45.7.2 Automation Bias

Automation Bias is the tendency for humans to favor suggestions made by automated decision-making systems while ignoring contradictory information.

In practice, the human decision maker may think:

"The AI knows best."

As a result, they follow the AI recommendation instead of relying on their own judgment.

45.8 Optional Display

With an **Optional Display**, the AI recommendation is hidden unless the human explicitly requests it.

The process works as follows:

1. We first review the decision case independently.
2. We form our own initial judgment.
3. We then request the AI recommendation if needed.

This approach helps reduce Automation Bias because:

- We have time to think independently.
- Our first impression is formed before seeing the AI recommendation.
- We are less likely to be overly influenced by the AI.

45.9 Trust in AI Recommendations

Trust also plays an important role.

Sometimes AI recommendations include an **accuracy percentage**, indicating how likely the prediction is to be correct.

Interestingly, when this accuracy percentage is displayed, humans are generally **less likely to incorporate the AI recommendation into their decision**, regardless of the actual percentage shown.

In simple terms:

People do not like recommendations that openly admit they might be wrong.

45.10 Key Takeaways

- Whether a decision should be made by a **human**, an **AI system**, or **both together** can be determined objectively rather than subjectively.
- AI performs best when it has **very high** or **very low confidence** in its prediction.

- Humans perform better when AI is **uncertain**, especially in rare or complex situations where additional context is valuable.
- **Augmented Intelligence** combines human expertise with AI recommendations and often produces the best overall decision-making performance.
- The way AI recommendations are presented is important.
- **Forced Display** can create **Automation Bias**, causing humans to over-rely on AI.
- **Optional Display** helps reduce Automation Bias by allowing humans to make an independent initial judgment before viewing AI recommendations.
- Human trust in AI can also be affected by how confidence or accuracy information is presented.

By combining human intelligence with AI algorithms while minimizing cognitive bias, we can significantly improve decision-making outcomes. Together, humans and AI form a powerful team—we simply need to know who should make the decision in each situation.

46 Module Summary

At this point in the module, we know the following concepts and topics.

46.1 AI Agents

- Agents in AI are software programs that engage with surroundings, process data, and autonomously carry out tasks to meet human-defined goals.

46.2 Robotics and Cobots

- Robotics involves designing, constructing, and operating robots to perform tasks with or without humans.
- Cobots collaborate with humans using advanced sensors and AI, communicating and coordinating tasks that require teamwork.
- Robotic Process Automation (RPA) automates repetitive tasks and helps create, use, and control virtual robots.

46.3 AI in Business

AI helps businesses automate repetitive tasks and improve operational efficiency.

46.3.1 Business Applications of AI

- AI automates repetitive tasks for businesses, such as data entry, scheduling, and report generation, allowing employees to focus more on creative and strategic work.
- AI enables businesses to swiftly analyze large volumes of data by identifying patterns and trends, leading to fewer errors and better predictions.
- AI generates innovative ideas and overcomes creative blocks by helping businesses resonate with the target audience, leading to higher engagement and sales.

46.4 Evolution of Artificial Intelligence

- Early AI used rule-based systems but lacked creativity.
- Decades later, deep learning enabled AI to learn and adapt like the human brain.
- Generative Adversarial Networks (GANs) (a deep learning technique where two neural networks compete to generate realistic content) marked a breakthrough by generating high-quality images, music, design, and medicine.
- GANs advanced generative AI in fields such as music, design, and medicine.

46.5 Key Areas Where Generative AI Contributes

- **Content Generation:** Analyzes existing content and target audience data to create unique content.
- **Data Analysis:** Analyzes complex datasets to identify patterns and generate actionable reports.
- **Customer Service:** Handles repetitive tasks by providing fast and personalized responses.
- **Product Development:** Produces multiple design variations for faster and more innovative product development.

46.6 Adopting AI in Business

AI can be adopted in businesses to:

- Increase efficiency.
- Enhance decision-making.
- Improve customer service.
- Drive innovation.

46.6.1 General Steps to Adopt AI in Business

- **Define business goals:** Identify the problems and set clear goals.
- **Identify suitable use cases:** Identify areas where AI can be applied.
- **Data readiness:** Collect data, followed by cleaning and organizing it.
- **Build AI capabilities:** Train employees and develop the necessary setup.
- **Deploy AI solutions:** Integrate AI with existing systems.
- **Monitor and optimize AI systems:** Ensure smooth functioning.

46.7 AI for Organizational Productivity

- AI analyzes vast amounts of data at incredible speeds, revolutionizing how we operate, innovate, and succeed.
- AI tools help organizations enhance efficiency by automating employee engagement, recruitment, onboarding, and performance management.
- AI chatbots have transformed customer interactions by automating responses and ensuring consistency.
- Professionals from different domains leverage AI tools and platforms to enhance productivity and efficiency.

46.8 AI Tools and Platforms

Category	AI Tools and Platforms	Purpose
General AI Assistants	ChatGPT, Gemini	Provide versatile applications
Marketing and Content Creation	Copy.ai, Jasper, Synthesia	Aid marketers and content creators
Writing and Communication	Grammarly, QuillBot	Assist with writing and communication tasks
Language Learning and Translation	Duolingo, Google Translate, Babel	Enable language learning and translation
Customer Service	AI chatbots, Zendesk, LivePerson	Enhance customer service
Data Analysis and Visualization	Tableau, Power BI	Support data analysis and visualization
Software Development	GitHub Copilot	Help developers with code completion
Task Management	Todoist, Microsoft To Do, Evernote	Assist in task management

46.9 Building a Career in AI

To switch to a career in an AI-related field, we should:

- Identify our existing skills that can be leveraged.
- Learn core AI concepts.
- Apply our skills practically.
- Stay updated with the latest trends.
- Delve deeper into specialized areas of AI.

46.10 AI Career Opportunities

- **AI Ethicist:** Ensures that AI systems are developed and used ethically while considering social and moral implications.
- **AI Product Manager:** Oversees the development and implementation of AI products, ensuring they meet market needs.
- **AI Strategist:** Develops long-term strategies for implementing AI to achieve business goals.
- **AI Marketing Specialist:** Utilizes AI to analyze consumer data and optimize marketing strategies.

Module 4: Issues, Concerns, and Ethical Considerations

47 Ethical Considerations and Responsible Use of AI

47.1 Learning Objectives

After completing this topic, we will be able to:

- Explain the key ethical principles in Artificial Intelligence (AI).
- Develop strategies for the responsible implementation and use of AI.

47.2 Case Study: Amazon's AI Recruitment Tool

In **2014**, Amazon developed an AI tool to automate the shortlisting of job applicants. However, the system exhibited **gender bias** by penalizing female-associated terms because the training data predominantly consisted of male candidates.

Despite efforts to correct the bias, Amazon discontinued the tool in **2015** to uphold its commitment to diversity and equality.

This case study highlights the need for ethical considerations in AI development to ensure these powerful technologies benefit society while minimizing potential harm.

47.3 Importance of Ethical AI

AI has transformed industries by enhancing efficiency, innovation, and decision-making processes.

From healthcare and education to finance and entertainment, AI systems can:

- Analyze vast amounts of data.
- Identify patterns.
- Make predictions.
- Save lives.
- Improve services.
- Drive economic growth.

However, with great power comes great responsibility. Therefore, understanding the ethical principles and responsible use of AI is essential.

47.4 Data Privacy and Security

One of the foremost ethical concerns in AI is **data privacy and security**.

AI systems rely heavily on data to function effectively, and this data often includes sensitive personal information.

Therefore, it is imperative to implement robust data protection measures to prevent unauthorized access and data breaches.

Organizations should comply with data protection regulations such as:

- **General Data Protection Regulation (GDPR)**
- **California Consumer Privacy Act (CCPA)**

These regulations help ensure that:

- User consent is obtained.
- Data is anonymized wherever possible.
- Personal information is handled responsibly.

47.4.1 Real-World Example: Clearview AI

According to a [New York Times](#) article published in **2020**, the facial recognition company **Clearview AI** collected billions of images from social media platforms and other websites without users' consent.

The company provided facial recognition software to law enforcement agencies, raising serious concerns about privacy violations.

This incident sparked significant debate regarding:

- The ethical use of AI.
- Privacy protection.
- The need for stronger data privacy regulations.

47.5 Bias and Fairness

AI systems can unintentionally amplify existing biases present in their training data.

These biases may result in unfair treatment and discrimination in areas such as:

- Hiring
- Lending
- Law enforcement

To mitigate bias, developers should:

- Use diverse datasets.
- Implement fairness-aware algorithms.
- Continuously audit AI systems for biased outcomes.

Ensuring fairness in AI is not only a technical challenge but also a societal imperative.

47.6 Transparency and Accountability

Transparency in AI operations is essential for building trust.

Users should understand:

- How AI systems make decisions.
- What data AI systems use.
- How their personal information is handled.

Clear accountability is equally important.

If an AI system causes harm or makes an incorrect decision, it should be clear:

- Who is responsible.
- How the issue will be rectified.

Establishing accountability frameworks is crucial for the ethical deployment of AI.

47.7 Autonomous Systems and Human Oversight

As AI systems become more autonomous, human oversight remains essential.

Fully autonomous systems, especially in critical sectors such as:

- Healthcare
- Transportation
- Military

should include **human-in-the-loop mechanisms** (systems where humans remain involved in reviewing or approving important AI decisions).

These mechanisms ensure that:

- Human judgment remains part of decision-making.
- Ethical considerations are incorporated.
- Potentially dangerous or unethical outcomes are prevented.

47.8 Access and Equality

Access to AI technology and its benefits should be equitable.

If AI is available only to privileged groups, it may worsen existing inequalities.

Efforts should be made to democratize AI by ensuring it is accessible and beneficial to:

- All segments of society.
- Underserved communities.
- Marginalized communities.

This also includes investing in:

- Education.
- Infrastructure.

These investments help bridge the digital divide.

47.9 Ethical Considerations Across Different AI Applications

Different AI applications present unique ethical challenges.

47.9.1 Healthcare

In healthcare, AI should prioritize:

- Patient welfare.
- Patient confidentiality.

47.9.2 Law Enforcement

In law enforcement, AI applications such as:

- Surveillance
- Predictive policing (using AI to predict where crimes may occur or who may be involved)

raise significant concerns regarding:

- Privacy.
- Civil liberties.

Each application area requires tailored ethical guidelines to address its unique challenges and ensure responsible AI use.

47.10 Environmental Impact of AI

Another critical ethical consideration is AI's environmental impact.

Training and running large AI models consume substantial amounts of energy, contributing to carbon emissions.

Developers and organizations should adopt sustainable practices by:

- [Creating energy-efficient algorithms.](#)
- [Leveraging advancements in hardware.](#)
- [Optimizing algorithms for energy efficiency.](#)
- [Utilizing renewable energy sources for data centers.](#)

These practices help reduce AI's environmental footprint.

47.11 Continuous Monitoring and Improvement

Ethical AI is not a one-time achievement but a continuous process.

AI systems should be monitored regularly to evaluate their impacts.

Organizations should establish feedback mechanisms to gather input from:

- Users.
- Stakeholders.

Continuous improvement practices should be implemented to update AI systems based on:

- New ethical insights.
- Technological advancements.
- Societal needs.

47.12 Responsible AI Implementation Strategies

Responsible AI implementation involves:

- Using diverse datasets.
- Implementing fairness-aware algorithms.
- Continuously auditing AI systems for biased outcomes.
- Designing fully autonomous systems with human-in-the-loop mechanisms.
- Ensuring transparency in AI decision-making.
- Establishing clear accountability frameworks.
- Protecting user privacy and complying with regulations such as GDPR and CCPA.
- Ensuring AI accessibility and benefits for all segments of society.
- Optimizing algorithms for energy efficiency.
- Utilizing renewable energy sources for data centers.
- Continuously monitoring AI systems and incorporating user and stakeholder feedback.

47.13 Key Takeaways

The ethical considerations of AI include:

- Data privacy and security.
- Bias and fairness.
- Transparency and accountability.
- Autonomous systems and human oversight.
- Access and equality.
- Environmental sustainability.
- Continuous monitoring and improvement.

Responsible AI requires organizations and developers to adopt ethical practices that ensure AI systems are fair, transparent, secure, sustainable, accountable, and beneficial to society while minimizing potential harm.

48 Considerations Around Generative AI

48.1 Learning Objectives

After studying this topic, we will be able to:

- Identify the complexities of copyright and content ownership in generative AI.
- Explain the importance of privacy and confidentiality of data while using AI models.
- Explore the accuracy of generative AI outputs and understand the potential for hallucinations.
- Identify ethical considerations when developing and using generative AI technologies.

Generative AI is transforming industries, pushing the boundaries of creativity and innovation. However, as with any powerful tool, it comes with its own set of challenges and considerations.

48.2 Copyright and Content Ownership

One of the major considerations in generative AI is the question of intellectual property and ownership.

In 2018, a piece of artwork titled **Edmond de Bellamy**, created by an AI algorithm called **GAN (Generative Adversarial Network)**, was sold at an auction for **over \$400,000**.

This raised several important questions:

- Who owned the rights to the artwork?
- Were the rights owned by the developers who programmed the AI?
- Did the buyers of the artwork own the rights?

The sale sparked debates and discussions around intellectual property rights in AI-generated creations.

As AI systems become more creative, the question of ownership of AI-generated content and innovations becomes increasingly important.

To address these challenges, it is recommended to:

- Establish clear guidelines and policies regarding AI-generated content.
- Balance the interests of AI developers, users, and broader society.
- Develop appropriate regulations that govern AI while continuing to promote innovation.

Clarifying intellectual property rights related to AI-generated content is a complex and evolving issue.

It is also important to ensure that AI systems comply with existing laws and regulations, including:

- Data protection laws
- Consumer rights regulations

48.3 Privacy and Confidentiality

Another critical concern in generative AI is privacy and confidentiality.

Generative AI models often require large amounts of data for training, and this data can sometimes include sensitive or personal information.

Strong data protection measures should be implemented to ensure that the information used to train AI models does not violate individual privacy.

48.3.1 Private AI

To address privacy concerns, organizations are increasingly turning to **Private AI**.

Private AI refers to an AI environment created specifically for a particular organization and intended for exclusive use by that organization.

Private AI environments are designed to:

- Keep data confidential.
- Keep data secure.
- Prevent unauthorized access.
- Prevent misuse of organizational data.

48.4 Accuracy of AI Outputs and Hallucinations

Generative AI is incredibly powerful, but it still has limitations.

Sometimes AI produces outputs that are inaccurate or completely fabricated. This phenomenon is known as **hallucination** (AI-generated information that appears believable but is false or fabricated).

48.4.1 Real-World Example

An attorney used **ChatGPT** to craft a legal motion filled with phony judicial opinions and legal citations.

The attorney, unaware of ChatGPT's ability to fabricate cases, faced sanctions for submitting the document.

This example highlights the importance of critical thinking when using AI-generated content.

We should always:

- [Validate AI-generated information.](#)
- [Fact-check AI-generated content.](#)
- [Avoid spreading misinformation.](#)

48.4.2 Reducing AI Hallucinations

Developers and users should implement strategies to minimize the risk of hallucinations, including:

- Improving the quality of training data.
- Refining AI algorithms.
- Establishing robust validation processes.

Organizations can also use **private GPT models** to address hallucination concerns.

For example, **CustomGPT.ai** is an advanced AI platform that develops tailored GPT chatbots using business content, delivering accurate and reliable **GPT-4** responses based on that content without fabricating information.

By leveraging private models, businesses can reduce hallucinations by ensuring that the training data is highly relevant and accurate, leading to more reliable and context-specific responses.

48.5 Model Performance Considerations

Several factors influence the performance of generative AI models.

High-quality, diverse datasets are necessary for training effective models.

The quantity and quality of data play a significant role in model performance.

More complex models can achieve better performance, but they also:

- Require more computational resources.
- Are more difficult to interpret.

Another important consideration is scalability.

AI models should:

- Scale effectively across different applications and environments.
- Be robust to various inputs.
- Handle unexpected scenarios gracefully.

48.6 Ethical Considerations

Ethics is a fundamental consideration in generative AI.

From continuing biases to potential misuse, the ethical landscape of AI is complex.

48.6.1 Bias in AI Systems

AI systems can unintentionally continue biases from their training data, leading to unfair results.

Developers should prioritize:

- Fairness
- Accountability
- Transparency

This involves:

- Actively identifying biases.

- Fixing biases.
- Ensuring AI systems are used responsibly.

48.6.2 Deepfakes

Generative AI can be misused to create **deepfakes**.

Deepfakes are videos, images, or audio recordings that have been manipulated using generative AI models to make them appear as if they are real.

Deepfakes can be highly convincing and may be used for harmful purposes, including:

- [Creating fake news.](#)
- [Creating propaganda.](#)
- [Spreading disinformation.](#)
- [Harassing people.](#)
- [Blackmailing people.](#)
- [Committing fraud.](#)

Deepfakes can also be used to create forged identity documents or photos for fraudulent purposes, such as:

- [Opening bank accounts.](#)
- [Obtaining loans.](#)
- [Gaining access to secure locations.](#)

These activities can lead to financial fraud and other security breaches.

48.6.3 Using AI for Positive Social Impact

Ethics in AI extends beyond preventing harm.

It is also about fostering positive outcomes and ensuring that AI technologies benefit society as a whole.

This includes using AI to:

- Address social challenges.
- Improve access to resources.
- Enhance quality of life.

Developers and organizations should create AI systems that support initiatives promoting:

- Social justice.
- Environmental sustainability.
- Human well-being.

48.7 Key Takeaways

In this topic, we learned that:

- Copyright and ownership of AI-generated content present complex legal challenges that continue to evolve.
- Developers and users should become familiar with legal frameworks and stay updated on changing laws related to AI-generated content.
- Privacy and confidentiality of user data can be maintained through Private AI models along with appropriate technical and legal measures.
- Critical thinking, validation, and fact-checking are essential for ensuring the accuracy of AI-generated outputs and reducing the impact of hallucinations.
- Organizations can reduce hallucinations by using private GPT models with highly relevant and accurate training data.
- High-quality datasets, model complexity, robustness, and scalability significantly influence AI model performance.
- Ethical AI development requires fairness, accountability, transparency, and responsible use.
- Developers and organizations should ensure that AI technologies create positive social impact while minimizing risks such as bias, deepfakes, fraud, and misinformation.

49 Understanding Hallucinations in Large Language Models (LLMs)

49.1 Introduction

Consider the following three statements:

- The distance from the Earth to the Moon is **54 million kilometers**.
- Before working at **IBM**, we worked at a major Australian airline.
- The **James Webb Space Telescope** took the very first pictures of an exoplanet outside our solar system.

At first glance, these statements appear to be related to **space and aviation**, but their real connection is different.

All three statements are examples of **hallucinations** generated by a **Large Language Model (LLM)**, such as **ChatGPT** or **Bing Chat**.

The correct facts are:

- **54 million kilometers** is approximately the distance from the **Earth to Mars**, not the Moon.
- It was the speaker's **brother**, not the speaker, who worked at the Australian airline.
- The first image of an **exoplanet** was actually captured in **2004**, not by the James Webb Space Telescope.

Although Large Language Models can generate fluent and coherent text across many topics, they are also capable of producing **plausible-sounding but incorrect information**, commonly known as **hallucinations**.

In this topic, we discuss:

- What hallucinations are
- Why hallucinations happen
- Strategies to minimize hallucinations when using LLMs

49.2 What Are Hallucinations?

Hallucinations are outputs generated by **Large Language Models (LLMs)** that **deviate from factual information or contextual logic**.

They can range from:

- Minor inconsistencies
- Completely fabricated statements
- Contradictory information

Hallucinations can also be categorized according to different levels of granularity.

49.3 Types of Hallucinations

49.3.1 Sentence Contradictions

This is the simplest form of hallucination.

The LLM generates a sentence that contradicts one of its previous statements.

Example:

- The sky is blue today.
- The sky is green today.

These two statements directly contradict each other.

49.3.2 Prompt Contradictions

A prompt contradiction occurs when the generated response contradicts the user's original prompt.

Example:

If we ask an LLM:

Write a positive review of a restaurant.

And it responds:

The food was terrible, and the service was rude.

The response directly contradicts the requested positive review.

49.3.3 Factual Contradictions

These hallucinations involve factual errors where the model presents incorrect information as fact.

Example:

Barack Obama was the first President of the United States.

This statement is factually incorrect.

49.3.4 Nonsensical or Irrelevant Information

Sometimes an LLM includes information that has no logical relevance to the discussion.

Example:

The capital of France is Paris. Paris is also the name of a famous singer.

The additional statement is unrelated to the topic and serves no meaningful purpose.

49.4 Why Do Hallucinations Happen?

The exact reason is difficult to explain because the internal decision-making process of Large Language Models is still largely a **black box** (a system whose internal workings are not fully understood), even to the engineers who build them.

However, several common causes are known.

49.4.1 Data Quality

LLMs are trained on extremely large collections of text that may contain:

- Noise
- Errors
- Biases
- Inconsistencies

For example, some LLMs have been trained using data scraped from:

- Wikipedia
- Reddit

Even if these sources were completely reliable, they still cannot cover every possible topic or domain.

As a result, LLMs often:

- Generalize from existing information.
- Cannot independently verify whether information is accurate or relevant.
- Occasionally generate incorrect facts.

As reasoning capabilities improve, hallucinations generally become less frequent.

49.4.2 Generation Methods

LLMs generate text using different methods and optimization objectives, including:

- Beam Search
- Sampling
- Maximum Likelihood Estimation
- Reinforcement Learning

These methods introduce trade-offs between:

- Fluency and diversity
- Coherence and creativity
- Accuracy and novelty

For example:

- **Beam Search** may favor high-probability but generic words over lower-probability, more specific words.

These trade-offs can contribute to hallucinations.

49.4.3 Input Context

Input context refers to the information provided to the model in the prompt.

Clear context helps the model generate more accurate and relevant responses.

However, unclear, inconsistent, or contradictory context can confuse the model.

Example:

If we ask:

Can cats speak English?

The expected answer is:

No.

However, if the conversation is about the **Garfield** cartoon strip, where Garfield speaks English, then the correct contextual answer becomes:

Yes.

Garfield would probably also ask for another serving of lasagna.

This example demonstrates how important context is.

Similarly, if we do not specify whether we want:

- An academic essay
- Creative writing
- Technical documentation

the model cannot reliably determine the appropriate style or context.

49.5 Strategies to Reduce Hallucinations

49.5.1 Provide Clear and Specific Prompts

The more precise and detailed the prompt, the more likely the LLM is to generate relevant and accurate responses.

Instead of asking:

What happened in World War II?

Ask something more specific, such as:

Can you summarize the major events of World War II, including the key countries involved and the primary causes of the conflict?

A detailed prompt gives the model a much clearer understanding of what information is expected.

49.5.2 Use Active Mitigation Strategies

Many LLMs provide settings that influence how responses are generated.

One important parameter is **Temperature** (which controls the randomness and creativity of the generated text).

- **Lower Temperature**
 - Produces more conservative responses.
 - Produces more focused outputs.
 - Generally reduces hallucinations.
- **Higher Temperature**
 - Produces more diverse responses.
 - Encourages creativity.

- Increases the likelihood of hallucinations.

Choosing the appropriate temperature depends on whether accuracy or creativity is the primary goal.

49.5.3 Use Multi-Shot Prompting

Multi-shot prompting provides the LLM with multiple examples of the desired output format before asking the actual question.

Unlike single-shot prompting, which supplies only one prompt, multi-shot prompting helps the model better understand:

- Expected output format
- Desired writing style
- Context
- Patterns to follow

This technique is particularly useful for tasks such as:

- Generating code
- Writing poetry
- Producing responses in a specific style
- Answering questions in a predefined format

Providing several examples helps the model recognize the intended pattern more effectively.

49.6 Key Takeaways

- Hallucinations are responses generated by LLMs that deviate from factual information or logical context.
- Hallucinations may appear convincing even when they are completely incorrect.
- Common types include:
 - Sentence contradictions
 - Prompt contradictions
 - Factual contradictions
 - Nonsensical or irrelevant information
- Major causes include:
 - Poor or incomplete training data
 - Limitations of text generation methods
 - Insufficient or unclear input context
- Hallucinations can be reduced by:
 - Writing clear and specific prompts
 - Adjusting generation settings such as temperature
 - Using multi-shot prompting to provide examples and context

Understanding why hallucinations occur and applying these strategies enables us to harness the capabilities of Large Language Models better while reducing the likelihood of inaccurate or fabricated responses.

50 Key Players Around AI Ethics

50.1 Learning Objectives

After studying this topic, we will be able to:

- Discuss the approaches and principles adopted by various organizations to ensure ethical AI adoption.
- Explain IBM's Pillars of Trust.
- Understand Microsoft's approach to ethical AI.
- Describe Google's AI Principles and governance approach.

50.2 Overview


Ensuring AI ethics is a complex issue that organizations are addressing through different approaches. Organizations such as IBM, Microsoft, and Google have developed frameworks, principles, and tools to address ethical concerns in AI development and deployment. Each organization uses its own methods and initiatives to ensure that AI systems are developed and used responsibly.

50.3 IBM's Pillars of Trust

IBM has made significant progress in promoting AI ethics through its **Pillars of Trust**. These pillars represent focus areas that help organizations build and use AI ethically.

The five pillars are:

- Explainability
- Fairness
- Robustness
- Transparency
- Privacy



50.3.1 Explainability

AI is **explainable** when it can show how and why it arrived at a particular outcome or recommendation.

50.3.2 Fairness

AI is **fair** when it treats individuals or groups equitably.

50.3.3 Robustness

AI is **robust** when it can effectively handle exceptional conditions such as abnormal input or adversarial attacks (malicious inputs intentionally designed to deceive AI systems).

50.3.4 Transparency


AI is **transparent** when appropriate information is shared with humans about how the AI system was designed and developed.

50.3.5 Privacy

Because AI ingests (collects and processes) large amounts of data, AI systems must prioritize and safeguard human privacy and data rights.

50.4 IBM's AI Ethics Toolkits

IBM has developed several toolkits to support its Pillars of Trust.



Toolkit	Purpose
AI Explainability 360	Helps understand and interpret AI models to ensure transparency.
AI Fairness 360	Provides metrics and algorithms to detect and mitigate bias in AI systems.
Adversarial Robustness 360	Enhances AI system security against adversarial attacks.
AI FactSheets 360	Promotes transparency by documenting AI models, processes, and performance.

These pillars and toolkits form the foundation of IBM's approach to creating trustworthy AI that aligns with ethical principles and societal values.

50.5 Microsoft's Approach to Ethical AI

Microsoft emphasizes that ethical AI requires the involvement of multiple stakeholders, including:

- Designers
- Developers
- Data providers
- Regulatory bodies

Microsoft follows several practices to ensure responsible AI development and deployment.

50.5.1 Human-in-the-Loop Systems

Microsoft deploys **human-in-the-loop systems** (AI systems that include human oversight and intervention when needed) to ensure humans can intervene whenever necessary.

50.5.2 Continuous Monitoring and Improvement

Microsoft continuously evaluates AI systems to identify and address ethical concerns throughout their lifecycle.

50.5.3 Regular Audits

Microsoft conducts regular audits to ensure compliance with ethical standards and regulations.

50.5.4 Algorithmic Impact Assessments (AIA)

Microsoft uses **Algorithmic Impact Assessments (AIA)** to identify and mitigate biases before AI systems are deployed.

50.5.5 Internal Review Bodies

Microsoft has established bodies such as **Microsoft's AETHER Committee** (AI, Ethics, and Effects in Engineering and Research), whose primary responsibility is to guide on issues, technologies, processes, and best practices for responsible AI.

50.5.6 Responsible AI Standard

Microsoft has developed a **Responsible AI Standard** that outlines accountability measures, including:

- Transparency

- Fairness
- Human oversight

50.6 Google's AI Principles

Google follows a set of AI principles for developing and deploying AI responsibly.

50.6.1 Be Socially Beneficial

AI should benefit society and remain aligned with widely accepted principles.

50.6.2 Avoid Creating or Reinforcing Bias

AI systems should ensure fairness and avoid unjust impacts.

50.6.3 Build and Test AI for Safety

AI should be built and tested for safety while taking appropriate steps to mitigate risks.

50.6.4 Be Accountable to People

AI systems should allow appropriate human direction and control.

50.6.5 Incorporate Privacy Design Principles

User privacy and data should be protected throughout AI development and deployment.

50.6.6 Uphold High Standards of Scientific Excellence

AI research should maintain scientific rigor and integrity.

50.6.7 Make AI Available Only for Appropriate Uses

AI should only be made available for applications that are consistent with Google's AI principles.

50.7 Google's AI Governance

Google's AI governance includes a review process that assesses AI projects for ethical risks and compliance with its AI principles.

This review process involves:

- A multidisciplinary team of experts.
- Evaluation of AI projects at multiple stages.
- Assessment of ethical risks before deployment.

50.8 Industry Collaboration for Ethical AI

In addition to following internal ethical principles, organizations collaborate with external stakeholders to promote responsible AI.

50.8.1 AI Alliance

The **AI Alliance**, co-launched by **IBM** and **Meta**, together with leading organizations from various sectors, is dedicated to fostering an open community that accelerates responsible AI innovation.

The initiative prioritizes:

- Scientific rigor
- Trust
- Safety
- Security
- Diversity
- Economic competitiveness

By bringing together a broad range of organizations, the AI Alliance helps ensure that AI education, research, development, deployment, and governance are conducted responsibly and inclusively, benefiting everyone involved in the AI ecosystem.

50.9 Consulting Firms Supporting Ethical AI

Several consulting firms help organizations develop and implement ethical AI strategies.

50.9.1 Deloitte

Deloitte works with clients and external experts to develop ethical AI strategies.

Its services include:

- Ethical AI consulting
- AI deployment guidance
- Workshops on AI ethics
- Resources to help organizations understand and address ethical challenges during AI deployment

50.9.2 PwC

PwC offers AI ethics training programs for employees.

These programs focus on:

- Educating the workforce about the ethical implications of AI.
- Teaching how to implement ethical practices in AI projects.

50.10 Industry-Wide Ethical AI Initiatives

Organizations also participate in industry-wide collaborations to establish ethical AI standards.

50.10.1 Adobe

Adobe participates in initiatives such as the **Partnership on AI**, collaborating with other organizations to develop guidelines and best practices for responsible AI use.

50.11 Key Takeaways

- IBM's **Pillars of Trust** focus on **Explainability, Fairness, Robustness, Transparency, and Privacy** to promote trustworthy AI.
- IBM supports these pillars with toolkits such as **AI Explainability 360, AI Fairness 360, Adversarial Robustness 360, AI FactSheets 360, and AI Privacy 360**.
- Microsoft's ethical AI strategy includes **human-in-the-loop systems, continuous monitoring, regular audits, Algorithmic Impact Assessments (AIA), internal review bodies**, and a **Responsible AI Standard**.
- Google's AI Principles emphasize:
 - Being socially beneficial.
 - Avoiding bias.
 - Building and testing AI for safety.
 - Accountability to people.
 - Privacy by design.
 - Scientific excellence.
 - Restricting AI use to applications aligned with its principles.
- Google evaluates AI projects through a multidisciplinary governance review process.
- The **AI Alliance**, co-founded by **IBM** and **Meta**, promotes responsible AI innovation through collaboration.
- Consulting firms such as **Deloitte** and **PwC** help organizations implement ethical AI through consulting, workshops, and employee training.
- Organizations such as **Adobe** contribute to industry-wide initiatives like the **Partnership on AI** to establish responsible AI standards and best practices.

51 AI Governance

Artificial intelligence (AI) is growing at an exponential pace. Every day, we hear about new AI use cases, applications, and capabilities that were unimaginable just a few years ago.

At the same time, we are also hearing about AI systems deployed into production that are not delivering the expected outcomes. Examples include:

- Chatbots misdirecting customers and employees into making incorrect decisions.
- Chatbots hallucinating responses for customers.
- AI models generating biased outcomes.

There is no denying the enormous potential of artificial intelligence. However, premature deployment and adoption of AI systems can expose organizations to significant reputational and financial risks.

This is why **Artificial Intelligence Governance** has become one of the most relevant and important topics today.

51.1 What Is AI Governance?

AI governance refers to a set of rules, standards, and processes established to ensure the responsible and ethical development and deployment of artificial intelligence systems.

We can think of AI governance as a set of guardrails that help ensure the ethical use of AI systems by minimizing risks while maximizing their potential benefits.

51.2 Benefits of Artificial Intelligence

Artificial intelligence provides several important benefits, including:

- Reduced costs
- Improved efficiency
- Automation of repetitive and manual tasks

Although these advantages make AI one of today's most important and sought-after technologies, the associated risks make AI governance an even more critical topic.

51.3 Understanding an AI System

An AI system is designed to:

- Accept inputs.
- Produce outputs that mimic, augment, or assist human decision-making.

At the heart of every AI system is the **AI model**.

The primary objective of an AI model is to analyze the input and generate an output similar to what a human would normally produce.

51.3.1 The Role of Data

To mimic, augment, or assist human decision-making, an AI model must be supplied with **human-generated data**.

This data can exist in multiple formats, including:

- Structured data (columns and values)
- Semi-structured data (such as XML files)
- Unstructured data, including:
 - PDF documents

- Text documents
- Audio files
- Video files

The AI model can be thought of as highly engineered code that uses complex mathematical algorithms. It analyzes the data, identifies patterns, and learns how to reproduce the expected human behavior.

51.4 Risks Associated with AI Systems

Although AI provides tremendous benefits, several risks arise from the way AI systems are built and trained.

51.4.1 Bias in AI Models

Since AI models learn from human-generated data, they also learn from human biases.

Humans are not devoid of bias. We naturally possess several cognitive biases.

For example, some people may place undue importance on certain factors while making decisions while completely ignoring other relevant factors.

These biases become embedded in the data.

The biases are usually **latent** (hidden and not directly visible). We cannot simply inspect the data and identify them. However, highly engineered mathematical algorithms can detect these hidden patterns.

In the worst-case scenario, AI models reproduce these biases in their outputs, making AI systems susceptible to biased outcomes.

51.4.2 Privacy and Copyright Infringement

AI systems rely heavily on data.

Data is:

- Sent as input.
- Used for training the model.

Without proper oversight, this data may contain:

- Private information
- Sensitive information

If appropriate controls are not in place, this information can become embedded in the model and later appear in the model's output, resulting in privacy infringement.

Similarly, unstructured data may contain copyrighted material, which can also be reproduced in AI-generated outputs, leading to copyright infringement.

51.4.3 Lack of Transparency and Trust

Some AI models are **black box models** (models whose internal decision-making process is difficult or impossible to interpret).

Black box models are often preferred over **glass box models** (models whose internal logic can be understood and explained) because they typically produce more accurate results.

However, using black box models introduces a significant challenge.

The people developing these models often have little control over or understanding of the model's internal decision-making process.

As a result, when asked why a model made a particular decision, they may be unable to provide a clear explanation.

This creates a lack of transparency.

Without transparency, organizations face a lack of trust because users may hesitate to trust a system that cannot explain how or why it reached its conclusions.

51.4.4 Model Deterioration and Continuous Monitoring

AI models are not systems that can be created once and expected to generate high-quality results indefinitely.

Over time, models can deteriorate because incoming data may become significantly different from the data used during training.

As performance degrades, continuous monitoring becomes essential.

Without ongoing monitoring, AI models may no longer produce consistently accurate or high-quality outcomes.

51.5 AI Regulations and Guidelines

Because of these risks, organizations around the world are developing regulations and guidelines for managing AI systems.

51.5.1 AI Guidelines

One example is the **NIST AI Risk Management Framework**.

These guidelines provide recommendations for the responsible deployment and management of AI systems.

51.5.2 AI Regulations

One well-known regulation is the **EU AI Act**.

Unlike guidelines, regulations are legally enforceable.

Organizations that fail to comply with regulatory requirements may face:

- Financial penalties
- Reputational damage
- Ethical challenges

Non-compliance can expose organizations to significant legal, financial, and reputational risks.

51.6 Why AI Governance Is Essential

Several factors make AI governance critically important, including:

- Bias in AI models
- Privacy infringement
- Copyright infringement
- Lack of transparency
- Lack of trust
- Lack of continuous monitoring
- Increasing AI regulations and compliance requirements

Proper governance helps organizations manage these risks while ensuring that AI systems are developed and deployed responsibly.

51.7 Key Takeaway

The promise of artificial intelligence is undeniable, but its risks are equally real.

To fully realize the potential of AI, organizations must ensure that their AI systems are properly governed through well-defined rules, standards, processes, continuous monitoring, and compliance with relevant regulations.

52 AI Ethics in Enterprise AI Systems

52.1 Introduction

AI ethics is an important consideration for every organization developing or using artificial intelligence. We need to determine whether our AI solutions are at risk of crossing ethical boundaries and establish clear principles to ensure AI is developed and used responsibly.

52.2 Ethical Guidelines for AI Systems

Before creating or interacting with AI systems, we should establish a set of ethical guidelines or rules to follow.

- **Artificial intelligence is meant to augment human intelligence.** It is designed to support and enhance human capabilities, not replace humans.
- **Data and insights belong to their creator.** If we use customers' data for any purpose, it remains their data, not ours.
- **Solutions must be transparent and explainable.** We need visibility into:
 - Who is training the AI system.
 - What data is being used to train the system.
 - How the training data influences the algorithm's recommendations to the end user.

52.3 Evaluating Whether AI Follows Ethical Guidelines

Once we establish ethical rules, we need a way to determine whether our AI system follows or violates them.

One useful design thinking activity is **dichotomy mapping**.

52.3.1 Dichotomy Mapping

Dichotomy mapping involves:

- Listing all the features of the AI solution.
- Listing the intended benefits of each feature.
- Evaluating whether any feature could potentially cause harm.

52.3.2 Example: Hotel Recommendation System

Consider a hotel that uses an AI recommendation system to personalize a guest's experience.

Possible recommendations include:

- Determining which room a guest receives.
- Leaving personalized treats in the guest's room.
- Assigning a guest going skydiving to a room on a higher floor.

These features are intended to improve the customer experience.

However, we must also evaluate potential ethical concerns, such as:

- Is customer data being sold to advertisers?
- Is the customer data secure?
- Can differently abled users access and use the interface?
- Are differently abled users included fairly in the algorithm?

52.4 Addressing Identified Ethical Issues

After identifying potential issues, the next step is to fix them.

52.4.1 Implement Guard Rails

Guard rails are rules that every AI system must follow.

For example:

- We do not sell customer data to advertisers.

This serves as an ethical guard rail for the AI system.

52.4.2 Use Diverse Training Data

The data used to train an AI system should be diverse.

If we do not use a diverse dataset, the AI system will not be able to accommodate all users effectively, resulting in unfair or biased outcomes.

52.4.3 Use Open-Source Fairness Tools

We can also use open-source tools to evaluate and improve AI fairness.

One example is **IBM AI Fairness 360**.

This tool helps us:

- Detect bias in machine learning models.
- Mitigate (reduce) bias in machine learning models.

Other tools may help with:

- Adhering to privacy regulations.
- Detecting uncertainty in AI models.

52.5 Key Takeaways

Responsible AI development requires more than building accurate models. We should:

- Establish clear ethical guidelines.
- Evaluate AI systems using methods such as dichotomy mapping.
- Identify potential ethical risks.
- Implement guard rails to prevent misuse.
- Train AI systems using diverse datasets.
- Use fairness and privacy tools to detect and mitigate bias.

AI is everyone's responsibility. We must ensure that the AI we create and use is:

- Safe
- Secure
- Transparent
- Explainable
- Built by humans with humans in mind

53 Lesson Summary

We have completed this lesson. At this point, we understand the ethical considerations, responsible use, governance, and risks associated with Artificial Intelligence (AI), including Generative AI.

53.1 Ethical Considerations and Responsible Use of AI

Ethical considerations and responsible use of AI refer to the principles, practices, and guidelines aimed at ensuring that artificial intelligence technologies are developed, deployed, and used in ways that are beneficial, fair, and respectful of human rights and societal values.

The ethical considerations of AI include:

- Data privacy and security
- Bias and fairness
- Transparency and accountability
- Autonomous systems and human oversight
- Access and equality

AI can be used responsibly by:

- Using diverse datasets.
- Designing fully autonomous systems that include **human-in-the-loop** mechanisms (systems where humans remain involved in reviewing, supervising, or making critical decisions).
- Ensuring AI's accessibility and benefits to all segments of society.
- Optimizing algorithms for energy efficiency.
- Utilizing renewable energy sources for data centers.

53.2 Ethical Considerations for Generative AI

Generative AI involves the use of algorithms to create new content, such as:

- Text
- Images
- Audio
- Video

This technology presents unique opportunities as well as challenges.

Important considerations when using Generative AI include:

- **Copyright and ownership:** It is important for developers and users to familiarize themselves with legal frameworks and stay updated on evolving laws.
- **Privacy and confidentiality of user data:** It is beneficial to use private AI models and a combination of technical and legal measures.
- **Deepfakes and misinformation:** It is essential to think critically to ensure the accuracy of Generative AI outputs and to mitigate AI flaws such as **hallucinations** (AI-generated responses that sound convincing but contain incorrect or fabricated information).
- **Ethical considerations:** Developers should prioritize fairness, accountability, and social benefits throughout AI development.

53.3 Perspectives of Major Organizations on AI Ethics

Examining the perspectives of key organizations helps us gain a comprehensive understanding of different approaches to AI ethics, the challenges they address, and the strategies they employ to promote the responsible use of AI technologies.

53.3.1 IBM

IBM's pillars of trust are central to its approach to ensuring that AI systems are:

- Explainable
- Fair
- Robust
- Transparent
- Privacy-preserving

53.3.2 Microsoft

Microsoft's approach to ethical AI includes:

- Human-in-the-loop systems
- Continuous monitoring and improvement
- Regular audits using the **Artificial Intelligence Act (AIA)**
- Internal review bodies
- Responsible AI standards

53.3.3 Google

Google's AI principles focus on:

- Benefiting society
- Ensuring fairness and accountability
- Upholding high standards of scientific excellence

53.4 Risks of Premature AI Adoption

Premature deployment and adoption of AI systems could put companies at significant risk of:

- Reputational loss
- Financial loss

53.5 AI Governance

AI governance refers to a set of rules, standards, and processes established to ensure the responsible and ethical development and deployment of artificial intelligence systems.

Proper AI governance helps organizations maximize AI's benefits while minimizing associated risks.

53.6 Risks Associated with AI

53.6.1 Bias in AI Models

- Human-generated data can contain hidden biases.
- AI models can learn and reflect these biases in their outcomes.

53.6.2 Privacy and Copyright Infringement

- Sensitive data can be inadvertently included in AI models.
- Unstructured data may contain copyrighted material.

53.6.3 Lack of Transparency

- **Black box models** (AI models whose internal decision-making process is difficult or impossible to interpret) can provide high accuracy but lack explainability.
- A lack of transparency can erode trust in AI systems.

53.6.4 Model Deterioration

- AI models require continuous monitoring and updates.
- Incoming data may differ from training data, causing performance degradation.
- Continuous monitoring ensures AI models produce consistent, high-quality outcomes and prevents deterioration in model performance over time.

53.7 AI Regulations and Guidelines

Organizations worldwide are creating regulations and guidelines for AI.

Examples include:

- **NIST AI Risk Management Framework**
- **EU AI Act**

These regulations can penalize non-compliant companies, emphasizing the importance of adhering to ethical guidelines.

53.8 Key Takeaways

- Ethical AI focuses on fairness, transparency, accountability, privacy, security, and equal access.
- Responsible AI development requires diverse datasets, human oversight, energy-efficient algorithms, and inclusive accessibility.
- Generative AI introduces additional challenges related to copyright, privacy, misinformation, hallucinations, and ethical development.
- IBM, Microsoft, and Google each have their own frameworks and principles for promoting trustworthy and responsible AI.
- Premature AI deployment can lead to reputational and financial risks.
- AI governance provides the rules and processes needed for responsible AI development and deployment.
- Major AI risks include bias, privacy and copyright infringement, lack of transparency, and model deterioration.
- Continuous monitoring and compliance with regulations such as the **NIST AI Risk Management Framework** and the **EU AI Act** are essential for maintaining trustworthy AI systems.
- AI offers significant benefits but also presents real risks. Proper governance is crucial to fully realize AI's potential while mitigating associated risks.

54 Project: Transforming Organizational Functions with Generative AI

54.1 Learning Objectives

After completing these exercises, we will be able to:

- Apply generative AI within various organizational functions.
- Use generative AI to enhance customer experience through customer segmentation and personalized recommendations.
- Use generative AI for supplier performance analysis and demand forecasting.
- Utilize generative AI to improve customer support through chatbots.

54.2 Introduction

GenKart is a mid-sized e-commerce company that offers electronics, fashion, home goods, and sports equipment. To remain competitive and meet evolving customer expectations, the company aims to transform its business processes using generative AI.

This topic demonstrates how generative AI can enhance customer experience, optimize operations, improve decision-making, and drive business growth across different organizational functions.

54.3 Transforming Business Functions

54.3.1 Purpose

In this exercise, we explore how generative AI can transform different business functions of an e-commerce company to enhance customer experience, optimize operations, and support business growth.

54.3.2 Task

Log in to ChatGPT and use the following prompt.

Prompt Instructions

Consider a mid-sized e-commerce company offering electronics, fashion, home goods, and sports equipment. The company faces increasing competition and evolving customer expectations. To stay ahead, it is exploring generative AI to transform its functions, aiming to enhance customer experience, optimize operations, and drive growth. Generate the output in text form.

Prompt

How can generative AI be used to transform different functions and aspects in this e-commerce company?

54.4 Personalizing Customer Experience

54.4.1 Purpose

In this exercise, we demonstrate how an e-commerce company can use generative AI to personalize customer experiences through customer segmentation and personalized product recommendations, ultimately improving customer engagement and sales.

54.4.2 Customer Segmentation

Generative AI can segment customers by analyzing information such as:

- Browsing behavior
- Purchase history
- Demographic information
- Pages visited
- Time spent on each page
- Products viewed

54.4.3 Prompt 1: Customer Segmentation

Use the following prompt.

Based on the following data, segment customers into groups such as "frequent buyers," "seasonal shoppers," and "bargain hunters."

Customer A visits electronics and home goods pages frequently.

Customer B buys sports equipment regularly, with an average spend of \$200 per purchase.

Customer C is a 30-year-old female from New York with a high income.

Customer D makes 5 purchases per month, primarily in fashion and accessories.

Customer E makes large purchases during holiday sales but is inactive during other times.

Customer F has been shopping on the platform for over 3 years, primarily in electronics.

Customer G makes 2-3 purchases per month, primarily in fashion and sports equipment.

54.4.4 Prompt 2: Personalized Product Recommendations

Use the following prompt.

Generate product recommendations for each customer segment.

54.5 Supplier Performance Analysis

54.5.1 Purpose

In this exercise, we demonstrate how generative AI can streamline and optimize operational processes by evaluating supplier performance and recommending the most suitable suppliers based on delivery time, quality, cost, and reliability.

54.5.2 Prompt

Use the following prompt.

Based on the following data about a product's suppliers, evaluate and compare supplier performance based on average delivery time, quality compliance, and cost per unit.

Supplier M

- Average delivery time: 5 days
- Quality: 98% of products meet quality standards
- Cost: \$15 per unit
- Reliability: 95% on-time delivery rate

Supplier N

- Average delivery time: 3 days
- Quality: 90% of products meet quality standards
- Cost: \$10 per unit
- Reliability: 90% on-time delivery rate

Supplier O

- Average delivery time: 7 days
- Quality: 95% of products meet quality standards
- Cost: \$20 per unit
- Reliability: 98% on-time delivery rate

Supplier P

- Average delivery time: 4 days
- Quality: 97% of products meet quality standards

- Cost: \$17 per unit
- Reliability: 92% on-time delivery rate

54.6 Demand Prediction

54.6.1 Purpose

In this exercise, we use generative AI to analyze historical sales data, seasonal trends, and customer feedback to predict future product demand and support inventory planning and business decision-making.

54.6.2 Prompt

Use the following prompt.

Analyze the historical sales data, seasonal trends, and customer reviews to predict future demand for products.

54.6.3 Historical Sales Data

Product A

Monthly sales (units sold) over the past year:

[120, 150, 180, 160, 200, 220, 210, 230, 240, 250, 260, 280]

Product B

Monthly sales (units sold) over the past year:

[80, 90, 85, 100, 110, 120, 130, 125, 140, 150, 160, 170]

Product C

Monthly sales (units sold) over the past year:

[50, 55, 60, 65, 70, 80, 75, 85, 90, 95, 100, 110]

Seasonal Trends

Product A

High sales in November and December due to the holiday season.

Product B

Increased sales in spring and summer (April to August).

Product C

High sales in August and September.

54.6.4 Customer Feedback and Reviews

Product A

Positive feedback: 95% of reviews are 4 stars or higher.

Common feedback: Requests for more color variations.

Product B

Positive feedback: 90% of reviews are 4 stars or higher.

Common feedback: Requests for additional sizes.

Product C

Positive feedback: 85% of reviews are 4 stars or higher.

Common feedback: Requests for more durable materials.

54.7 Improving Customer Support

54.7.1 Purpose

In this exercise, we demonstrate how an AI-powered chatbot can improve customer support by assisting customers with product recommendations, answering product-related questions, and referring users to customer care when sufficient information is unavailable.

54.7.2 Task

Log in to ChatGPT and use the following prompt.

Prompt Instructions

Imagine you are a Shopping AI assistant for an e-commerce portal, GenKart. You help users find products, compare prices, and make purchasing decisions. Consider the data of the laptops offered by the portal and respond to the user's questions. If you do not have enough information to answer the user's questions, request them to contact the customer care team.

54.7.3 Prompt 1

I'm looking for a new laptop, but I have some very specific requirements for the software I need to use. Can you help?

54.7.4 Prompt 2

The software I need to use runs on Intel Core i9 processor. Does GenKart offer a laptop with this processor?

54.7.5 Prompt 3

Will GenKart offer any laptop with Intel Core i9 processor in the near future?

54.8 Summary

In this topic, we explored how generative AI can transform organizational functions by:

- Improving business operations.
- Personalizing customer experiences.
- Analyzing supplier performance.
- Forecasting product demand.
- Enhancing customer support through AI-powered chatbots.

These practical applications demonstrate how generative AI can help organizations make data-driven decisions, optimize workflows, and deliver better customer experiences.

55 Course Summary and Key Takeaways

Now that we thoroughly understand AI and its transformative potential, we are well-equipped to apply these skills to boost our careers.

This topic reviews the key concepts covered throughout the course.

55.1 Artificial Intelligence and Generative AI

Artificial Intelligence (AI) is a technology that performs tasks that typically require human intelligence, such as:

- Problem-solving
- Decision-making
- Recognizing speech
- Recognizing images

Generative AI is a branch of Artificial Intelligence that creates new content in different formats. It uses:

- Large Language Models (LLMs) to generate text

- Advanced algorithms to create images, voices, and videos

55.2 AI in Everyday Life

AI influences our daily lives in many ways by improving convenience and personalization.

Examples include:

- Personalizing recommendations on services like **Netflix** and **Spotify**
- Improving efficiency with smart home devices, such as:
 - Smart thermostats
 - Security systems

AI-powered chatbots have also evolved from simple rule-based systems into advanced personal assistants capable of natural conversations.

Some key benefits of AI-powered chatbots include:

- 24/7 availability
- Scalability
- Personalized services
- Multilingual capabilities

55.3 AI Applications Across Industries

AI is used across a wide range of industries, including:

- Manufacturing
- Healthcare
- Finance
- Retail
- Many other sectors

Businesses leverage AI to improve efficiency, automate tasks, and support decision-making.

55.4 Core Concepts of Artificial Intelligence

Some of the core concepts of AI include:

- Machine Learning
- Deep Learning
- Neural Networks

55.4.1 Machine Learning

Machine Learning is a subset of AI that uses algorithms to analyze data and make decisions.

55.4.2 Deep Learning

Deep Learning is a specialized subset of Machine Learning that uses layered algorithms to create neural networks.

This enables advanced tasks such as:

- Natural language understanding
- Image recognition

55.4.3 Neural Networks

Neural networks are inspired by the structure of the human brain.

A neural network consists of:

- An input layer
- One or more hidden layers
- An output layer

55.5 Generative AI Model Architectures

Generative AI models are deep learning models that employ algorithms to generate fresh content.

Common generative AI architectures include:

- Variational Autoencoders (VAEs)
- Generative Adversarial Networks (GANs)
- Autoregressive Models
- Transformers

55.6 AI Domains

The development of AI has driven innovation across several domains, including:

- Natural Language Processing (NLP)
- Computer Vision
- Robotics

55.6.1 Natural Language Processing (NLP)

Natural Language Processing allows machines to interpret and produce human language.

55.6.2 Computer Vision

Computer Vision enables machines to understand visual data by:

- Analyzing images
- Analyzing videos
- Drawing meaningful insights

55.6.3 Robotics

Robotics involves creating and using robots to perform tasks, either with or without human assistance.

Robots typically include:

- Sensors to observe the environment
- Actuators to perform movements
- Controllers to make decisions

Robotics often incorporates:

- Artificial Intelligence
- Machine Learning
- Natural Language Processing (NLP)

to enable automation.

55.7 AI in Business

Artificial Intelligence, including Generative AI, enhances business operations and creates new career opportunities.

Businesses use AI to:

- Create content
- Analyze data for patterns
- Enhance customer service
- Improve product development

55.8 AI Career Opportunities

Popular AI-related careers include:

- AI Engineer
- Data Scientist
- Robotics Engineer
- AI Strategist
- AI Application Developer
- AI Marketing Specialist

To transition into an AI career, we should:

- Build on our current skills

- Learn key AI concepts
- Gain practical experience
- Stay updated with AI trends
- Specialize in emerging AI fields

55.9 Ethical Considerations in AI

Although AI offers numerous benefits, responsible adoption requires careful consideration of ethical issues.

Some essential ethical considerations include:

- Data privacy and security
- Bias and fairness
- Transparency and accountability
- Autonomous systems and human oversight
- Access and equity

56 Glossary of Terms

This alphabetized glossary contains many terms. Understanding these terms is essential when working in the industry and participating in user groups.

56.1 A

Term	Definition
AI algorithms	Programming that tells the computer how to learn to operate on its own, thereby enabling it to attain Artificial Intelligence (AI).
AI-driven predictive maintenance systems	Using AI in production control to monitor and detect missing materials or quality issues.
AI-driven robotics	Robots powered with AI that are augmented with a variety of sensors.
AI-driven energy management systems	Using AI in production control to manage and optimize energy production, distribution, and consumption within energy systems.
AI-powered quality control systems	Using AI in production control to manage quality by enhancing efficiency, accuracy, and decision-making capabilities in various fields.
Augmented intelligence	Technology that enhances human capabilities by providing AI-powered tools and assistance.
Authorization	The process of determining whether an authenticated user has the right to perform an operation.
Automation	The techniques involving technology, programs, robotics, or processes to achieve minimal human intervention.

56.2 B

Term	Definition
Big data stores	A larger, more complex data set, especially from new data sources.

Big data	A dynamic, large, and disparate volume of data being created by people, tools, and machines.
Biometric	Application of statistical analysis to biological data of individuals by means of unique physical characteristics.

56.3 C

Term	Definition
Carbon footprints	The greenhouse gases produced by digital technology resources, devices, tools, and platforms.
Chatbots	Computer programs that simulate human conversation to provide information, answer questions, and perform tasks via text or voice interactions.
Claude models	Advanced AI language models designed for natural language understanding and generation. They are similar to OpenAI's GPT series and are used in applications such as chatbots, content creation, and other AI-driven tasks requiring language processing capabilities.
Cloud computing	A technology that stores and uses data and services over the internet instead of keeping everything on a local computer.
Cobots (Collaborative robots)	AI-driven robots that can work alongside humans in a shared workspace to enhance efficiency and increase productivity. Unlike traditional robots, cobots are designed to interact directly with people, making automation more accessible and versatile.
Cognitive	The term involves intellectual activities such as thinking, reasoning, and problem-solving.
Cognitive computing	A technology that can evaluate an individual's performance and offer personalized recommendations. It provides enhanced functionality, adaptability, and intelligence across various domains.
Cyberattack	An attempt targeting to damage a computer network, computer information system, or personal digital devices.
Cybersecurity	Protection of information technology of individuals and organizations from cyberattacks.

56.4 D

Term	Definition
Data analysis	Process that involves cleaning, transforming, and modeling data to uncover useful information to aid business decision-making.
Data augmentation	The process of generating data from the existing data to train machine learning models.
Dashboards	A tool to present a bird's-eye view of the complete picture while also allowing us to drill down into the next level of information for each parameter. These are easy to comprehend for an average user, make collaboration easy between teams, and allow us to generate reports on the go.
DALL·E model	A multimodal model developed by OpenAI that exhibits the ability to generate images that precisely match the input text or prompt.
Data scientist	Data professionals who develop algorithms, build predictive models, and uncover patterns and trends in large data sets. They apply statistical analysis, especially inferential statistics, machine learning, and predictive modeling, to extract insights from data and make predictions.

Data readiness	The technique involves removing inconsistencies, filling gaps, and ensuring the data is relevant to the problem.
Deepfake	Altering of an image or recording to misrepresent someone as doing or saying something that was not actually done or said.
Deep learning	A specialized subset of machine learning. Deep learning layers algorithms to create a neural network, which is an artificial replication of the brain's structure and functionality.
Digital landscape	The digital landscape encompasses all hardware, software, and content involved in digital advertising, where businesses and customers interact.
Digital data source	The digital location where the data is held in a data table, object, or other storage format.

56.5 E

Term	Definition
E-commerce	A platform for buying and selling products and services, transmitting funds and data over the internet.
Edge computing	The practice of processing data closer to the source of generation rather than relying on centralized data centers.
Edge AI	A type of AI that lives on the device itself rather than relying on the cloud.
Electronic Health Records (EHRs)	A digital version of patient-centered records that makes information available in real time.
Encryption	A method of encoding data so that only authorized individuals can comprehend the information.

56.6 G

Term	Definition
Generative AI	AI that is capable of creating new content (text, images, music, audio, and video) and responding to human conversations.
Generative Variational Autoencoders (VAEs) AI:	VAEs are a type of generative AI model that works by transforming input data through encoding and decoding. They have three main parts: an encoder network, a latent space, and a decoder network.
Generative Adversarial Networks (GANs) AI:	GANs involve two neural networks: the generator and the discriminator. The generator creates new data samples, and the discriminator checks whether the data is real or fake.
Generative Autoregressive models AI:	Autoregressive models create data sequentially, considering the context of earlier generated elements. These models predict the next element in the sequence based on the previous one.
Generative Transformers AI:	Transformers are generally used in Natural Language Processing (NLP). They consist of encoder and decoder layers, enabling the model to effectively generate text sequences or perform cross-language translations.
Generative Unimodal models AI:	Models that process inputs and generate outputs within the same modality (same type of data).
Generative Multimodal models AI:	Models that handle inputs from one modality and produce outputs in a different modality (different types of data).

56.7 H

Term	Definition
------	------------

Hallucinations	A phenomenon where AI can produce outputs that are inaccurate or completely fabricated. This highlights the importance of critical thinking when using AI-generated content.
-----------------------	--

56.8 I

Term	Definition
Internet of Things (IoT)	A network of physical devices connected to the internet that collect and share data for processing and analysis. These devices can be sensors, cameras, or other devices that generate data.
Inventory management	The process of tracking, controlling, and usage of inventory from purchase to the sale of goods.
Interactive learning environment	A dynamic learning setup where learners actively engage with content, instructors, and peers through digital or physical interactions to enhance learning.

56.9 L

Term	Definition
Large Language Model (LLM)	A type of AI program that can recognize and generate text. LLMs are pretrained on large amounts of data.

56.10 M

Term	Definition
Machine learning	The branch of AI and computer science that focuses on using data and algorithms to imitate how humans learn, gradually improving accuracy. It helps systems learn and improve at forecasting, much like people learn from experience.
Machine learning models	Computer programs that train computers to identify patterns within new data and make predictions.
Machine learning engineers	Programmers who construct the algorithms, systems, models, and frameworks that enable machines to learn and perform functions independently and effectively.
Meta's Llama models	An accessible, open-source LLM designed for developers, researchers, and businesses to build, experiment, and responsibly scale their generative AI ideas.

56.11 N

Term	Definition
Natural Language Processing (NLP)	Machine learning technology that enables computers to interpret and comprehend human language.
Neural networks	Computational models influenced by the human brain's neural structure.
Neurons in AI	An artificial neural network consists of interconnected nodes known as neurons. The neurons take incoming data, like the human brain's neurological network, and learn to make decisions over time.
Neural network: Perceptron	The simplest type of artificial neural network consisting of only input and output layers.
Neural network: Feed-forward	A type of artificial neural network in which information flows in one direction, that is, in the forward direction.

Neural network: Deep feed-forward	A type of neural network similar to the feed-forward network but with more than one hidden layer.
Neural network: Modular	A type of neural network that combines two or more neural networks to arrive at the output.
Neural network: Convolutional Neural Network (CNN)	A type of neural network particularly well-suited for analyzing visual data.
Neural network: Recurrent Neural Networks (RNNs)	A type of neural network where the neurons in hidden layers receive an input with a specific delay in time.

56.12 O

Term	Definition
IBM watsonx Assistant	A conversational AI solution that empowers a broader audience, including non-technical business users, to effortlessly build generative AI assistants.
OpenAI	A private AI research laboratory that aims to develop and direct AI in ways that produce services.

56.13 P

Term	Definition
Predictive analytics	A method that predicts future outcomes by using historical data combined with statistical modeling, data mining techniques, and machine learning.

56.14 R

Term	Definition
Rehabilitation robots	Robots that help patients recover mobility and strength, offering personalized therapy.
Reinforcement learning	A category of machine learning technology that trains software to make decisions to achieve the most optimal results.
Robo-advisors	Online applications that use AI algorithms to offer automated, algorithm-driven investment suggestions, portfolio management, and financial planning services.
Robots	Machines with complex systems made up of several key components such as sensors, actuators, and controllers.
Robotics	A technology that involves designing, constructing, and operating robots, where the machines can perform tasks by themselves or with some help.
Robotic Process Automation (RPA)	A type of computer software that helps create, use, and control virtual robots. These robots act like humans when they work with digital systems and software.

56.15 S

Term	Definition
Salesforce Marketing Cloud	An AI-driven marketing automation platform that uses machine learning algorithms to segment customers, target them with personalized messages, and optimize campaign performance.
Sentiment analysis	Processing digital text using AI to analyze and determine whether the emotional tone of the message is positive, negative, or neutral.
Smart grids	An electricity network that uses digital technologies to monitor and manage the transport of electricity from sources to meet the varying electricity demands of end users.

Smart home	A technology that allows homeowners to control appliances, thermostats, lights, and other devices remotely using a smartphone or tablet through an internet connection.
Speech-to-text (STT)	A technology that changes spoken words into written text using neural networks.
Supervised learning	A category of machine learning technology that uses labeled datasets to train algorithms to predict outcomes and recognize patterns.
Strong AI (Generalized AI)	A type of AI system that replicates advanced human functions, such as reasoning, planning, and problem-solving.
Super AI (Conscious AI)	A type of AI system that is superintelligent by being self-aware and intelligent enough to surpass the cognitive abilities of humans.
Statistical analysis	The process of collecting large volumes of data and then using statistics and other data analysis techniques to identify trends, patterns, and insights.
Statistical methods	Methods useful to ensure that data is interpreted correctly and apparent relationships are meaningful and not just happening by chance.

56.16 T

Term	Definition
Text-to-speech (TTS)	A type of assistive "read aloud" technology that reads digital text aloud.

56.17 U

Term	Definition
Unsupervised learning	A category of machine learning technology that uses algorithms to analyze and cluster unlabeled data sets.

56.18 V

Term	Definition
Virtual assistance	AI-powered software applications designed to assist users with various tasks, such as setting reminders, providing information, and managing schedules.
Voice assistance	AI-powered assistants that can be interacted with using voice commands. Users can speak to these assistants to perform tasks, ask questions, and control other devices.
Video analytics	The use of advanced AI and machine learning algorithms to monitor, analyze, and manage large volumes of video.

56.19 W

Term	Definition
Weak AI (Narrow AI)	AI that is focused on one narrow task.